UNIVERSITY OF WATERLOO

Big Data Infrastructure

CS 489/698 Big Data Infrastructure (Winter 2017)

Week 4: Analyzing Text (1/2)

January 24, 2017

Jimmy Lin
David R. Cheriton School of Computer Science
University of Waterloo

These slides are available at http://lintool.github.io/bigdata-2017w/

Structure of the Course

"Core" framework features
and algorithm design

Data-Parallel Dataflow Languages

We have a collection of records,
want to apply a bunch of operations
to compute some result

What are the dataflow operators?
Spark is a better MapReduce with a few more "niceties"!

Moving forward: generic reference to "mapper" and "reducers"

Structure of the Course

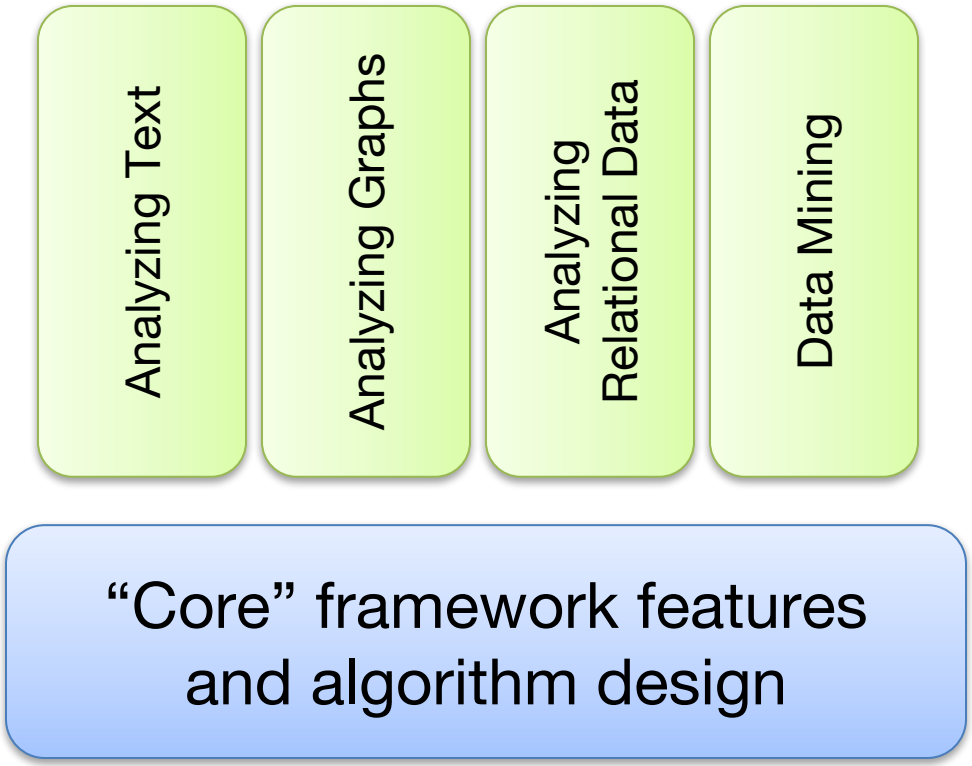

Source: http://www.flickr.com/photos/guvnah/7861418602/

Count
(Efficiently)

```
class Mapper {
  def map(key: Long, value: Text, context: Context) = {
    for (word <- tokenize(value)) {
      context.write(word, 1)
    }
  }
}

class Reducer {
  def reduce(key: Text, values: Iterable[Int], context: Context) = {
    for (value <- values) {
      sum += value
    }
    context.write(key, sum)
  }
}
```

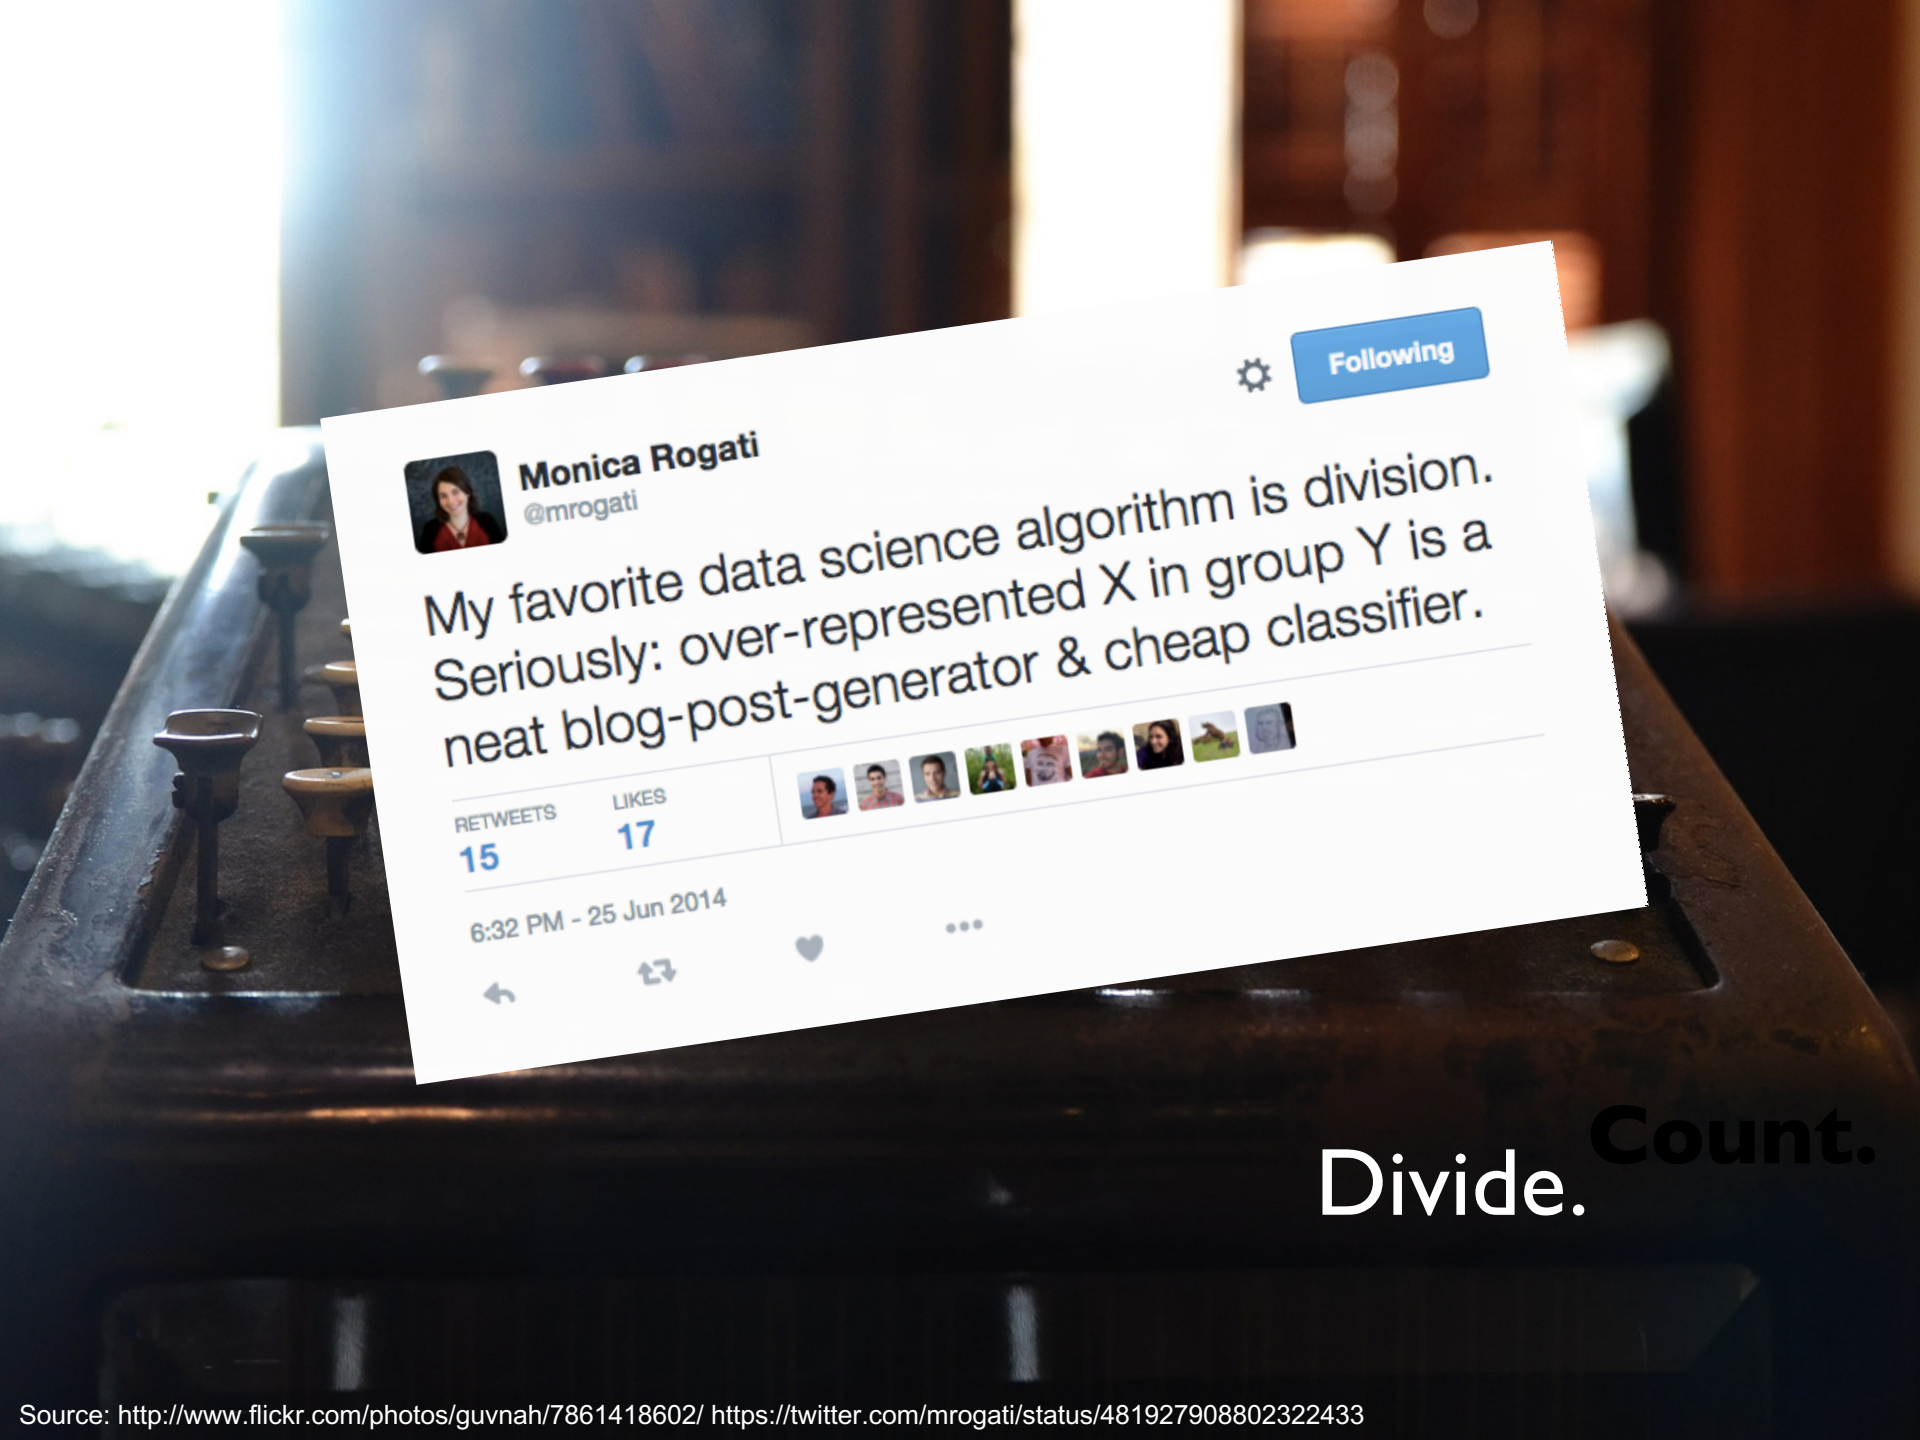

Source: http://www.flickr.com/photos/guvnah/7861418602/ https://twitter.com/mrogati/status/481927908802322433

Pairs. Stripes.

Seems pretty trivial…

More than a “toy problem”?

Answer: language models

Language Models

$$P(w_1, w_2, \ldots, w_T)$$

What are they?
How do we build them?
How are they useful?

Language Models

$$P(w_1, w_2, \ldots, w_T)$$

$$= P(w_1)P(w_2|w_1)P(w_3|w_1, w_2) \ldots P(w_T|w_1, \ldots, w_{T-1})$$

[chain rule]

Is this tractable?

Approximating Probabilities: *N*-Grams

Basic idea: limit history to fixed number of (*N* – 1) words (Markov Assumption)

$$P(w_k|w_1,\ldots,w_{k-1}) \approx P(w_k|w_{k-N+1},\ldots,w_{k-1})$$

N=1: Unigram Language Model

$$P(w_k|w_1,\ldots,w_{k-1}) \approx P(w_k)$$

$$\Rightarrow P(w_1,w_2,\ldots,w_T) \approx P(w_1)P(w_2)\ldots P(w_T)$$

Approximating Probabilities: *N*-Grams

Basic idea: limit history to fixed number of (*N* – 1) words (Markov Assumption)

$$P(w_k|w_1,\ldots,w_{k-1}) \approx P(w_k|w_{k-N+1},\ldots,w_{k-1})$$

N=2: Bigram Language Model

$$P(w_k|w_1,\ldots,w_{k-1}) \approx P(w_k|w_{k-1})$$

$$\Rightarrow P(w_1,w_2,\ldots,w_T) \approx P(w_1|<\mathrm{S}>)P(w_2|w_1)\ldots P(w_T|w_{T-1})$$

Approximating Probabilities: *N*-Grams

Basic idea: limit history to fixed number of (*N* – 1) words (Markov Assumption)

$$P(w_k|w_1,\ldots,w_{k-1}) \approx P(w_k|w_{k-N+1},\ldots,w_{k-1})$$

N=3: Trigram Language Model

$$P(w_k|w_1,\ldots,w_{k-1}) \approx P(w_k|w_{k-2},w_{k-1})$$

$$\Rightarrow P(w_1,w_2,\ldots,w_T) \approx P(w_1|<\mathrm{S}><\mathrm{S}>)\ldots P(w_T|w_{T-2}w_{T-1})$$

Building *N*-Gram Language Models

Compute maximum likelihood estimates (MLE) for Individual *n*-gram probabilities

Unigram $$P(w_i) = \frac{C(w_i)}{N}$$

Fancy way of saying: count + divide

Bigram $$P(w_i, w_j) = \frac{C(w_i, w_j)}{N}$$

$$P(w_j|w_i) = \frac{P(w_i, w_j)}{P(w_i)} = \frac{C(w_i, w_j)}{\sum_w C(w_i, w)} \stackrel{?}{=} \frac{C(w_i, w_j)}{C(w_i)}$$

Minor detail here…

Generalizes to higher-order n-grams
State of the art models use ~5-grams

We already know how to do this in MapReduce!

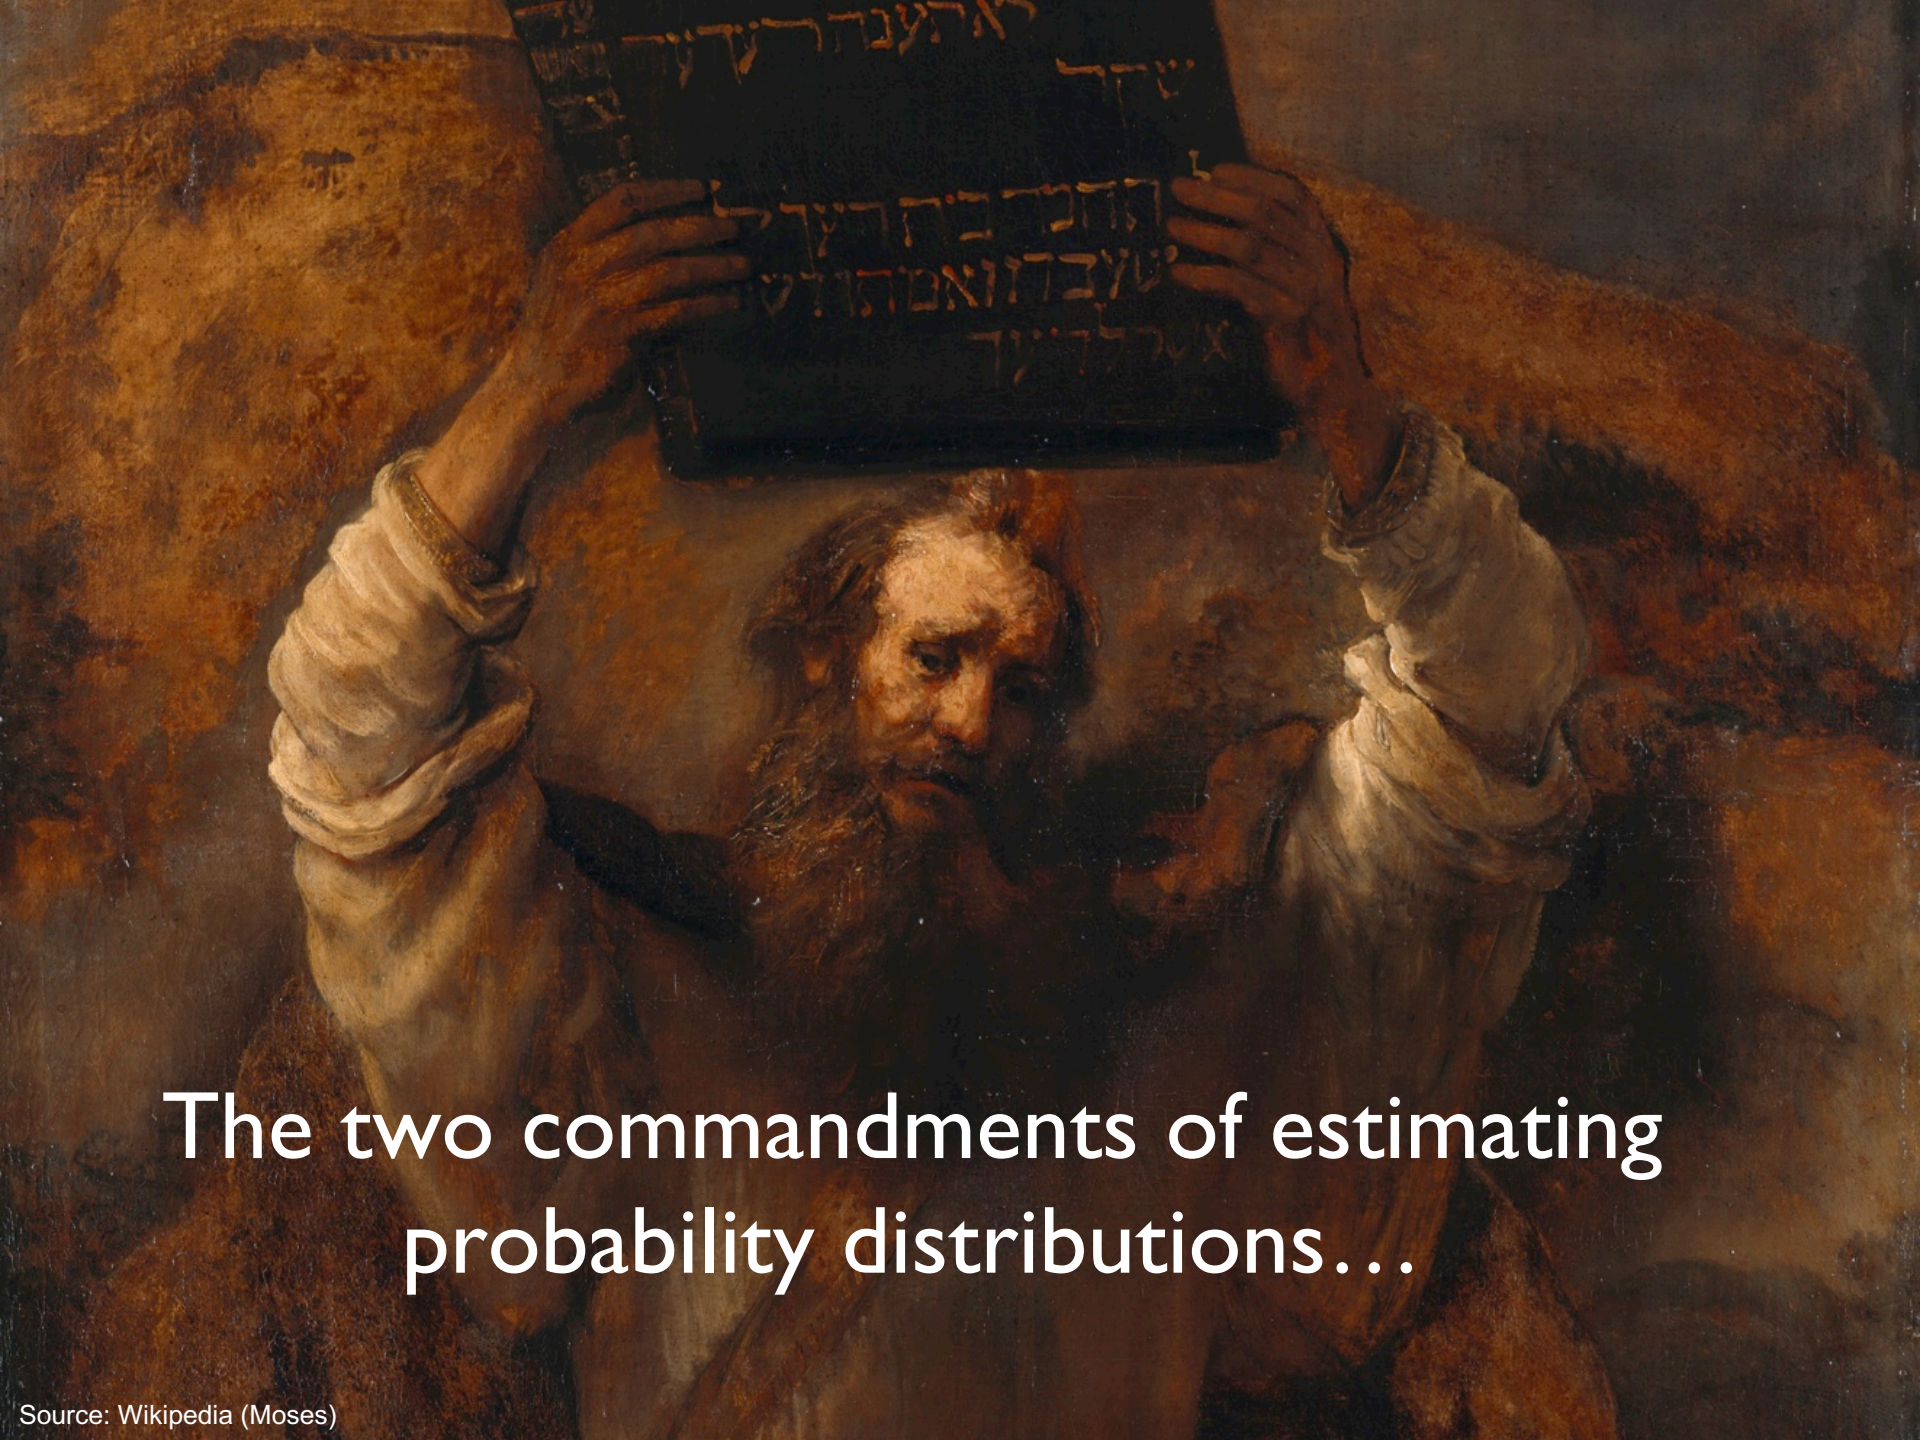

The two commandments of estimating probability distributions…

Source: Wikipedia (Moses)

Probabilities must sum up to one

Source: http://www.flickr.com/photos/37680518@N03/7746322384/

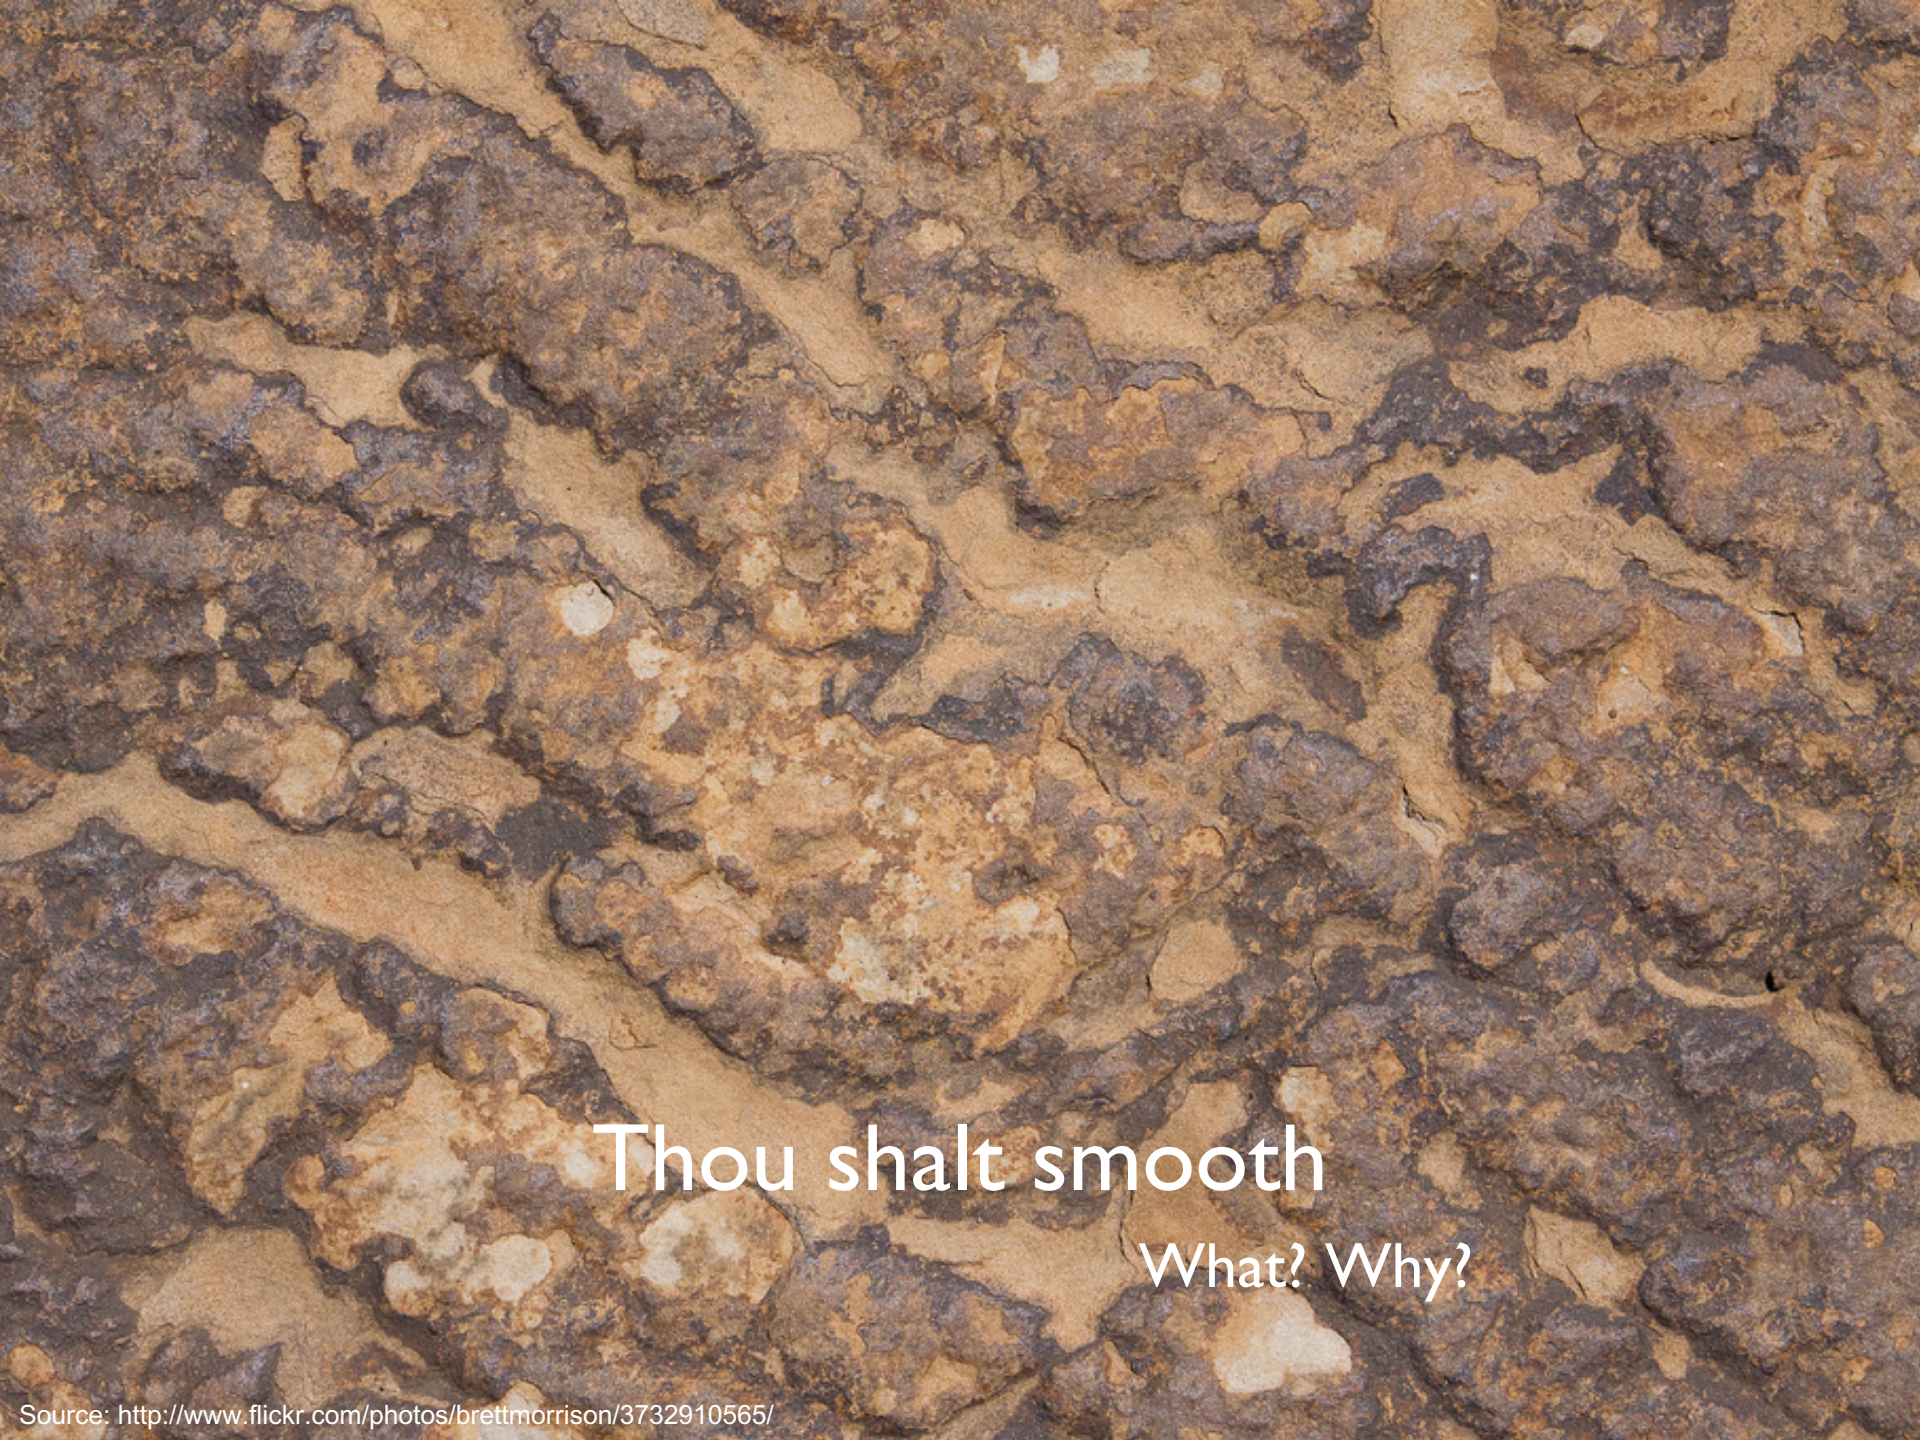

Source: http://www.flickr.com/photos/brettmorrison/3732910565/

Source: https://www.flickr.com/photos/avlxyz/6898001012/

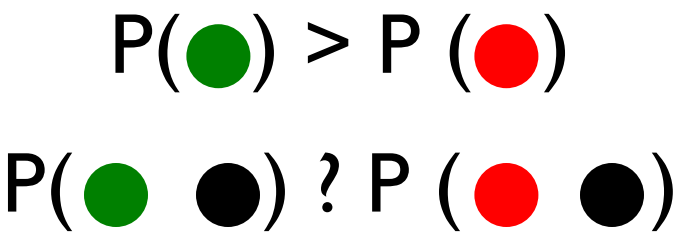
P( ) > P ( )
P(●) ? P (● ●)

Example: Bigram Language Model

<s> I am Sam </s>
<s> Sam I am </s>
<s> I do not like green eggs and ham </s>

Training Corpus

P(I | <s>) = 2/3 = 0.67
P(am | I) = 2/3 = 0.67
P(</s> | Sam)= 1/2 = 0.50
...

P(Sam | <s>) = 1/3 = 0.33
P(do | I) = 1/3 = 0.33
P(Sam | am) = 1/2 = 0.50

Bigram Probability Estimates

Note: We don't ever cross sentence boundaries

Data Sparsity

P(I | <s>) = 2/3 = 0.67
P(am | I) = 2/3 = 0.67
P(</s> | Sam)= 1/2 = 0.50
...

P(Sam | <s>) = 1/3 = 0.33
P(do | I) = 1/3 = 0.33
P(Sam | am) = 1/2 = 0.50

Bigram Probability Estimates

P(I like ham)

= P(I | <s>) P(like | I) P(ham | like) P(</s> | ham)

= 0

Why is this bad?

Issue: Sparsity!

Thou shalt smooth!

Zeros are bad for any statistical estimator

Need better estimators because MLEs give us a lot of zeros

A distribution without zeros is "smoother"

The Robin Hood Philosophy: Take from the rich (seen *n*-grams) and give to the poor (unseen *n*-grams)

Need better estimators because MLEs give us a lot of zeros

A distribution without zeros is "smoother"

Lots of techniques:

Laplace, Good-Turing, Katz backoff, Jelinek-Mercer

Kneser-Ney represents best practice

Laplace Smoothing

Learn fancy words
for simple ideas!

Simplest and oldest smoothing technique

Just add 1 to all *n*-gram counts including the unseen ones

So, what do the revised estimates look like?

Laplace Smoothing

Unigrams

$$P_{MLE}(w_i) = \frac{C(w_i)}{N} \longrightarrow P_{LAP}(w_i) = \frac{C(w_i) + 1}{N + V}$$

Bigrams

$$P_{MLE}(w_i, w_j) = \frac{C(w_i, w_j)}{N} \longrightarrow P_{LAP}(w_i, w_j) = \frac{C(w_i, w_j) + 1}{N + V^2}$$

Careful, don't confuse the *N*'s!

What if we don't know *V*?

Jelinek-Mercer Smoothing: Interpolation

Mix higher-order with lower-order models to defeat sparsity

Mix = Weighted Linear Combination

$$P(w_k|w_{k-2}w_{k-1}) =$$

$$\lambda_1 P(w_k|w_{k-2}w_{k-1}) + \lambda_2 P(w_k|w_{k-1}) + \lambda_3 P(w_k)$$

$$0 <= \lambda_i <= 1 \qquad \sum_i \lambda_i = 1$$

Kneser-Ney Smoothing

Interpolate discounted model with a special "continuation" *n*-gram model

Based on appearance of *n*-grams in different contexts

Excellent performance, state of the art

$$P_{KN}(w_k|w_{k-1}) = \frac{C(w_{k-1}w_k) - D}{C(w_{k-1})} + \beta(w_k)P_{CONT}(w_k)$$

$$P_{CONT}(w_i) = \frac{N(\bullet\ w_i)}{\sum_{w'} N(\bullet\ w')}$$

$N(\bullet\ w_i)$ = number of different contexts w_i has appeared in

Kneser-Ney Smoothing: Intuition

I can't see without my _____

"San Francisco" occurs a lot

I can't see without my Francisco?

Stupid Backoff

Let's break all the rules:

$$S(w_i|w_{i-k+1}^{i-1}) = \begin{cases} \frac{f(w_{i-k+1}^{i})}{f(w_{i-k+1}^{i-1})} & \text{if } f(w_{i-k+1}^{i}) > 0 \\ \alpha S(w_i|w_{i-k+2}^{i-1}) & \text{otherwise} \end{cases}$$

$$S(w_i) = \frac{f(w_i)}{N}$$

But throw *lots* of data at the problem!

Source: Brants et al. (EMNLP 2007)

Source: Wikipedia (Moses)

Stupid Backoff Implementation: Pairs!

Straightforward approach: count each order separately

A B ← remember this value
A B C S(C|A B) = f(A B C)/f(A B)
A B D S(D|A B) = f(A B D)/f(A B)
A B E S(E|A B) = f(A B E)/f(A B)
… …

More clever approach: count all orders together

A B ← remember this value
A B C ← remember this value
A B C P
A B C Q
A B D ← remember this value
A B D X
A B D Y
…

Stupid Backoff: Additional Optimizations

Replace strings with integers

Assign ids based on frequency (better compression using vbyte)

Partition by bigram for better load balancing

Replicate all unigram counts


State of the art smoothing (less data)
vs. Count and divide (more data)
EVERLAST
Source: Wikipedia (Boxing)

Statistical Machine Translation

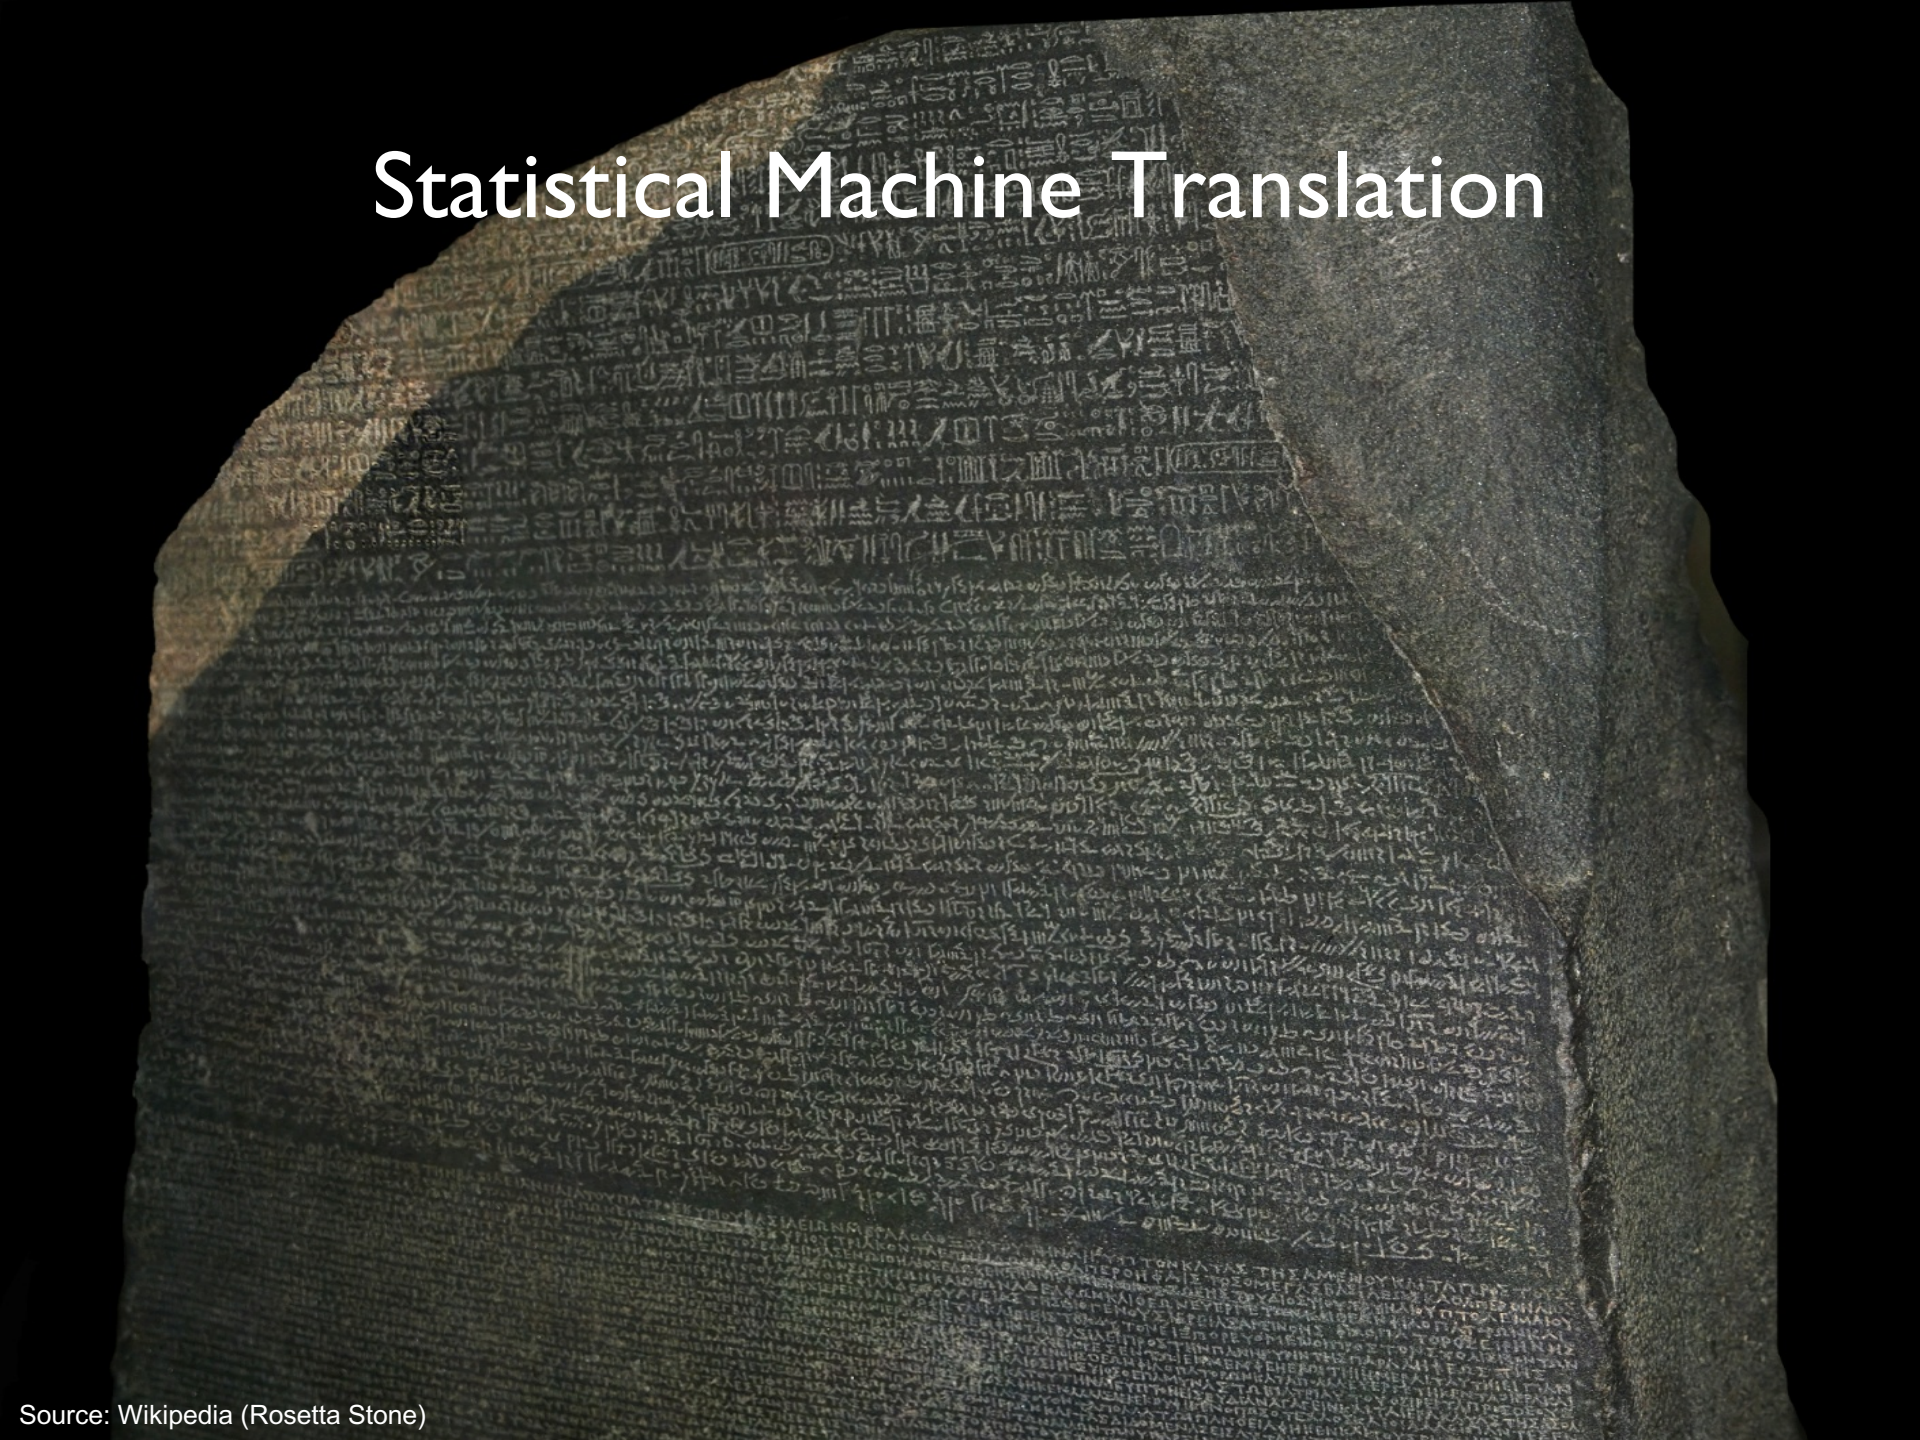

Source: Wikipedia (Rosetta Stone)

Statistical Machine Translation

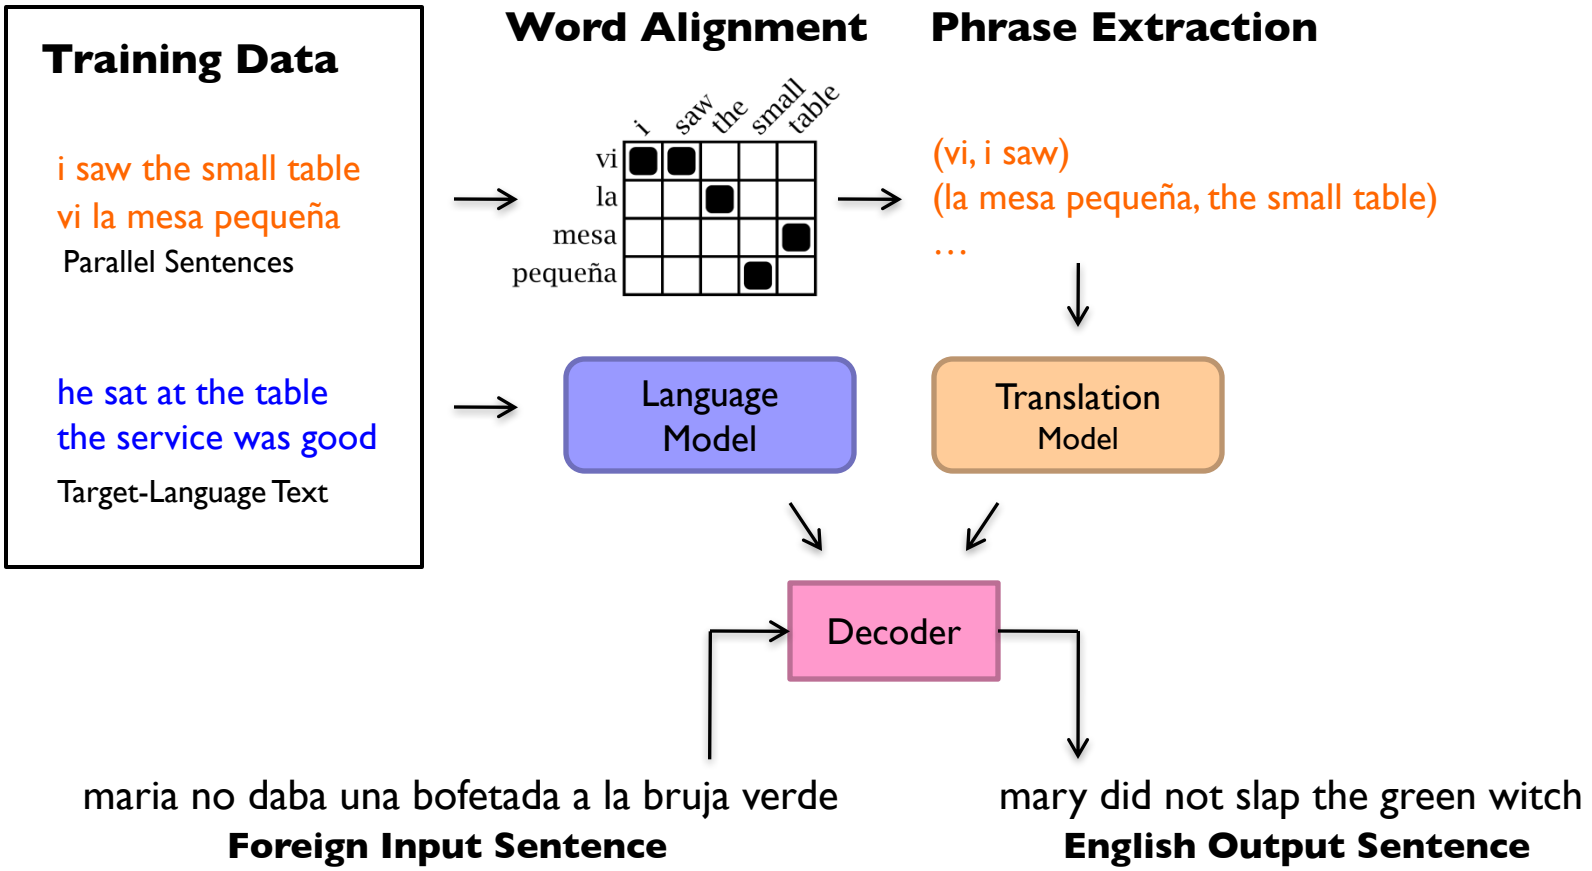

$$\hat{e}_1^I = \arg\max_{e_1^I} \left[P(e_1^I \mid f_1^J) \right] = \arg\max_{e_1^I} \left[P(e_1^I) P(f_1^J \mid e_1^I) \right]$$

Translation as a Tiling Problem

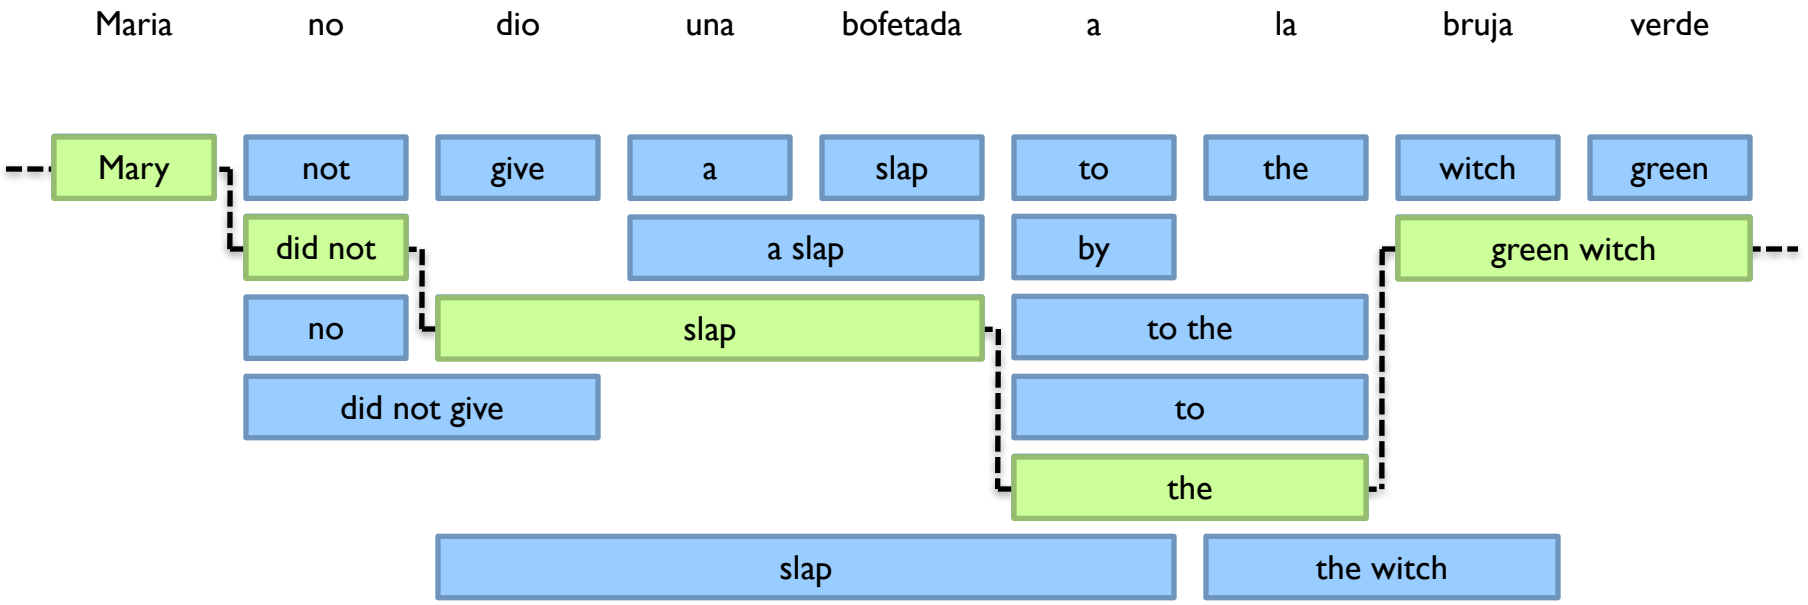

$$\hat{e}_1^I = \arg\max_{e_1^I} \left[P(e_1^I \mid f_1^J) \right] = \arg\max_{e_1^I} \left[P(e_1^I) P(f_1^J \mid e_1^I) \right]$$

Results: Running Time

	target	***webnews***	***web***
# tokens	237M	31G	1.8T
vocab size	200k	5M	16M
# n-grams	257M	21G	300G
LM size (SB)	2G	89G	1.8T
time (SB)	20 min	8 hours	1 day
time (KN)	2.5 hours	2 days	–
# machines	100	400	1500

Source: Brants et al. (EMNLP 2007)

Results: Translation Quality

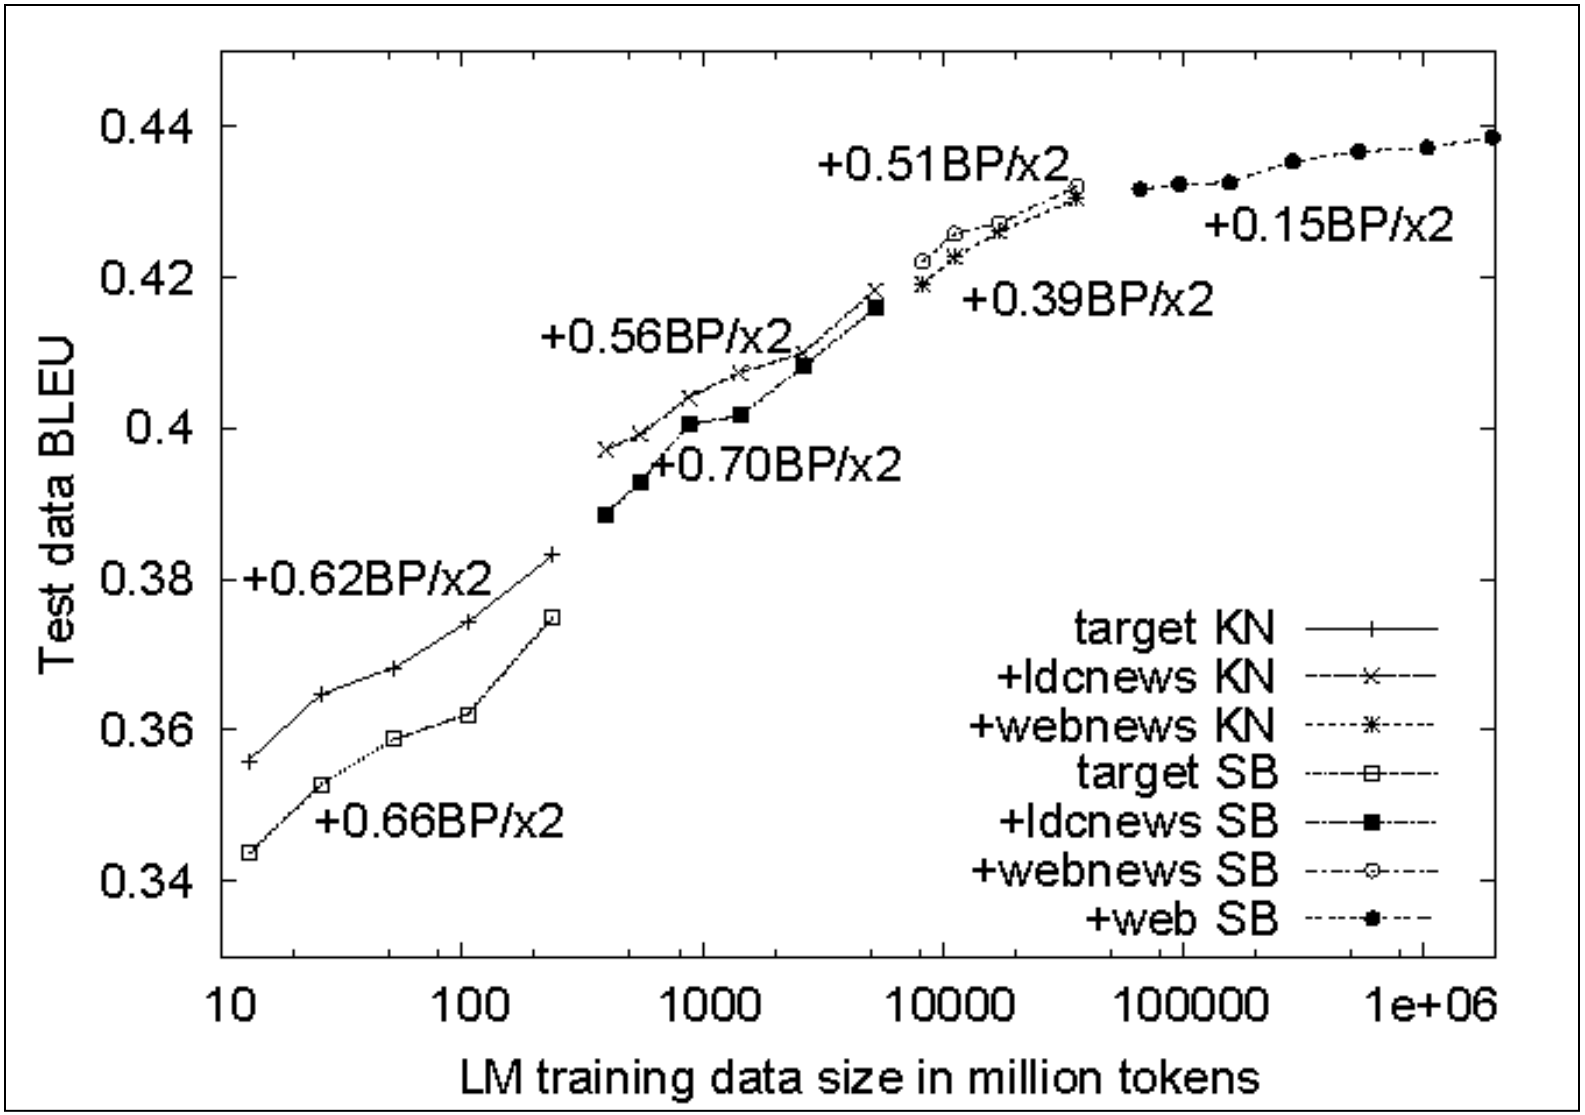

Source: Brants et al. (EMNLP 2007)

What's actually going on?

$$P(e|f) = \frac{P(e) \cdot P(f|e)}{P(f)}$$

$$\hat{e} = \arg\max_{e} P(e)P(f|e)$$

Source: http://www.flickr.com/photos/johnmueller/3814846567/in/pool-56226199@N00/

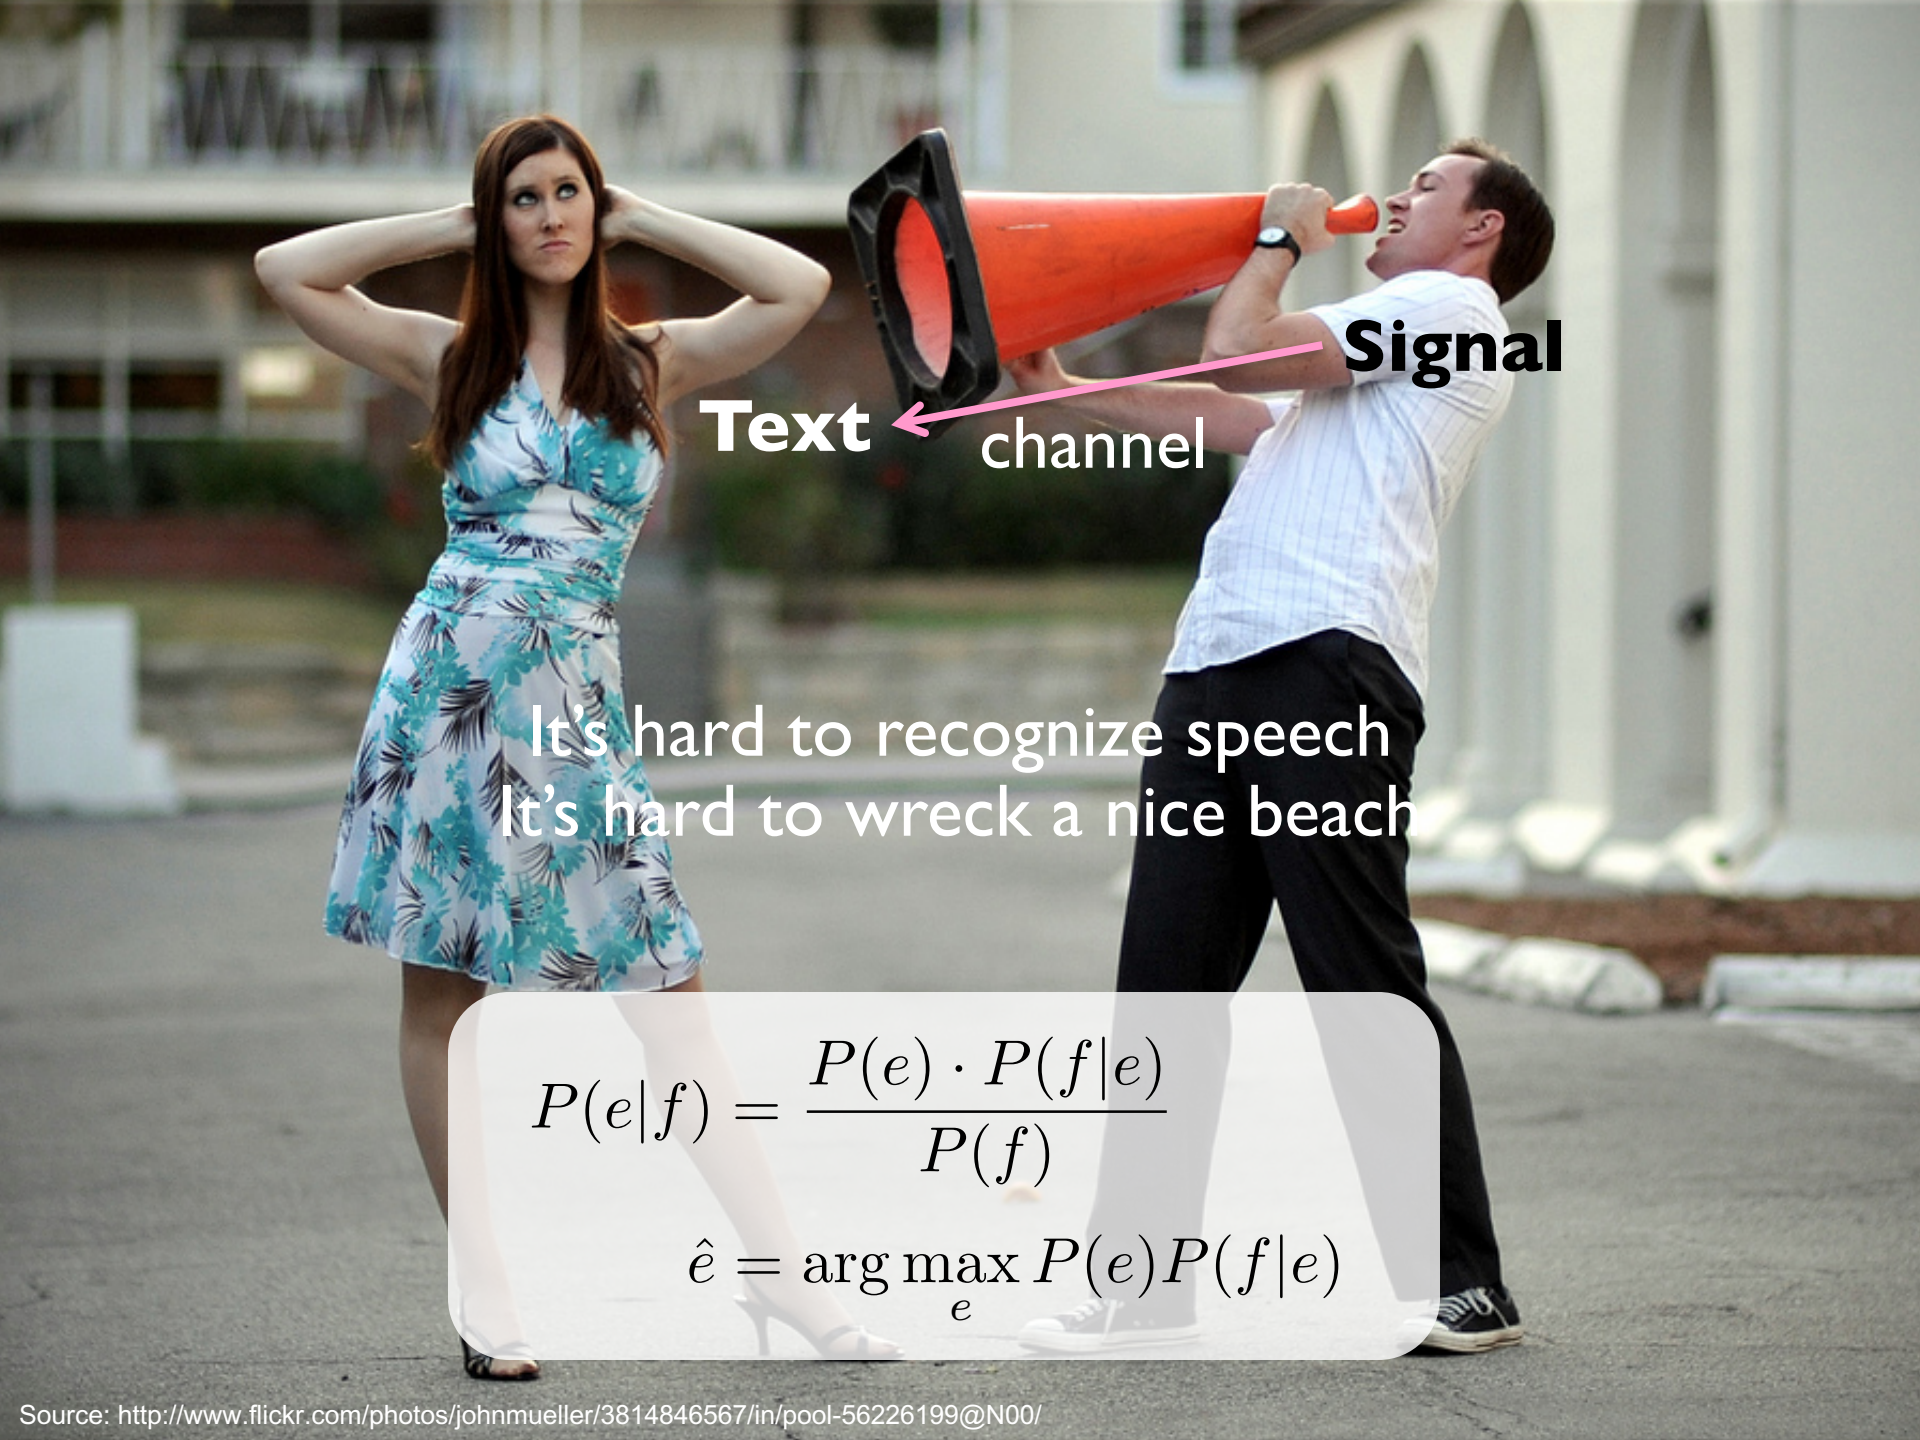

Source: http://www.flickr.com/photos/johnmueller/3814846567/in/pool-56226199@N00/

$$P(e|f) = \frac{P(e) \cdot P(f|e)}{P(f)}$$

$$\hat{e} = \arg\max_{e} P(e)P(f|e)$$

Source: http://www.flickr.com/photos/johnmueller/3814846567/in/pool-56226199@N00/

Neural Networks

Have taken over…

Source: http://www.flickr.com/photos/guvnah/7861418602/

First, nomenclature…

Search and information retrieval (IR)

Focus on textual information (= text/document retrieval)

Other possibilities include image, video, music, …

What do we search?

Generically, "collections"

Less-frequently used, "corpora"

What do we find?

Generically, "documents"

Though "documents" may refer to web pages, PDFs, PowerPoint, etc.

The Central Problem in Search

Searcher

Author

Concepts

Concepts

Query Terms

Document Terms

"tragic love story"

"fateful star-crossed romance"

Do these represent the same concepts?

Abstract IR Architecture

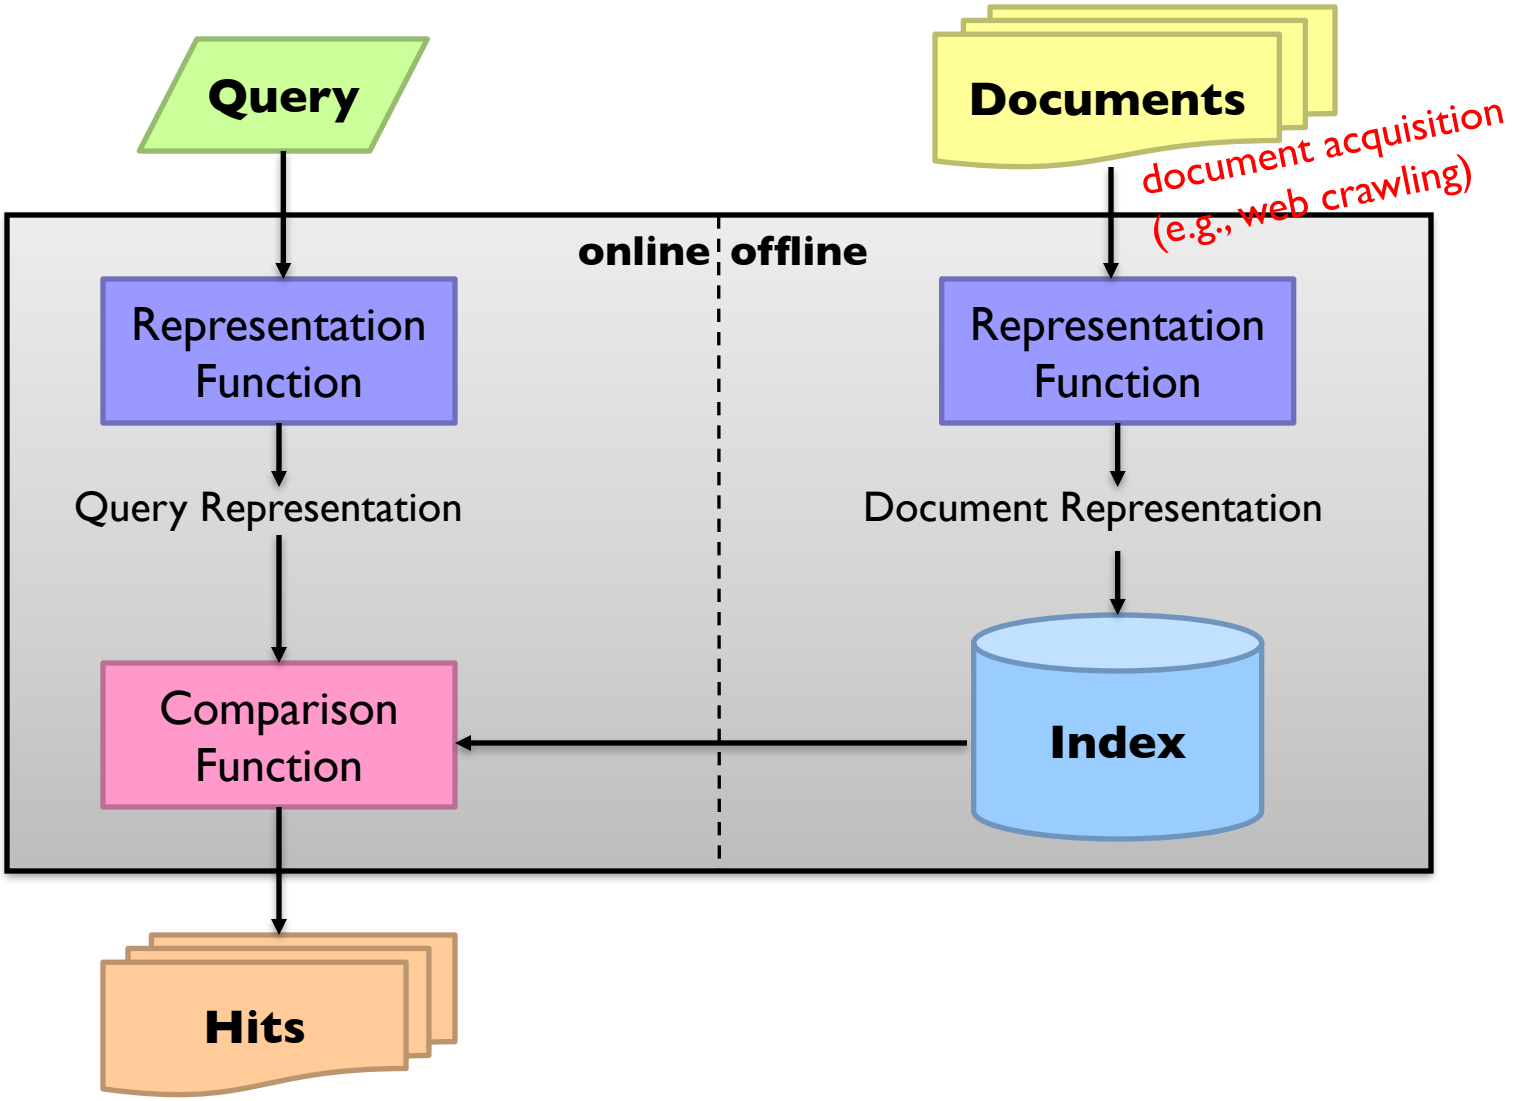

How do we represent text?

Remember: computers don't "understand" anything!

"Bag of words"

Treat all the words in a document as index terms
Assign a "weight" to each term based on "importance"
(or, in simplest case, presence/absence of word)
Disregard order, structure, meaning, etc. of the words
Simple, yet effective!

Assumptions

Term occurrence is independent
Document relevance is independent
"Words" are well-defined

What's a word?

天主教教宗若望保祿二世因感冒再度住進醫院。
這是他今年第二度因同樣的病因住院。

الناطق باسم - وقال مارك ريجيف
إن شارون قبل - الخارجية الإسرائيلية
الدعوة وسيقوم للمرة الأولى بزيارة
تونس، التي كانت لفترة طويلة المقر
1982. الرسمي لمنظمة التحرير الفلسطينية بعد خروجها من لبنان عام

Выступая в Мещанском суде Москвы экс-глава ЮКОСа заявил не совершал ничего противозаконного, в чем обвиняет его генпрокуратура России.

भारत सरकार ने आर्थकि सर्वेक्षण में वत्तिीय वर्ष 2005-06 में सात फ़ीसदी वकिास दर हासलि करने का आकलन कयिा है और कर सुधार पर ज़ोर दयिा है

日米連合で台頭中国に対処…アーミテージ前副長官提言

조재영 기자= 서울시는 25일 이명박 시장이 `행정중심복합도시" 건설안에 대해 `군대라도 동원해 막고싶은 심정"이라고 말했다는 일부 언론의 보도를 부인했다.

Sample Document

McDonald's slims down spuds

Fast-food chain to reduce certain types of fat in its french fries with new cooking oil.

NEW YORK (CNN/Money) - McDonald's Corp. is cutting the amount of "bad" fat in its french fries nearly in half, the fast-food chain said Tuesday as it moves to make all its fried menu items healthier.

But does that mean the popular shoestring fries won't taste the same? The company says no. "It's a win-win for our customers because they are getting the same great french-fry taste along with an even healthier nutrition profile," said Mike Roberts, president of McDonald's USA.

But others are not so sure. McDonald's will not specifically discuss the kind of oil it plans to use, but at least one nutrition expert says playing with the formula could mean a different taste.

Shares of Oak Brook, Ill.-based McDonald's (MCD: down $0.54 to $23.22, Research, Estimates) were lower Tuesday afternoon. It was unclear Tuesday whether competitors Burger King and Wendy's International (WEN: down $0.80 to $34.91, Research, Estimates) would follow suit. Neither company could immediately be reached for comment.

…

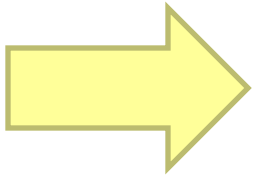

"Bag of Words"

14 × McDonalds

12 × fat

11 × fries

8 × new

7 × french

6 × company, said, nutrition

5 × food, oil, percent, reduce, taste, Tuesday

…

Counting Words…

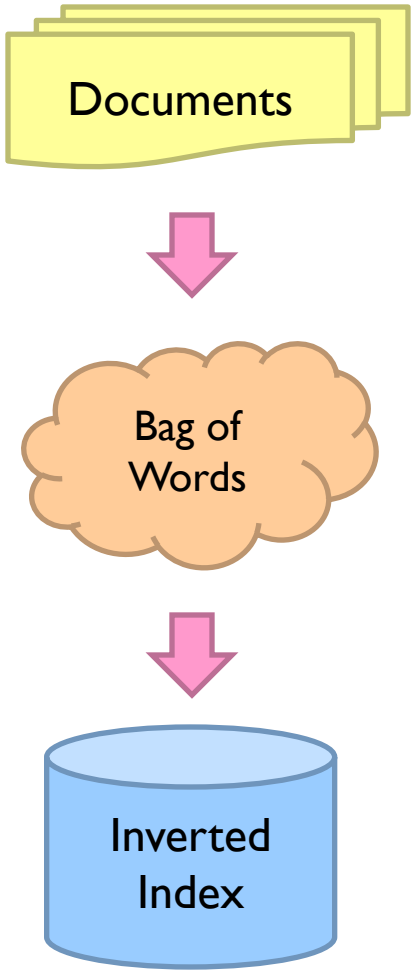

case folding, tokenization, stopword removal, stemming

Source: http://www.flickr.com/photos/guvnah/7861418602/

Doc 1
one fish, two fish

Doc 2
red fish, blue fish

Doc 3
cat in the hat

Doc 4
green eggs and ham

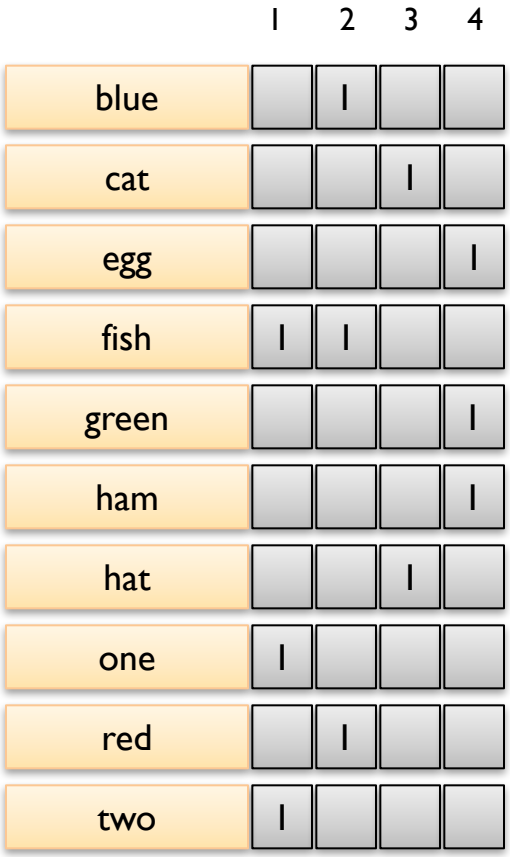

What goes in each cell?

- boolean
- count
- positions

Abstract IR Architecture

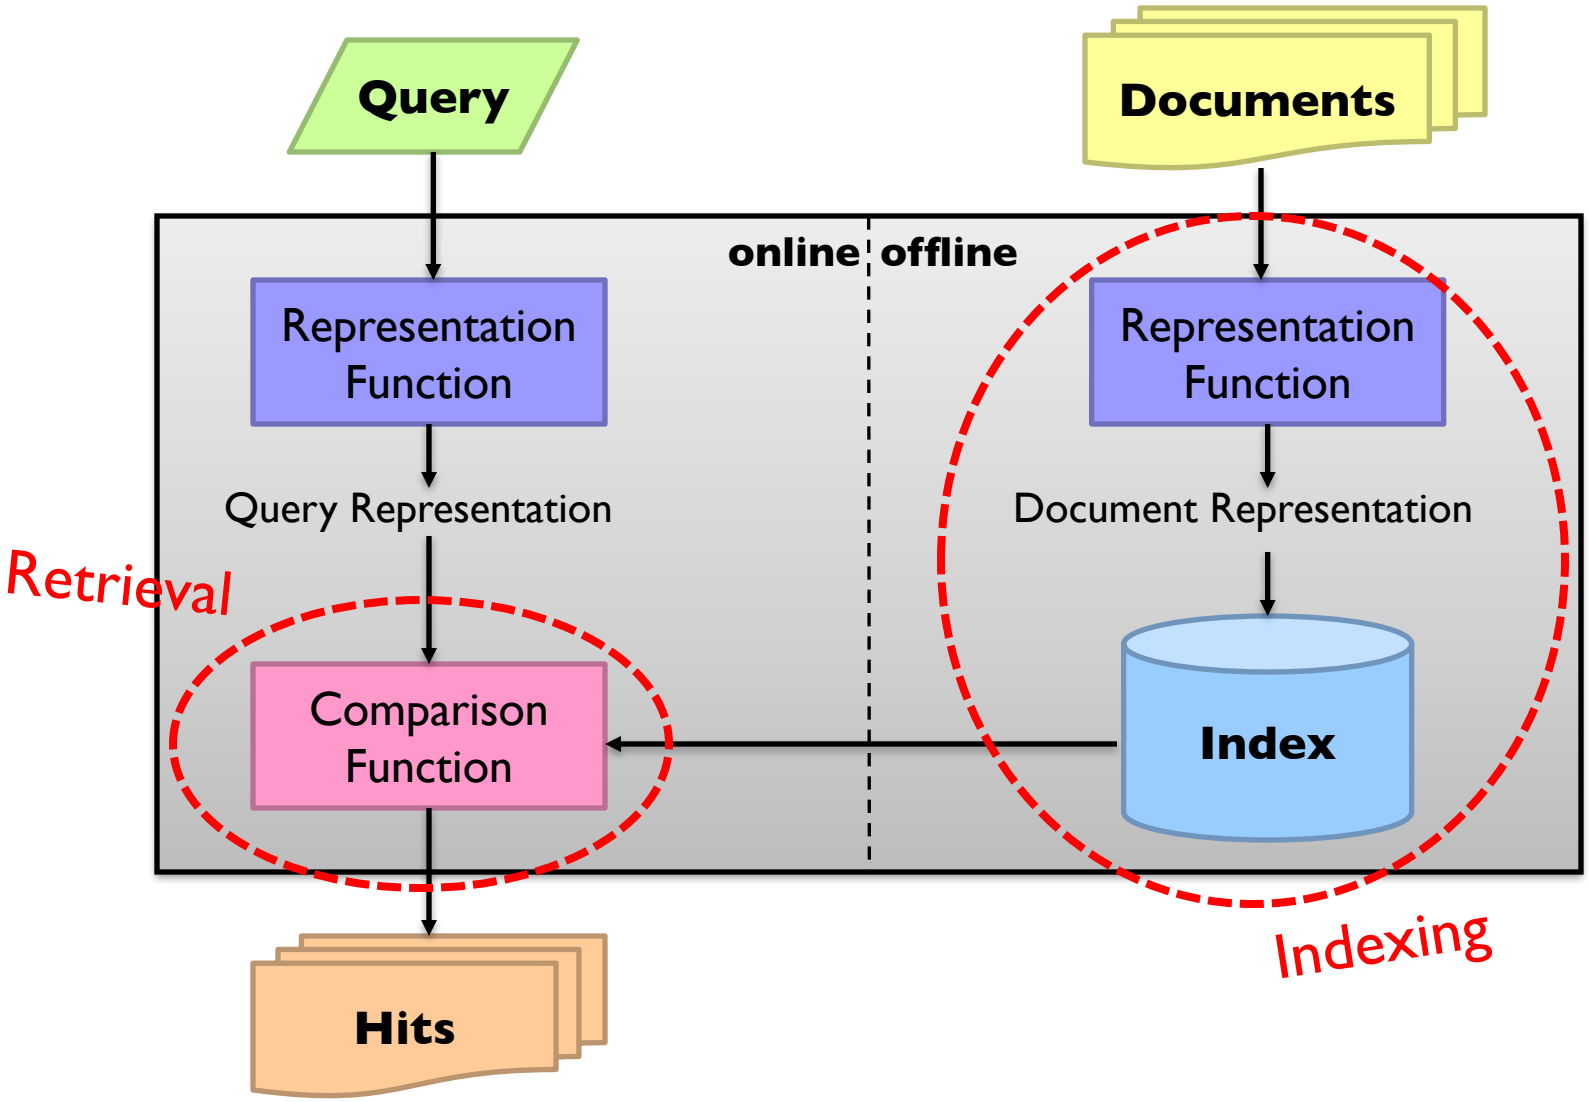

Doc 1 one fish, two fish

Doc 2 red fish, blue fish

Doc 3 cat in the hat

Doc 4 green eggs and ham

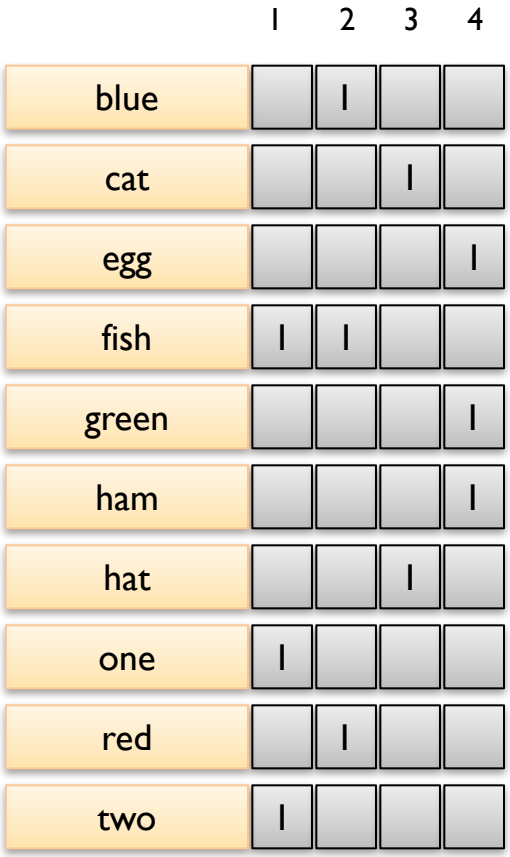

Indexing: building this structure

Retrieval: manipulating this structure

Where have we seen this before?

Doc 1
one fish, two fish

Doc 2
red fish, blue fish

Doc 3
cat in the hat

Doc 4
green eggs and ham

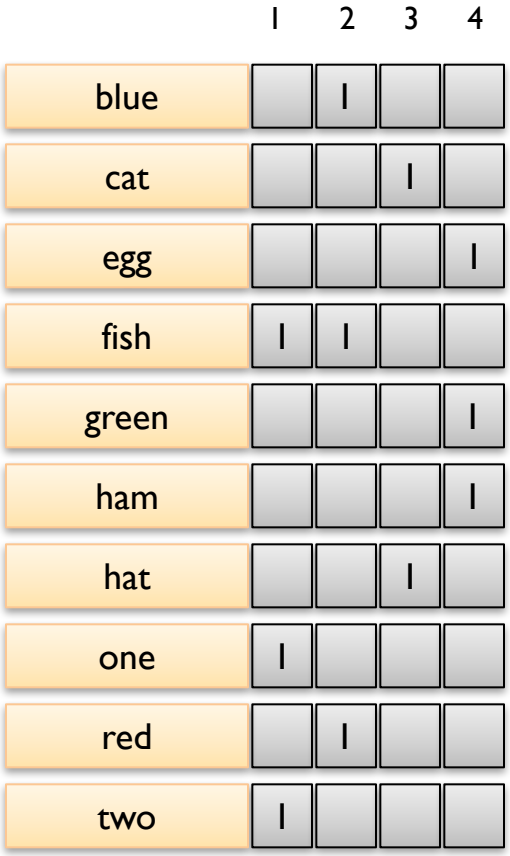

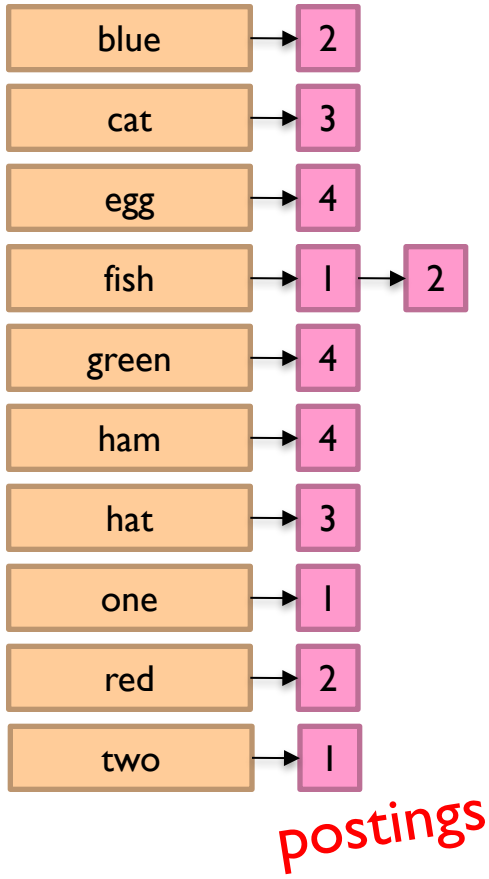

postings lists

Indexing: Performance Analysis

Fundamentally, a large sorting problem

- Terms usually fit in memory
- Postings usually don't

How is it done on a single machine?

How can it be done with MapReduce?

First, let's characterize the problem size:

- Size of vocabulary
- Size of postings

Vocabulary Size: Heaps' Law

$$M = kT^b$$

M is vocabulary size
T is collection size (number of documents)
k and *b* are constants

Typically, *k* is between 30 and 100, *b* is between 0.4 and 0.6

Heaps' Law: linear in log-log space

Surprise: Vocabulary size grows unbounded!

Heaps' Law for RCV1

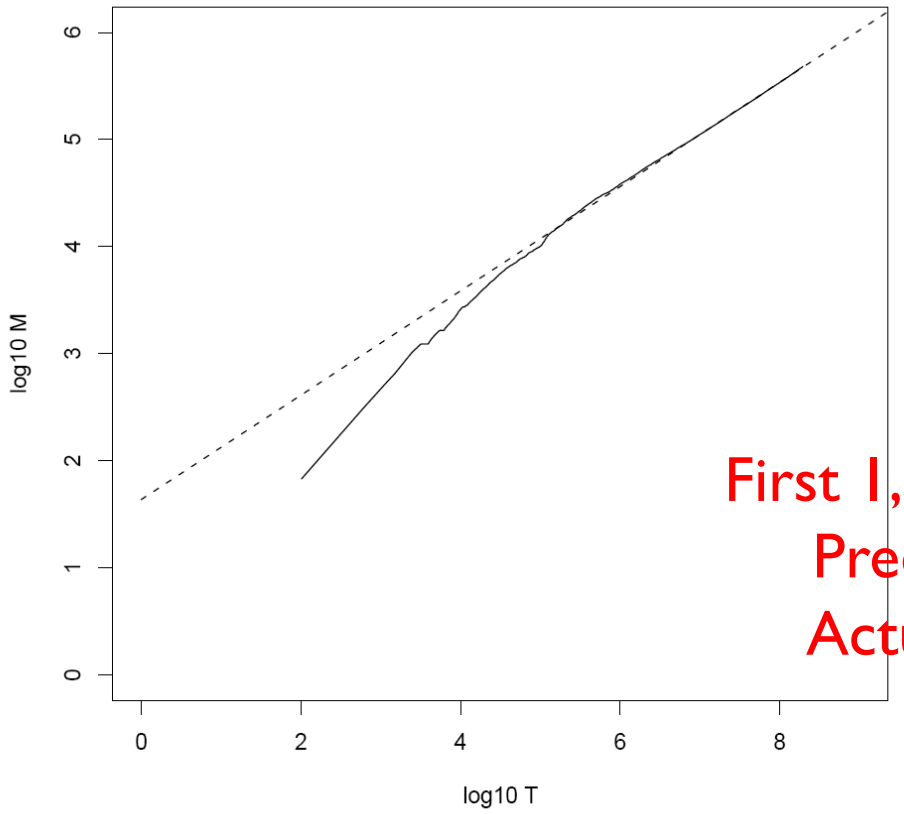

k = 44
b = 0.49

First 1,000,020 terms:
Predicted = 38,323
Actual = 38,365

Reuters-RCV1 collection: 806,791 newswire documents (Aug 20, 1996-August 19, 1997)

Postings Size: Zipf's Law

$$f(k; s, N) = \frac{1/k^s}{\sum_{n=1}^{N}(1/n^s)}$$

N	number of elements
k	rank
s	characteristic exponent

Zipf's Law: (also) linear in log-log space

Specific case of Power Law distributions

In other words:

A few elements occur very frequently

Many elements occur very infrequently

Zipf's Law for RCV1

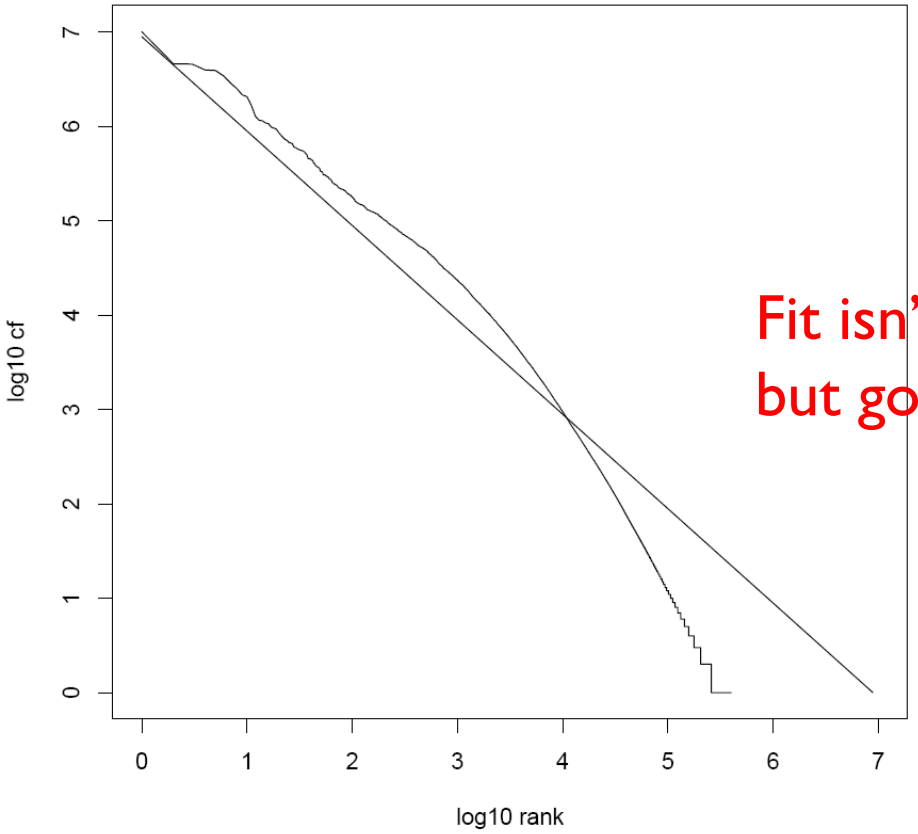

Fit isn't that good…
but good enough!

Reuters-RCV1 collection: 806,791 newswire documents (Aug 20, 1996-August 19, 1997)

Zipf's Law for Wikipedia

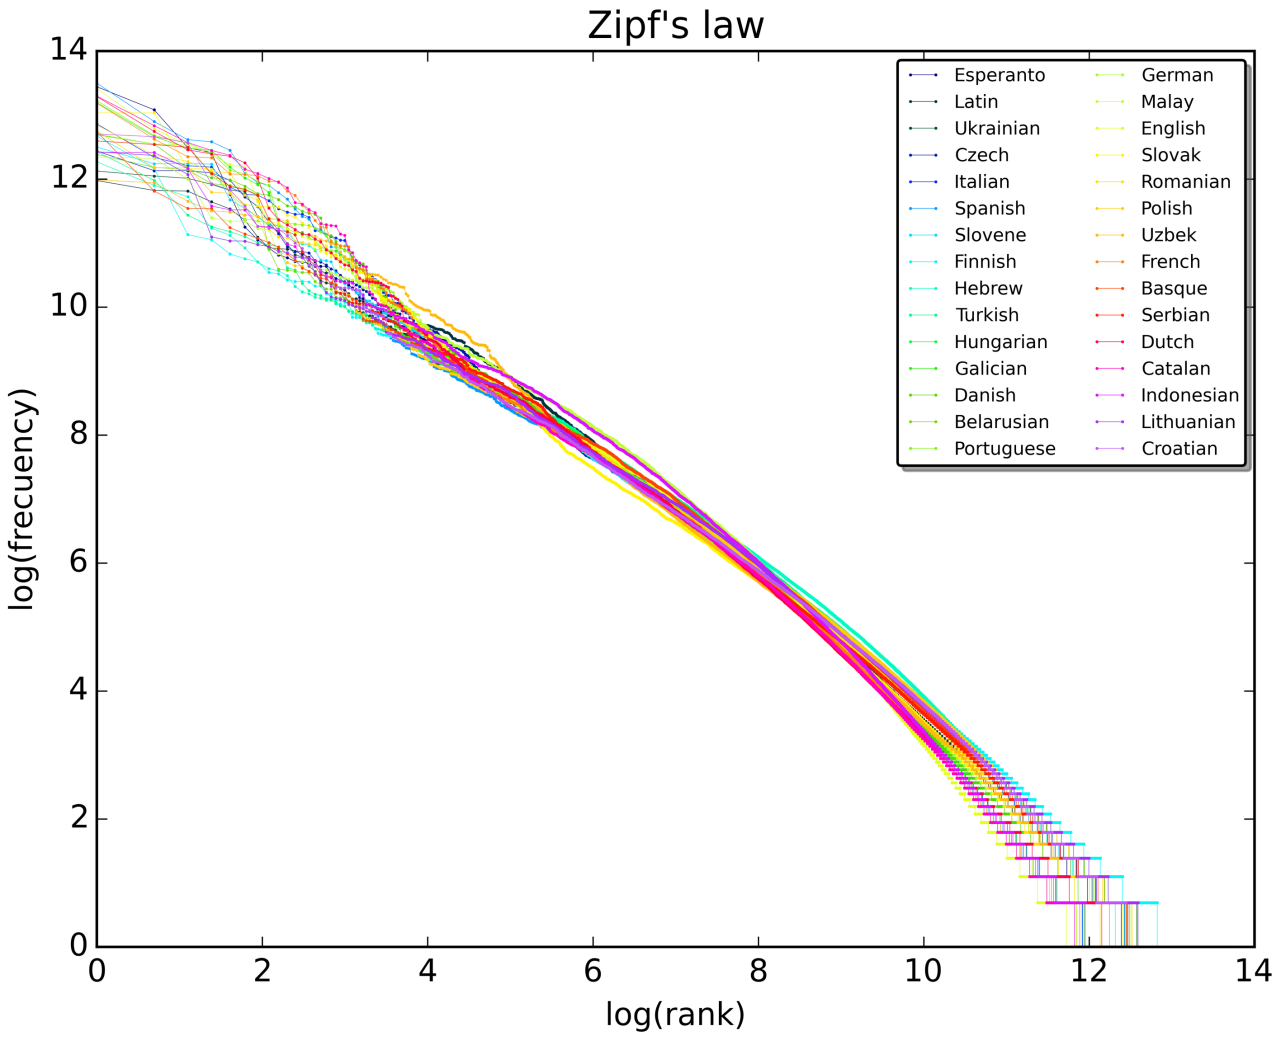

Rank versus frequency for the first 10m words in 30 Wikipedias (dumps from October 2015)

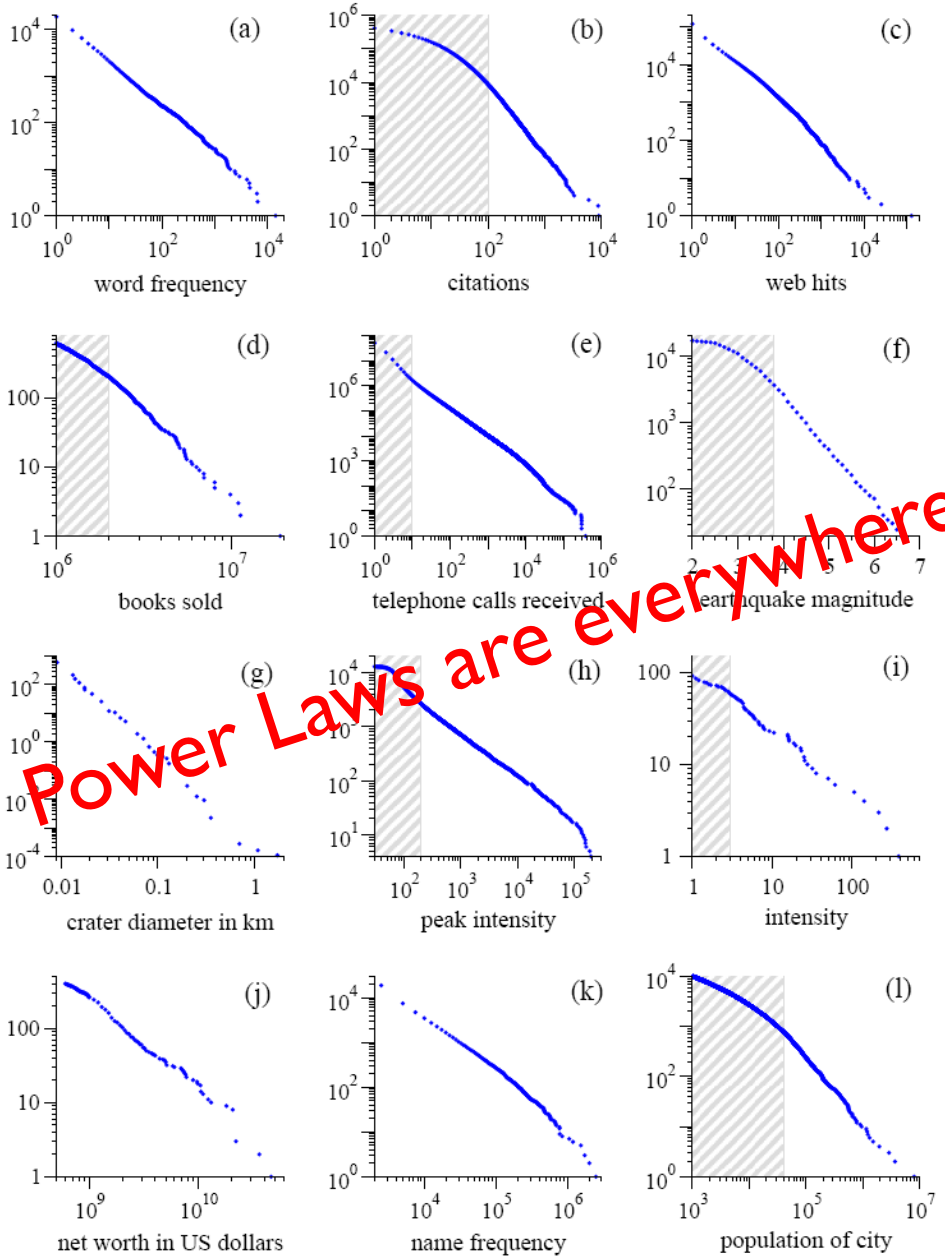

Figure from: Newman, M. E. J. (2005) “Power laws, Pareto distributions and Zipf's law.” Contemporary Physics 46:323–351.

MapReduce: Index Construction

Map over all documents

Emit *term* as key, (*docno*, *tf*) as value

Emit other information as necessary (e.g., term position)

Sort/shuffle: group postings by term

Reduce

Gather and sort the postings (e.g., by *docno* or *tf*)

Write postings to disk

MapReduce does all the heavy lifting!

Inverted Indexing with MapReduce

Doc 1
one fish, two fish

Doc 2
red fish, blue fish

Doc 3
cat in the hat

Map

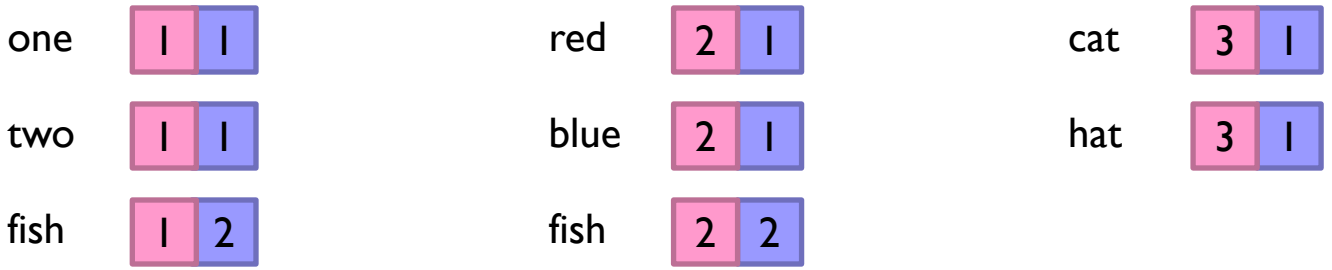

Shuffle and Sort: aggregate values by keys

Reduce

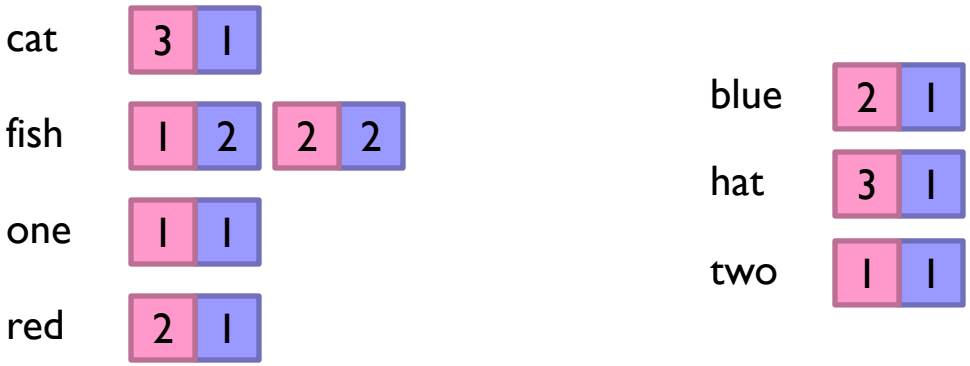

Inverted Indexing: Pseudo-Code

```
1: class Mapper
2:     method Map(docid n, doc d)
3:         H ← new AssociativeArray                    ▷ histogram to hold term frequencies
4:         for all term t ∈ doc d do      ▷ processes the doc, e.g., tokenization and stopword removal
5:             H{t} ← H{t} + 1
6:         for all term t ∈ H do
7:             Emit(term t, posting ⟨n, H{t}⟩)                ▷ emits individual postings
```

```
1: class Reducer
2:     method Reduce(term t, postings [⟨n1, f1⟩ ...])
3:         P ← new List
4:         for all ⟨n, f⟩ ∈ postings [⟨n1, f1⟩ ...] do
5:             P.Append(⟨n, f⟩)                          ▷ appends postings unsorted
6:         P.Sort()                                      ▷ sorts for compression
7:         Emit(term t, postingsList P)
```

What's the problem?

Stay tuned…

Source: Wikipedia (Japanese rock garden)