



UNIVERSITY OF
WATERLOO

Big Data Infrastructure

CS 489/698 Big Data Infrastructure (Winter 2017)

Week 10: Mutable State (1/2)

March 14, 2017

Jimmy Lin

David R. Cheriton School of Computer Science
University of Waterloo

These slides are available at <http://lintool.github.io/bigdata-2017/>



This work is licensed under a Creative Commons Attribution-Noncommercial-Share Alike 3.0 United States
See <http://creativecommons.org/licenses/by-nc-sa/3.0/us/> for details

Structure of the Course

Analyzing Text

Analyzing Graphs

Analyzing
Relational Data

Data Mining

“Core” framework features
and algorithm design

The Fundamental Problem

We want to keep track of *mutable* state in a *scalable* manner

Assumptions:

State organized in terms of logical records

State unlikely to fit on single machine, must be distributed

MapReduce won't do!

Want more? Take a *real* distributed systems course!

The Fundamental Problem

We want to keep track of *mutable* state in a *scalable* manner

Assumptions:

State organized in terms of logical records

State unlikely to fit on single machine, must be distributed

Uh... just use an RDBMS?

What do RDBMSes provide?

Relational model with schemas

Powerful, flexible query language

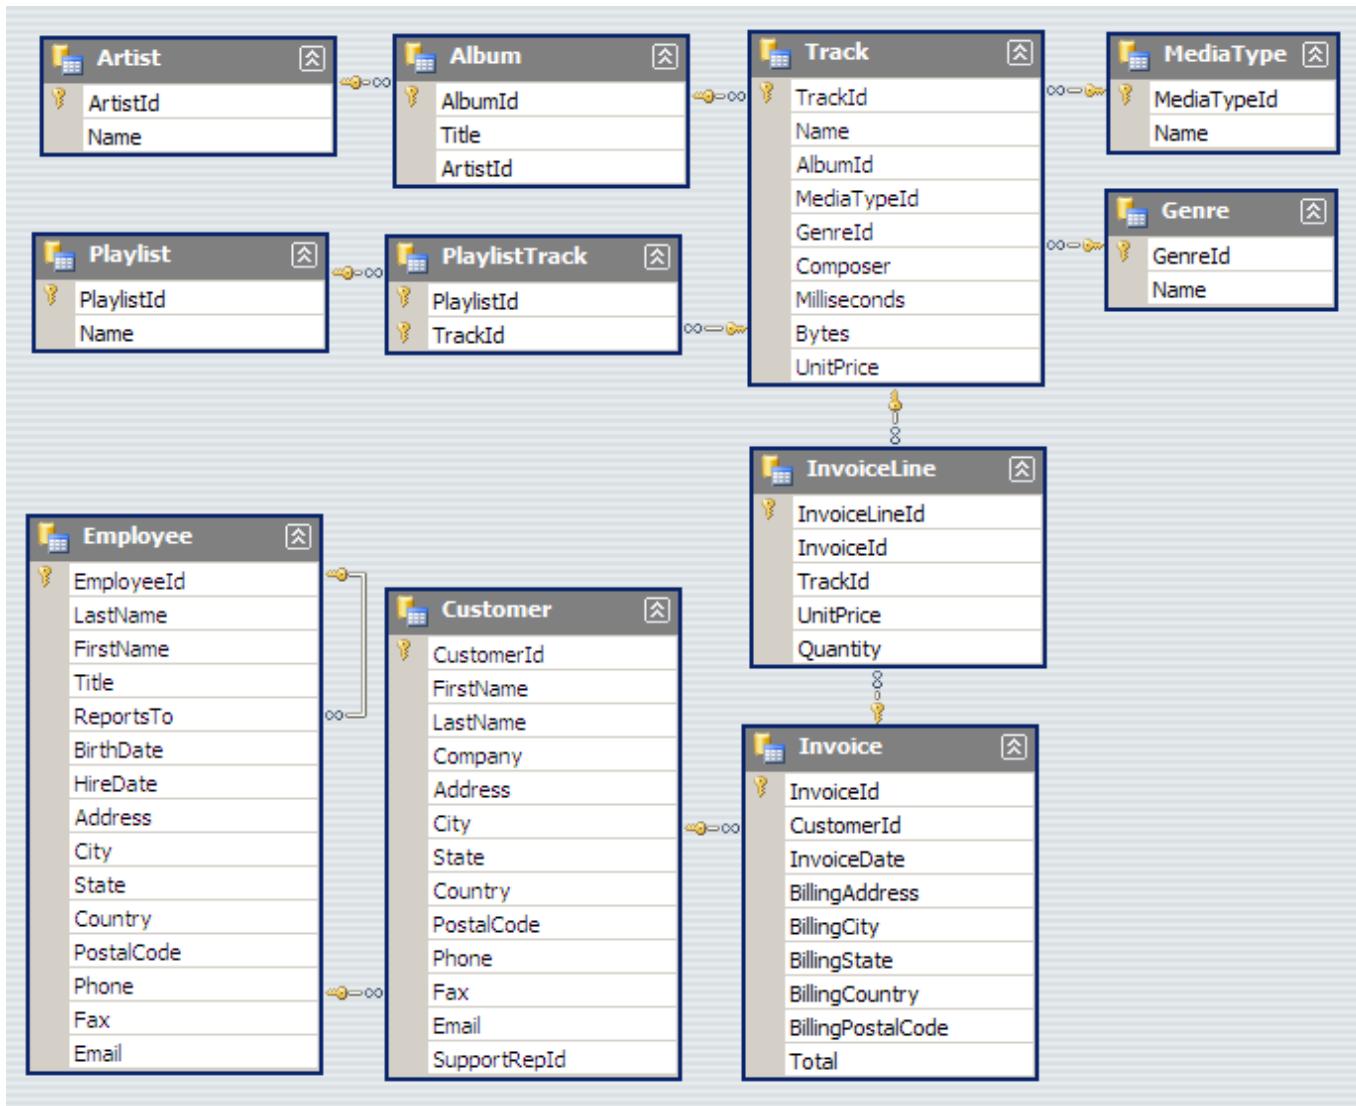
Transactional semantics: ACID

Rich ecosystem, lots of tool support

RDBMSes: Pain Points



#1: Must design up front, painful to evolve



Consistent field names?

{

```
"token": 945842,  
"feature_enabled": "super_special",  
"userid": 229922,  
"page": "null",  
"info": { "email": "my@place.com" }}
```

}

This should really be a list...



Is this really an integer?



Is this really null?



What keys? What values?

JSON to the Rescue!
Flexible design doesn't mean *no* design!

#2: Pay for ACID!



#3: Cost!



What do RDBMSes provide?

Relational model with

Powerful, flexible query

Transactional semantics

Rich ecosystem, lots of

What if we want *a la carte*?



Features *a la carte*?

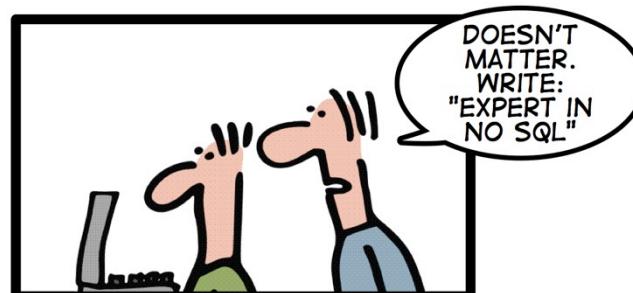
What if I'm willing to give up consistency for scalability?

What if I'm willing to give up the relational model for flexibility?

What if I just want a cheaper solution?

Enter... NoSQL!

HOW TO WRITE A CV



Leverage the NoSQL boom

NoSQL (Not only SQL)

1. Horizontally scale “simple operations”
2. Replicate/distribute data over many servers
3. Simple call interface
4. Weaker concurrency model than ACID
5. Efficient use of distributed indexes and RAM
6. Flexible schemas

But, don't blindly follow the hype...
Often, MySQL is what you really need!

“web scale”



(Major) Types of NoSQL databases

Key-value stores

Column-oriented databases

Document stores

Graph databases

Three Core Ideas

Keeping track of the partitions?

Partitioning (sharding)

To increase scalability and to decrease latency

Consistency?

Replication

To increase robustness (availability) and to increase throughput

Consistency?

Caching

To reduce latency

Key-Value Stores



Key-Value Stores: Data Model

Stores associations between keys and values

Keys are usually primitives

For example, ints, strings, raw bytes, etc.

Values can be primitive or complex: often opaque to store

Primitives: ints, strings, etc.

Complex: JSON, HTML fragments, etc.

Key-Value Stores: Operations

Very simple API:

Get – fetch value associated with key

Put – set value associated with key

Optional operations:

Multi-get

Multi-put

Range queries

Secondary index lookups

Consistency model:

Atomic puts (usually)

Cross-key operations: who knows?

Key-Value Stores: Implementation

Non-persistent:

Just a big in-memory hash table

Examples: Redis, memcached

Persistent

Wrapper around a traditional RDBMS

Examples: Voldemort

What if data doesn't fit on a single machine?

Simple Solution: Partition!

Partition the key space across multiple machines

Let's say, hash partitioning

For n machines, store key k at machine $h(k) \bmod n$

Okay... But:

How do we know which physical machine to contact?

How do we add a new machine to the cluster?

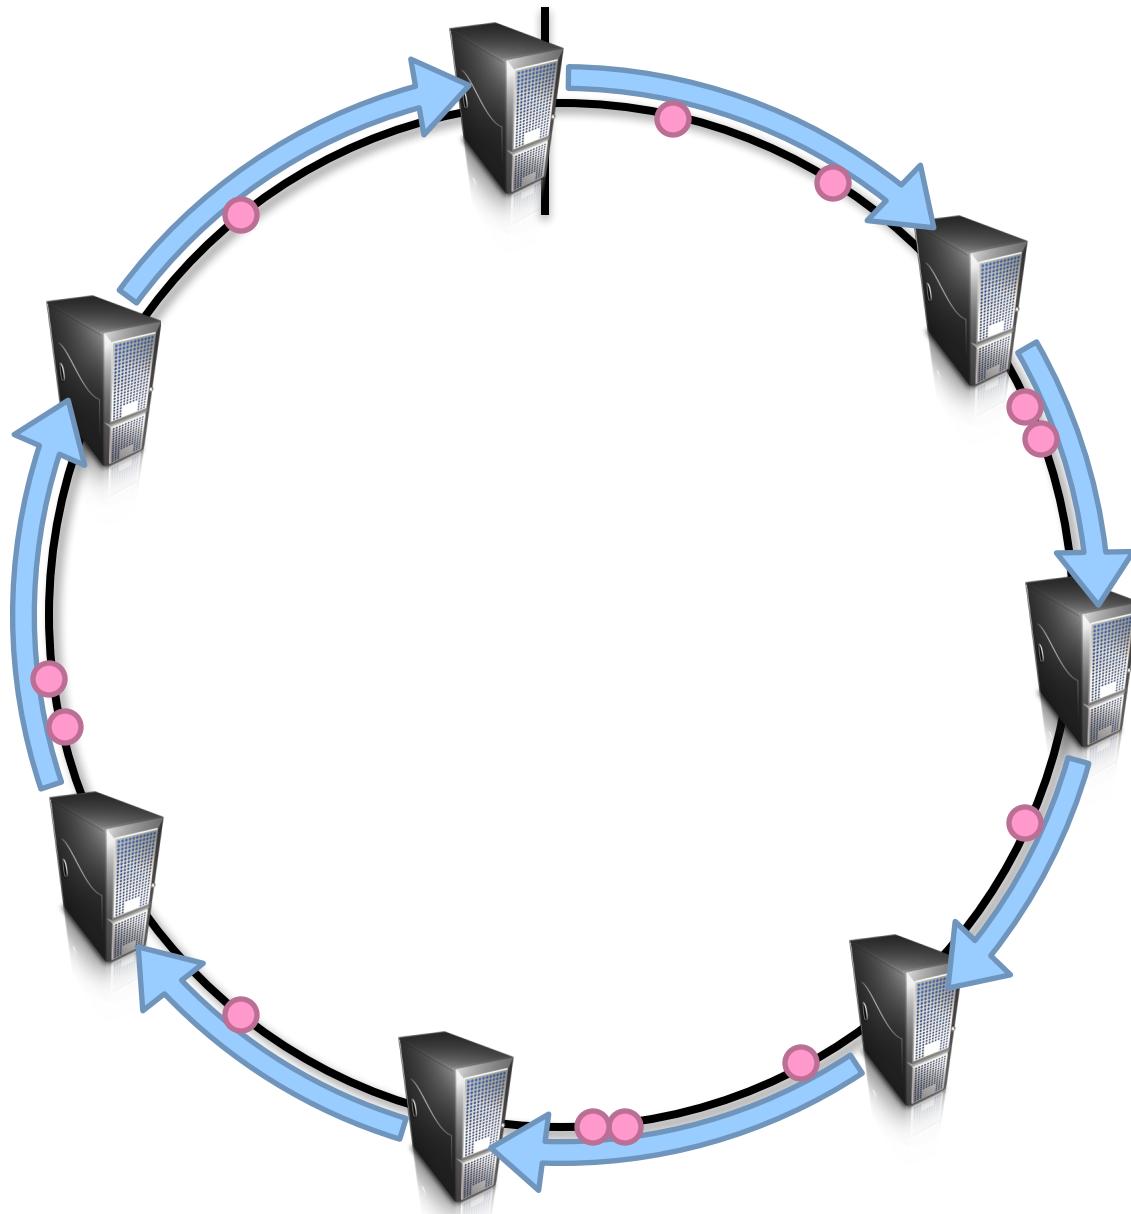
What happens if a machine fails?

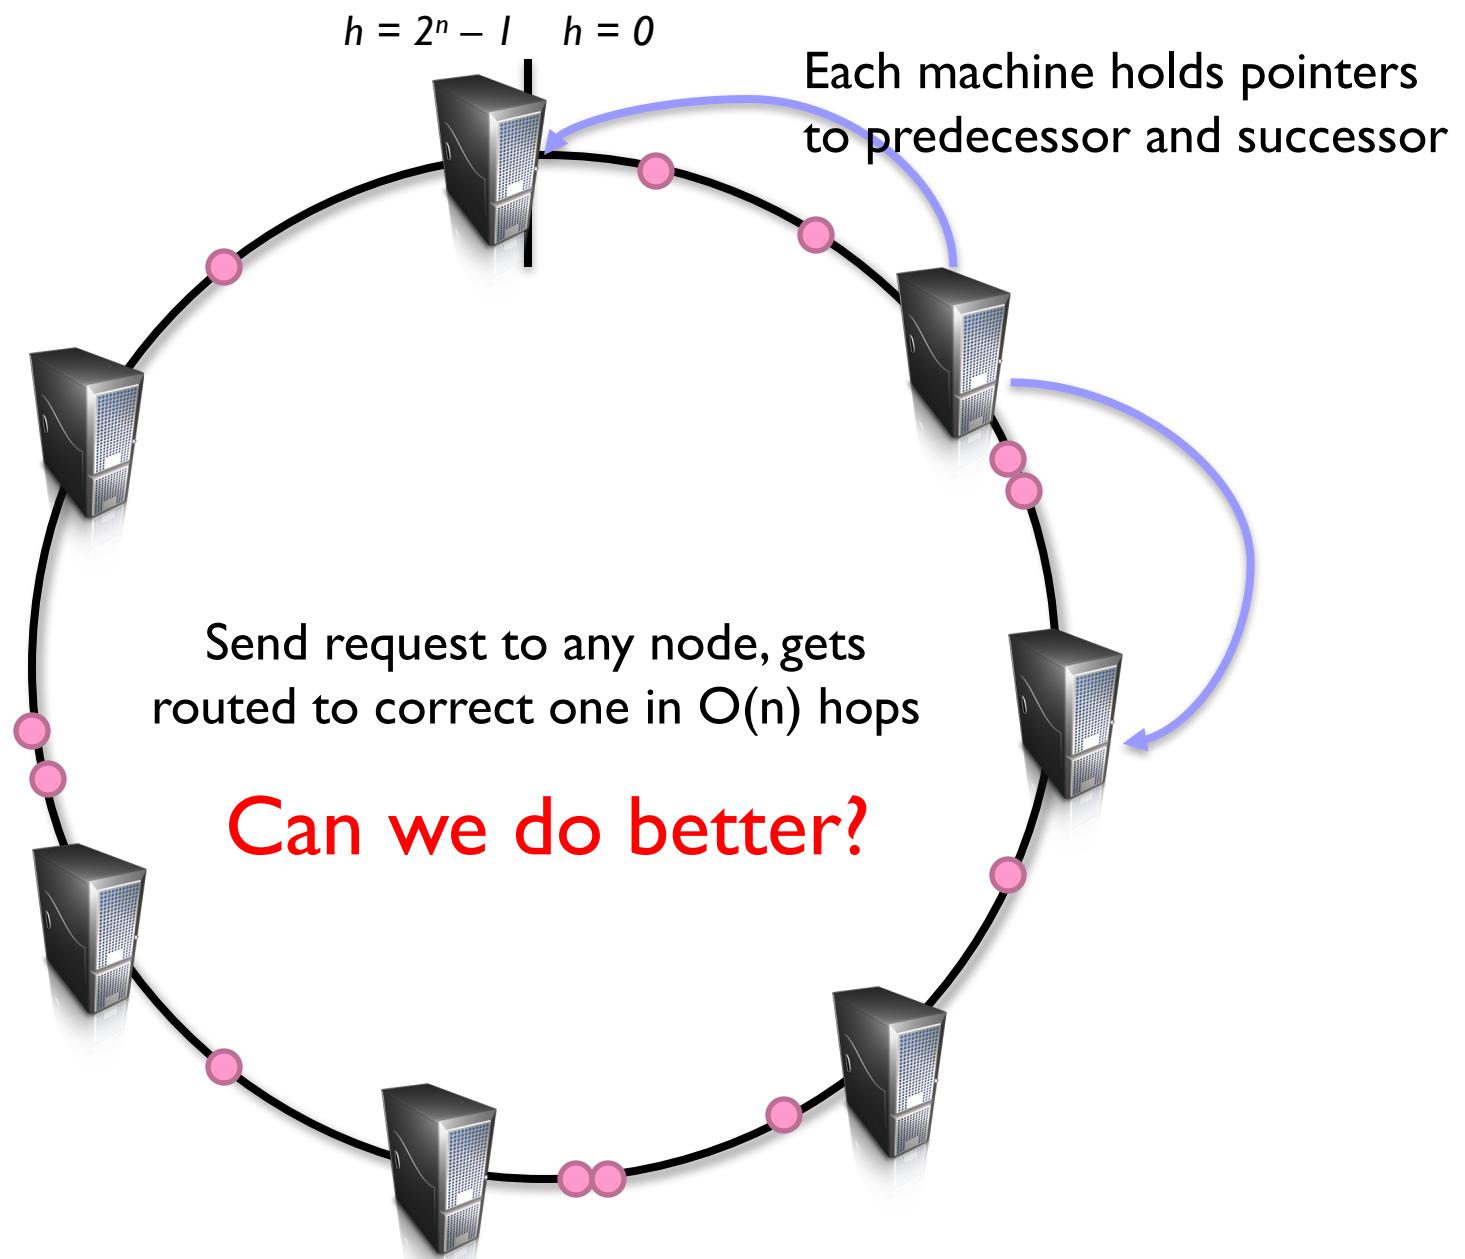
Clever Solution

Hash the keys
Hash the machines also!

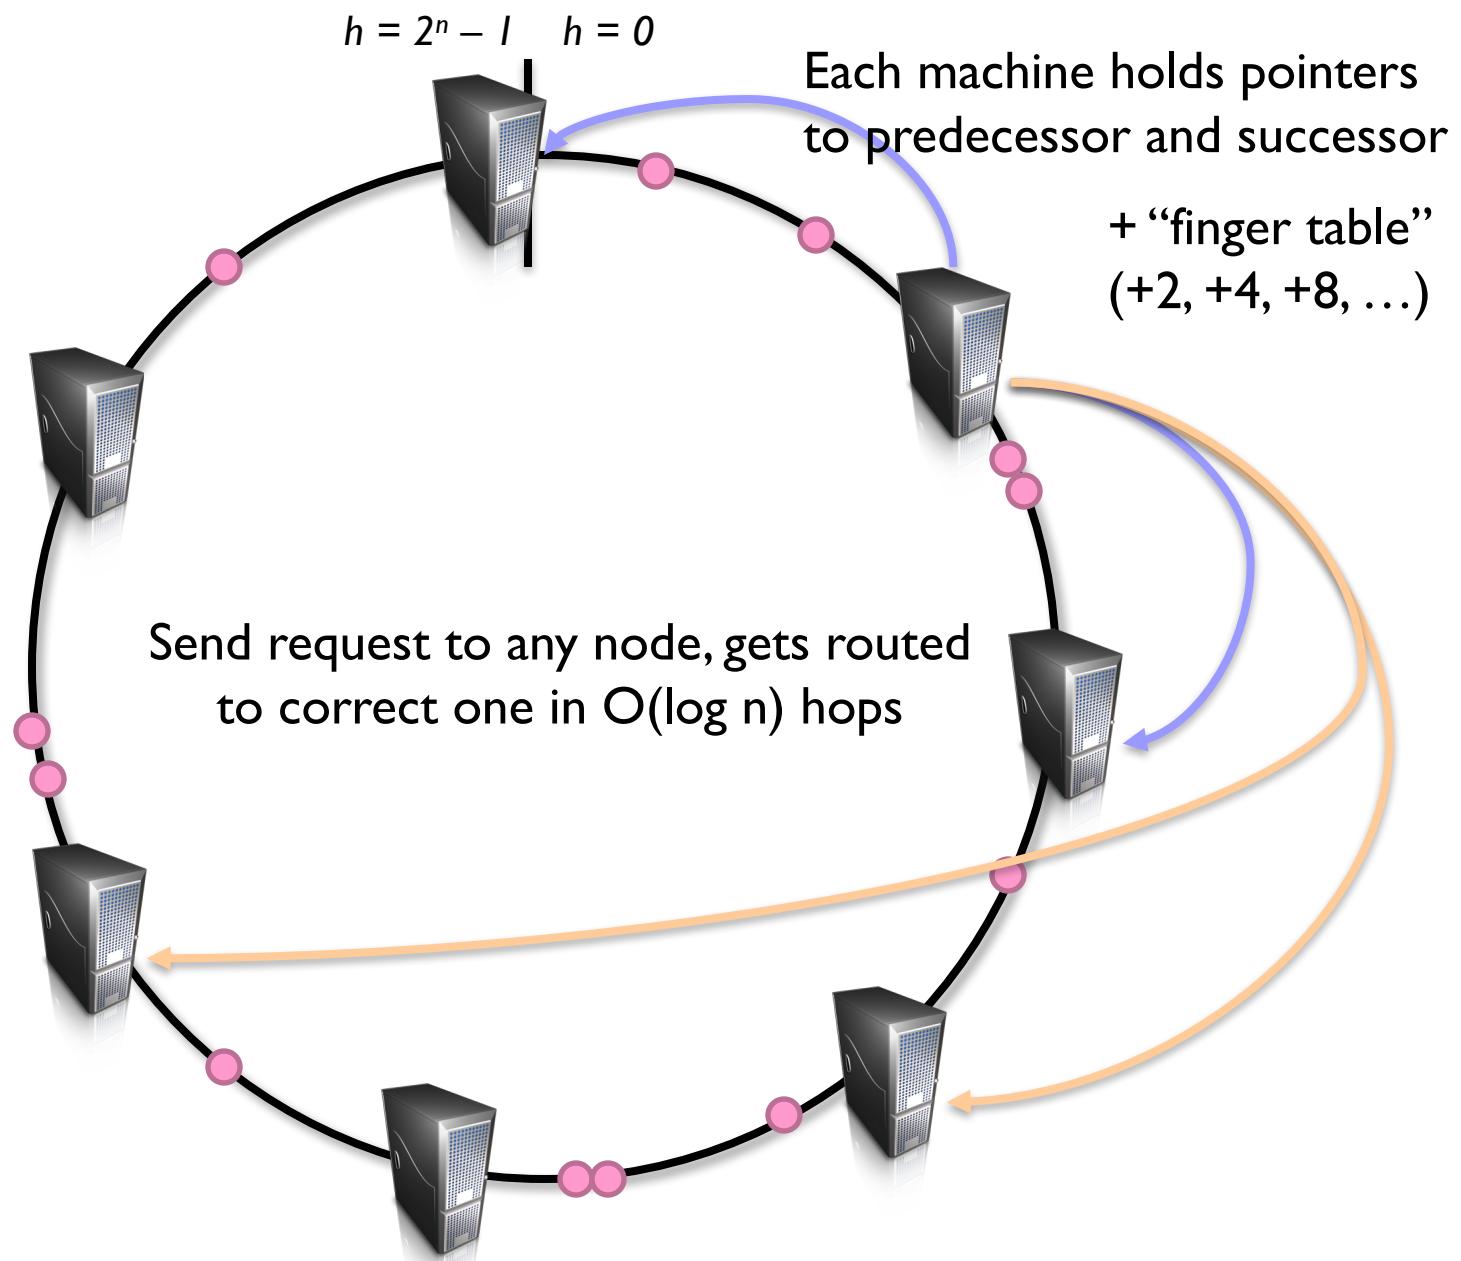
Distributed hash tables!
(following combines ideas from several sources...)

$$h = 2^n - l \quad h = 0$$





Routing: Which machine holds the key?



Routing: Which machine holds the key?

$$h = 2^n - 1 \quad h = 0$$

Simpler Solution

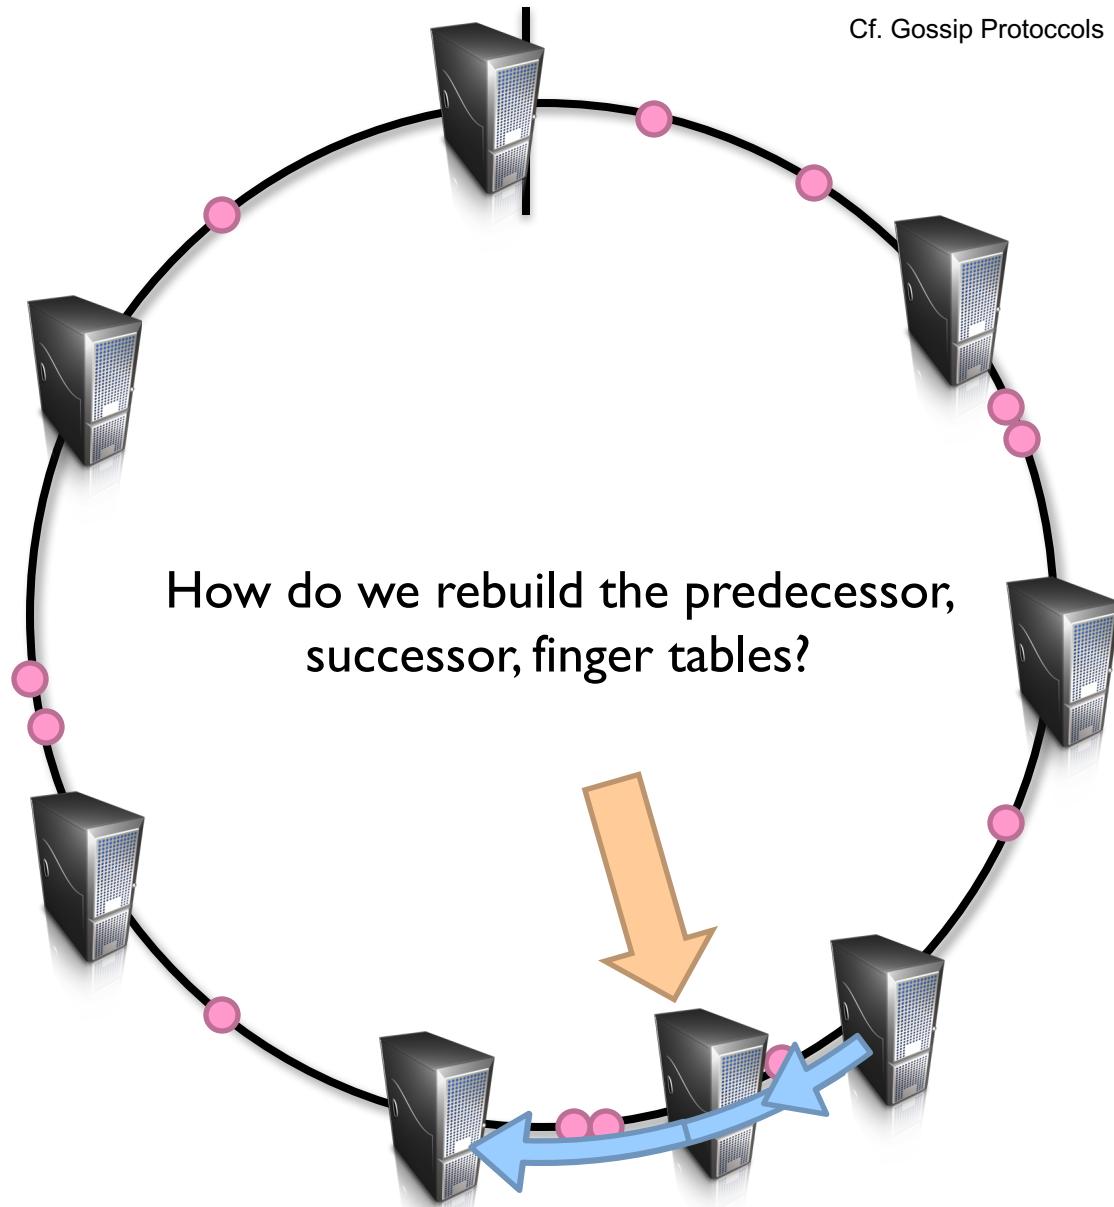
Service
Registry



Routing: Which machine holds the key?

Cf. Gossip Protocols

$$h = 2^n - 1 \quad h = 0$$

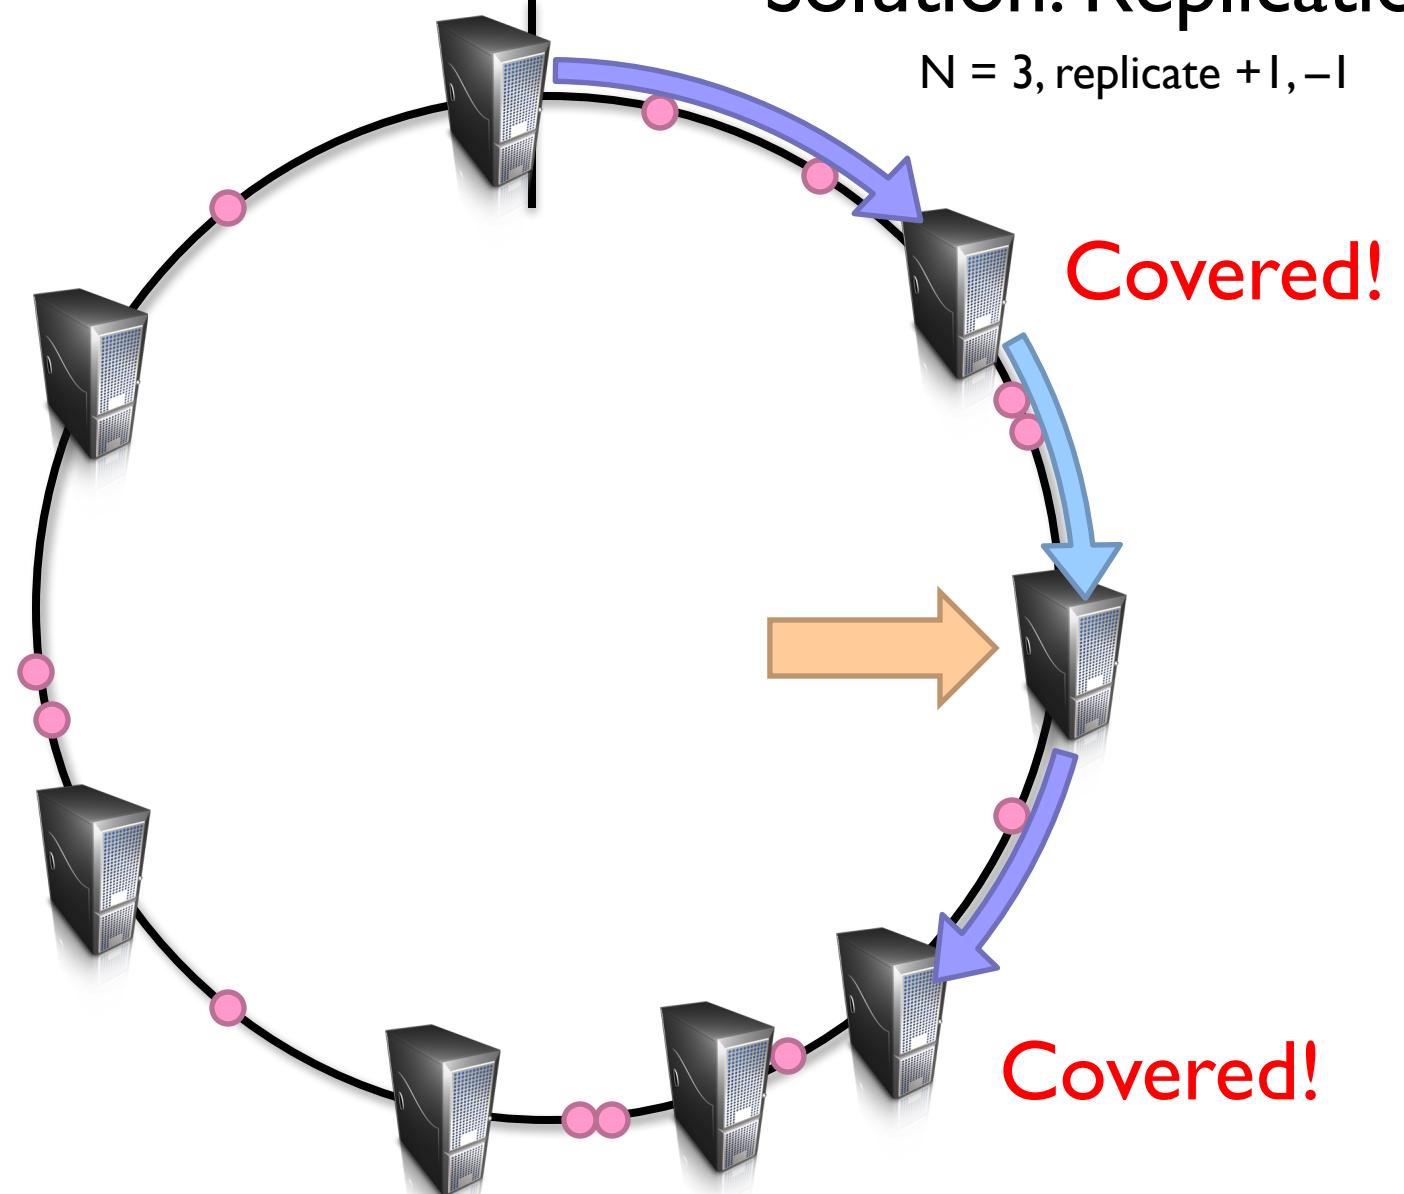


New machine joins: What happens?

$$h = 2^n - 1 \quad h = 0$$

Solution: Replication

N = 3, replicate +1, -1



Machine fails: What happens?

Three Core Ideas

Keeping track of the partitions?

Partitioning (sharding)

To increase scalability and to decrease latency

Consistency?

Replication

To increase robustness (availability) and to increase throughput

Caching

To reduce latency

Another Refinement: Virtual Nodes

Don't directly hash servers

Create a large number of virtual nodes, map to physical servers

Better load redistribution in event of machine failure

When new server joins, evenly shed load from other servers

Bigtable



Bigtable Applications

Gmail

Google's web crawl

Google Earth

Google Analytics

Data source and data sink for MapReduce

HBase is the open-source implementation...

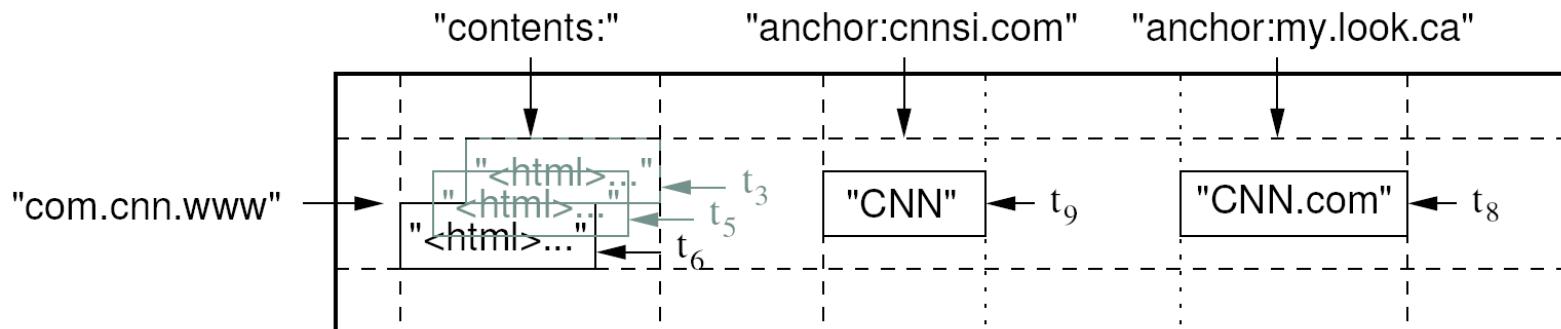
Data Model

A table in Bigtable is a sparse, distributed, persistent multidimensional sorted map

Map indexed by a row key, column key, and a timestamp
 $(\text{row:string}, \text{column:string}, \text{time:int64}) \rightarrow \text{uninterpreted byte array}$

Supports lookups, inserts, deletes

Single row transactions only



Rows and Columns

Rows maintained in sorted lexicographic order

Applications can exploit this property for efficient row scans

Row ranges dynamically partitioned into tablets

Columns grouped into column families

Column key = family:qualifier

Column families provide locality hints

Unbounded number of columns

At the end of the day, it's all key-value pairs!

Key-Values

row, column family, column qualifier, timestamp

value

Okay, so how do we build it?

In Memory

On Disk

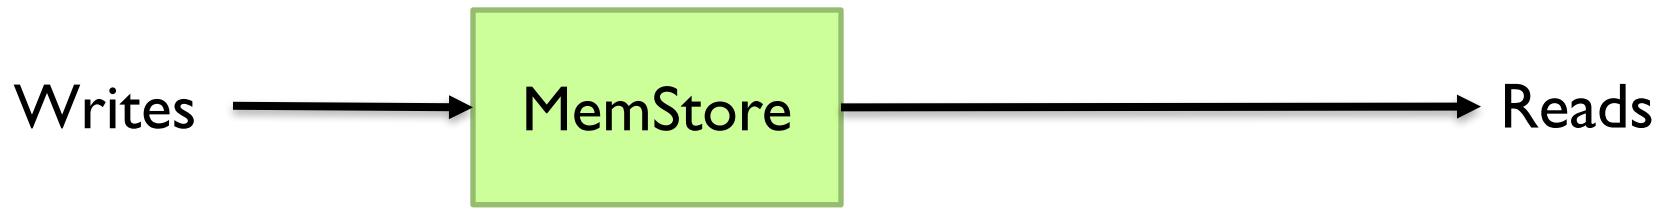
Mutability Easy

Mutability Hard

Small

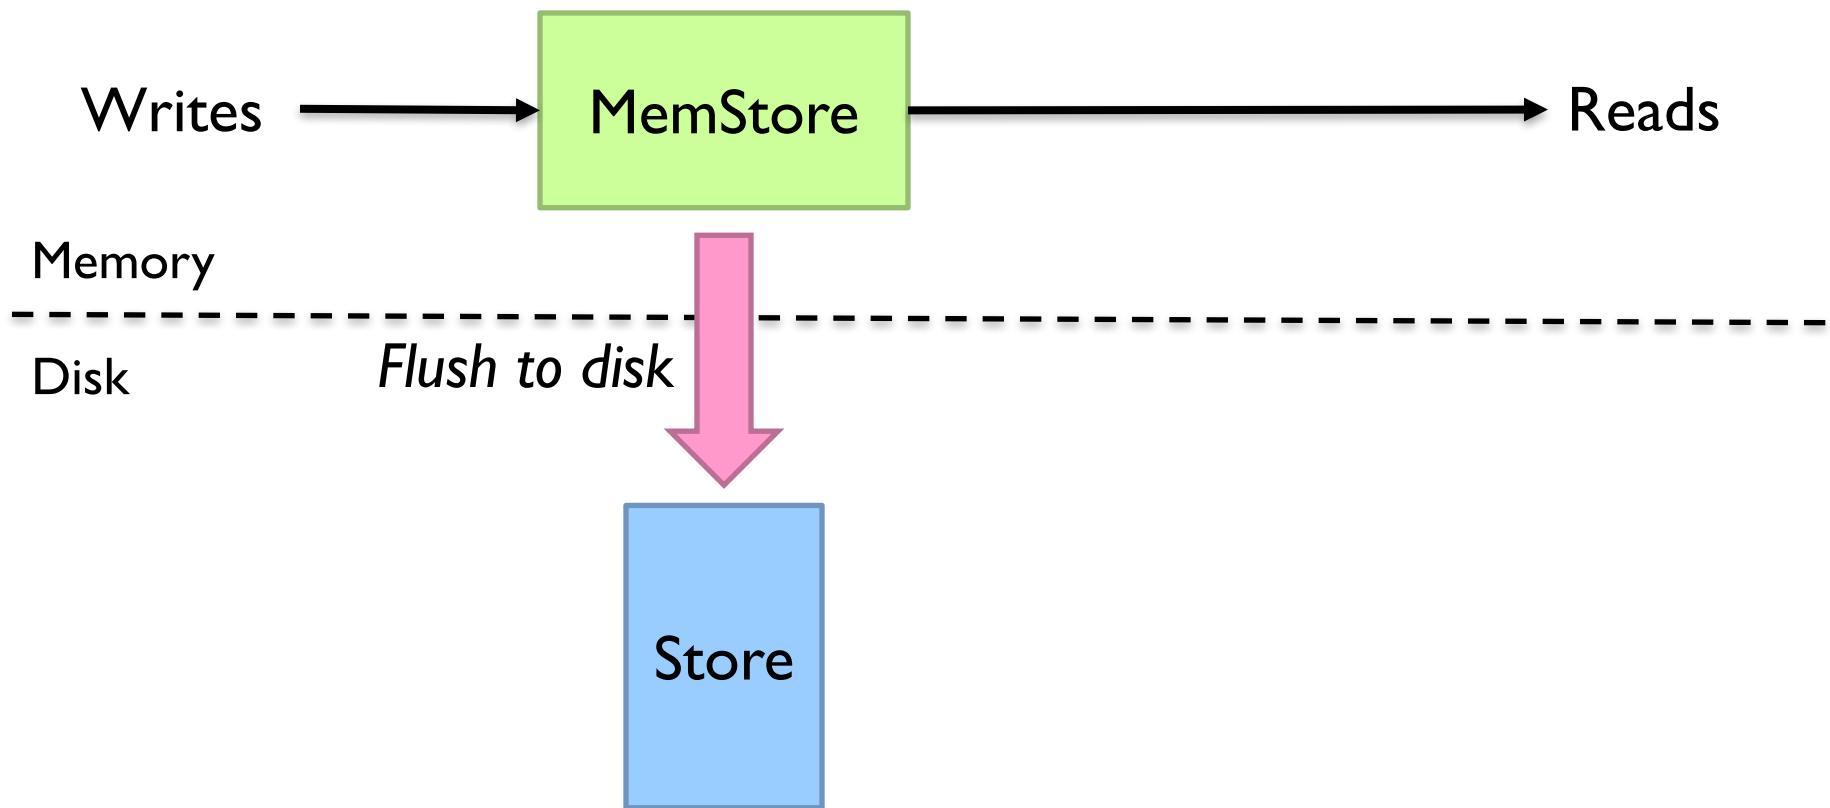
Big

Log Structured Merge Trees



What happens when we run out of memory?

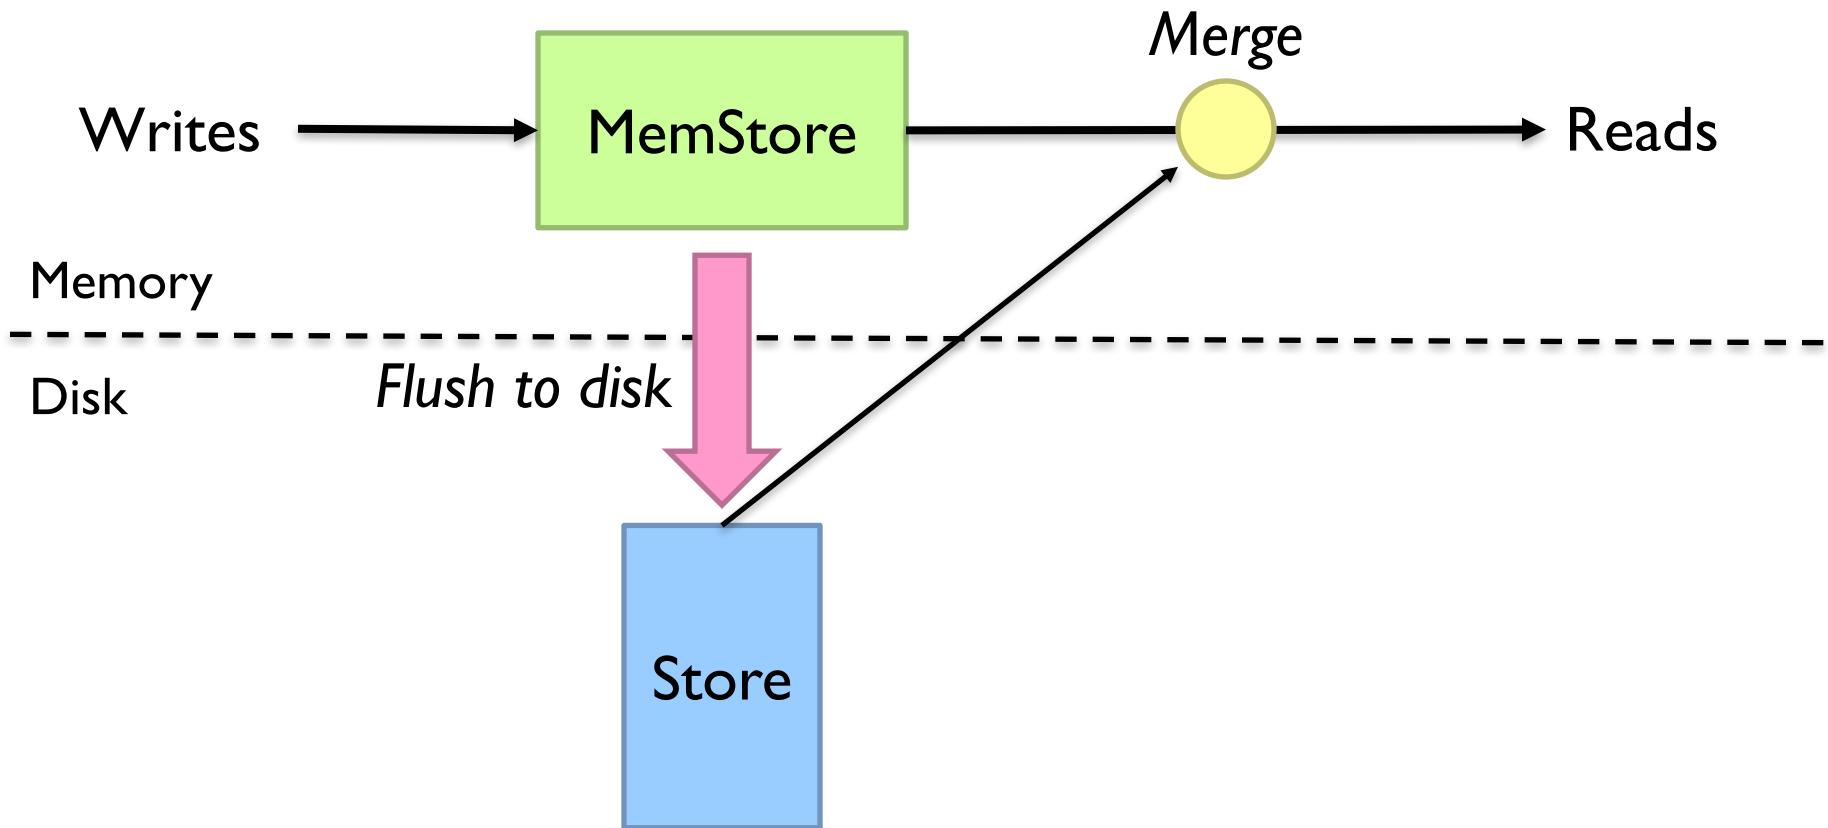
Log Structured Merge Trees



Immutable, indexed, persistent, key-value pairs

What happens to the read path?

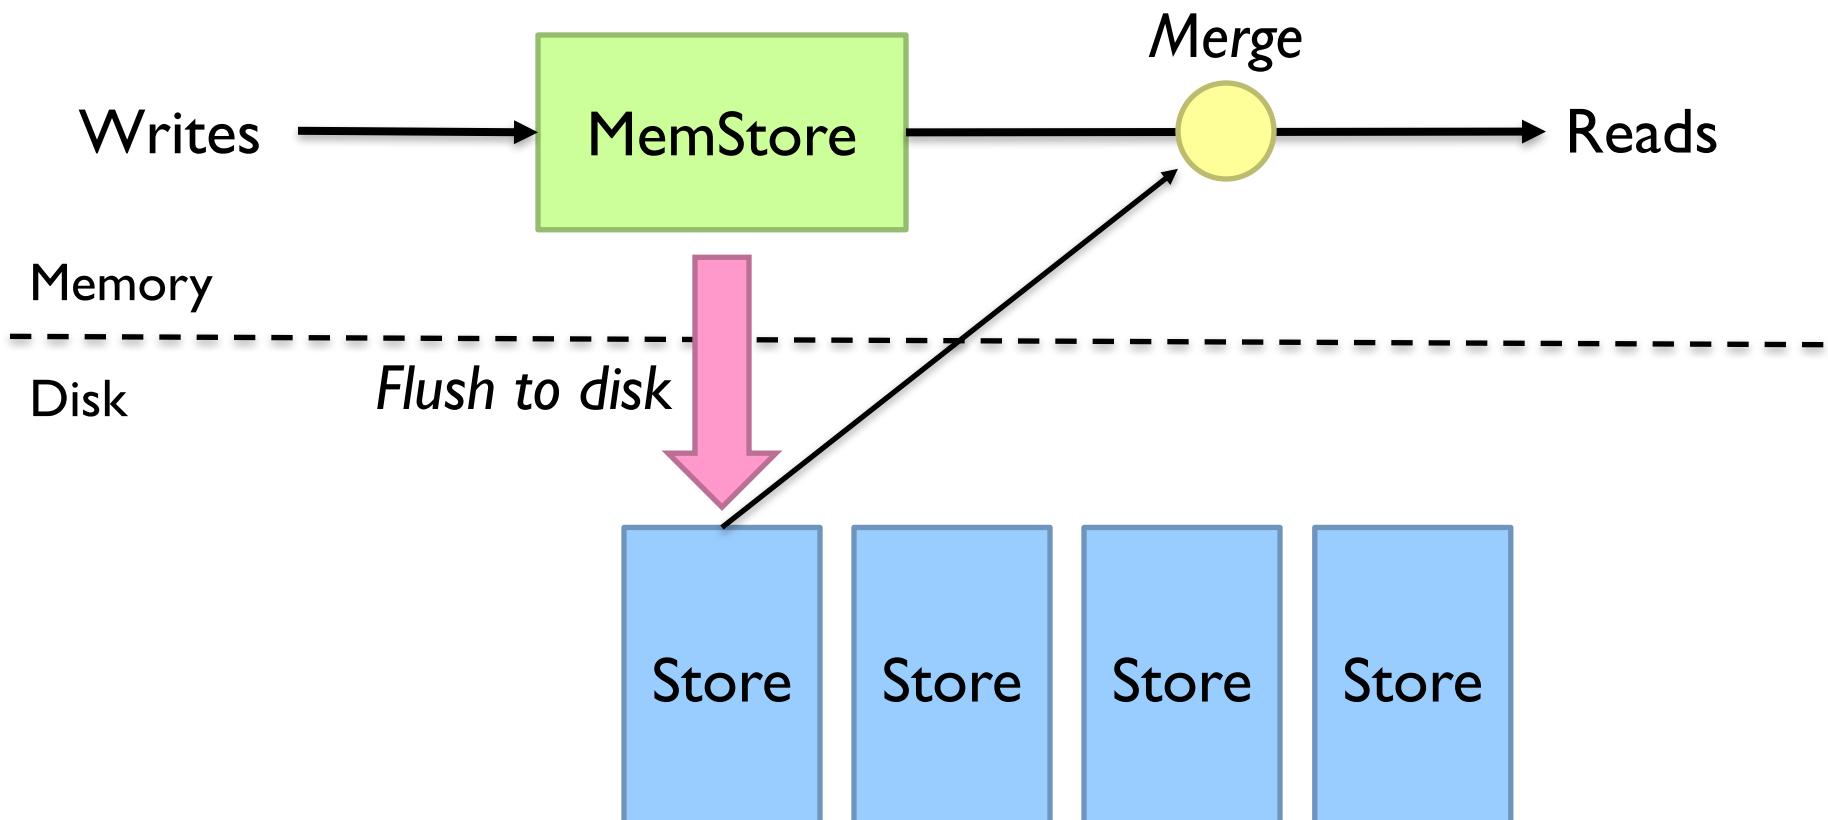
Log Structured Merge Trees



Immutable, indexed, persistent, key-value pairs

What happens as more writes happen?

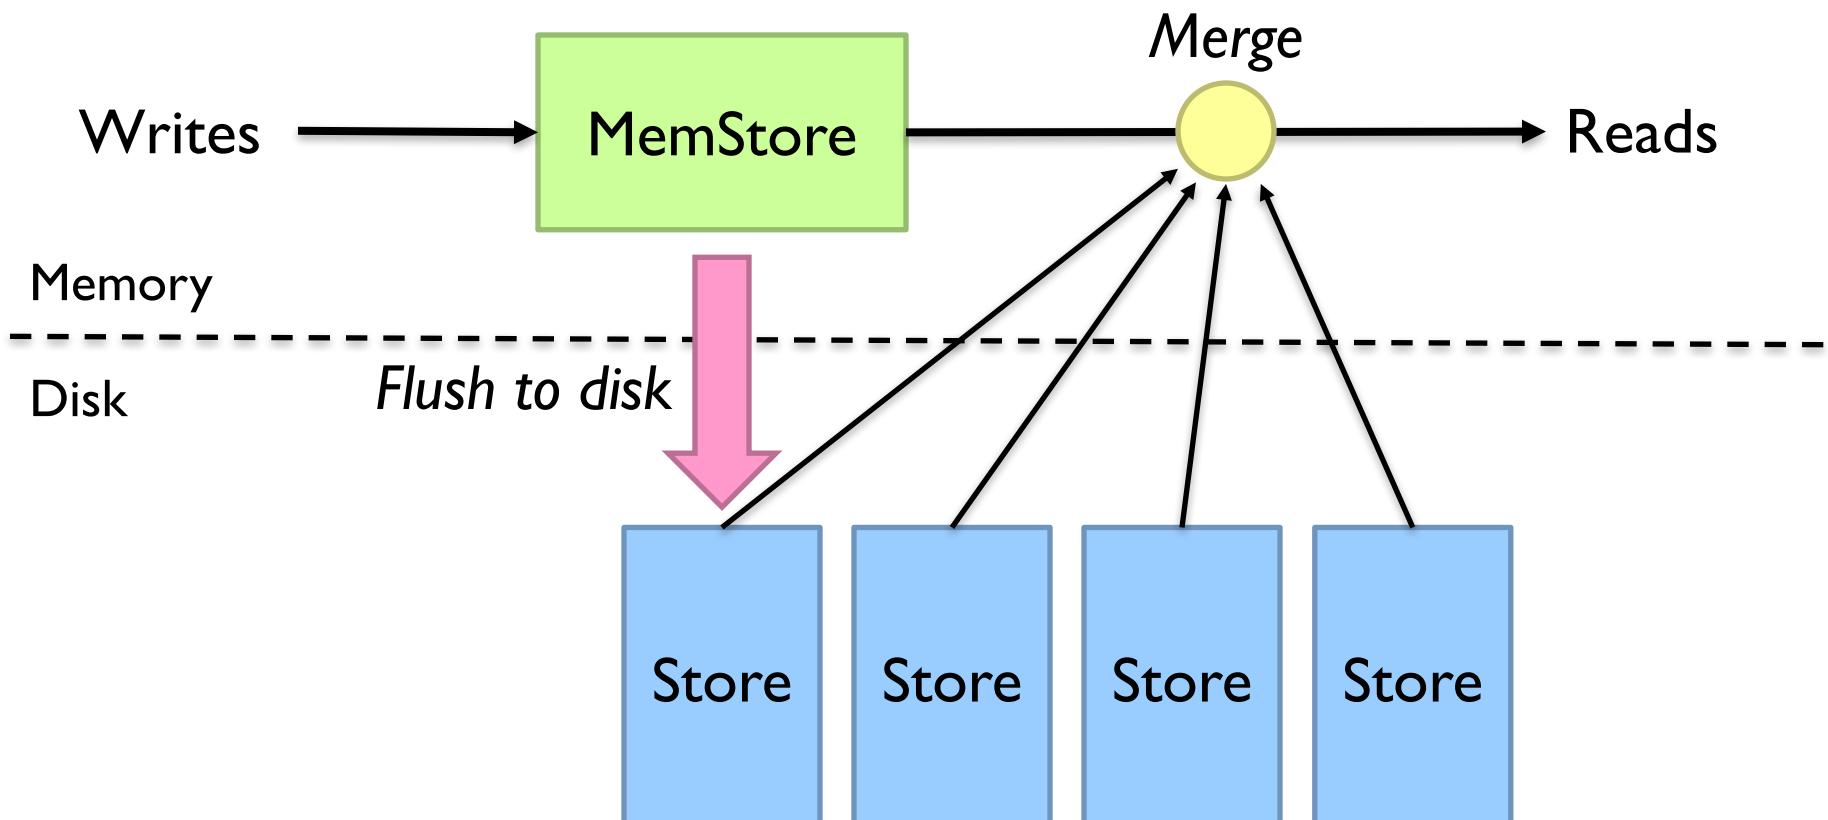
Log Structured Merge Trees



Immutable, indexed, persistent, key-value pairs

What happens to the read path?

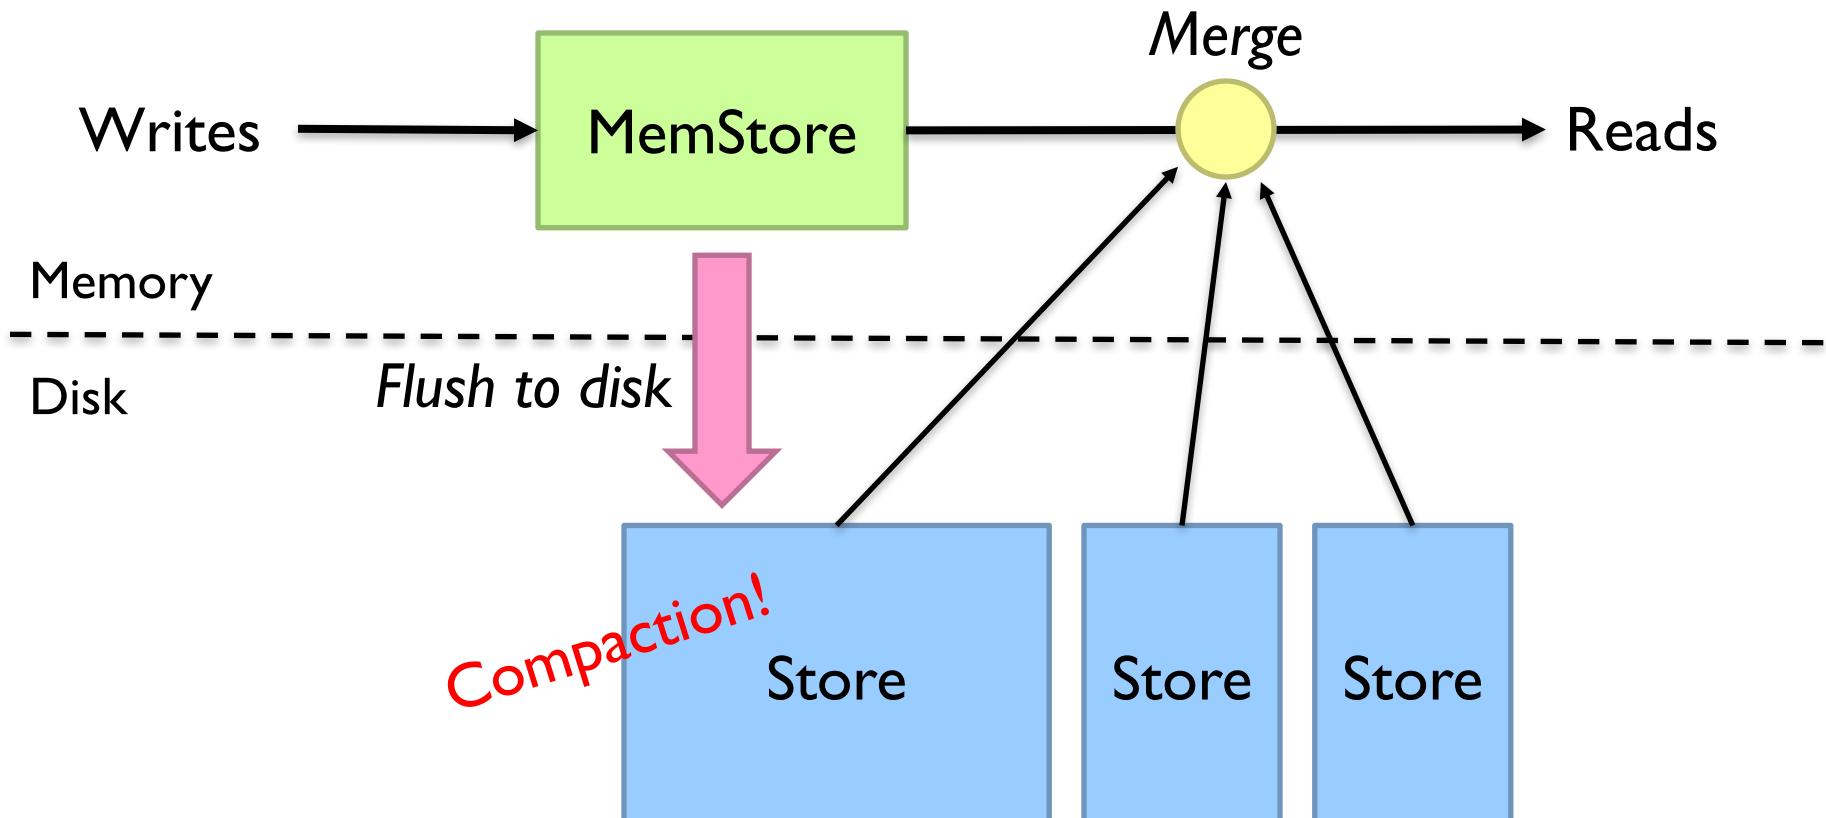
Log Structured Merge Trees



Immutable, indexed, persistent, key-value pairs

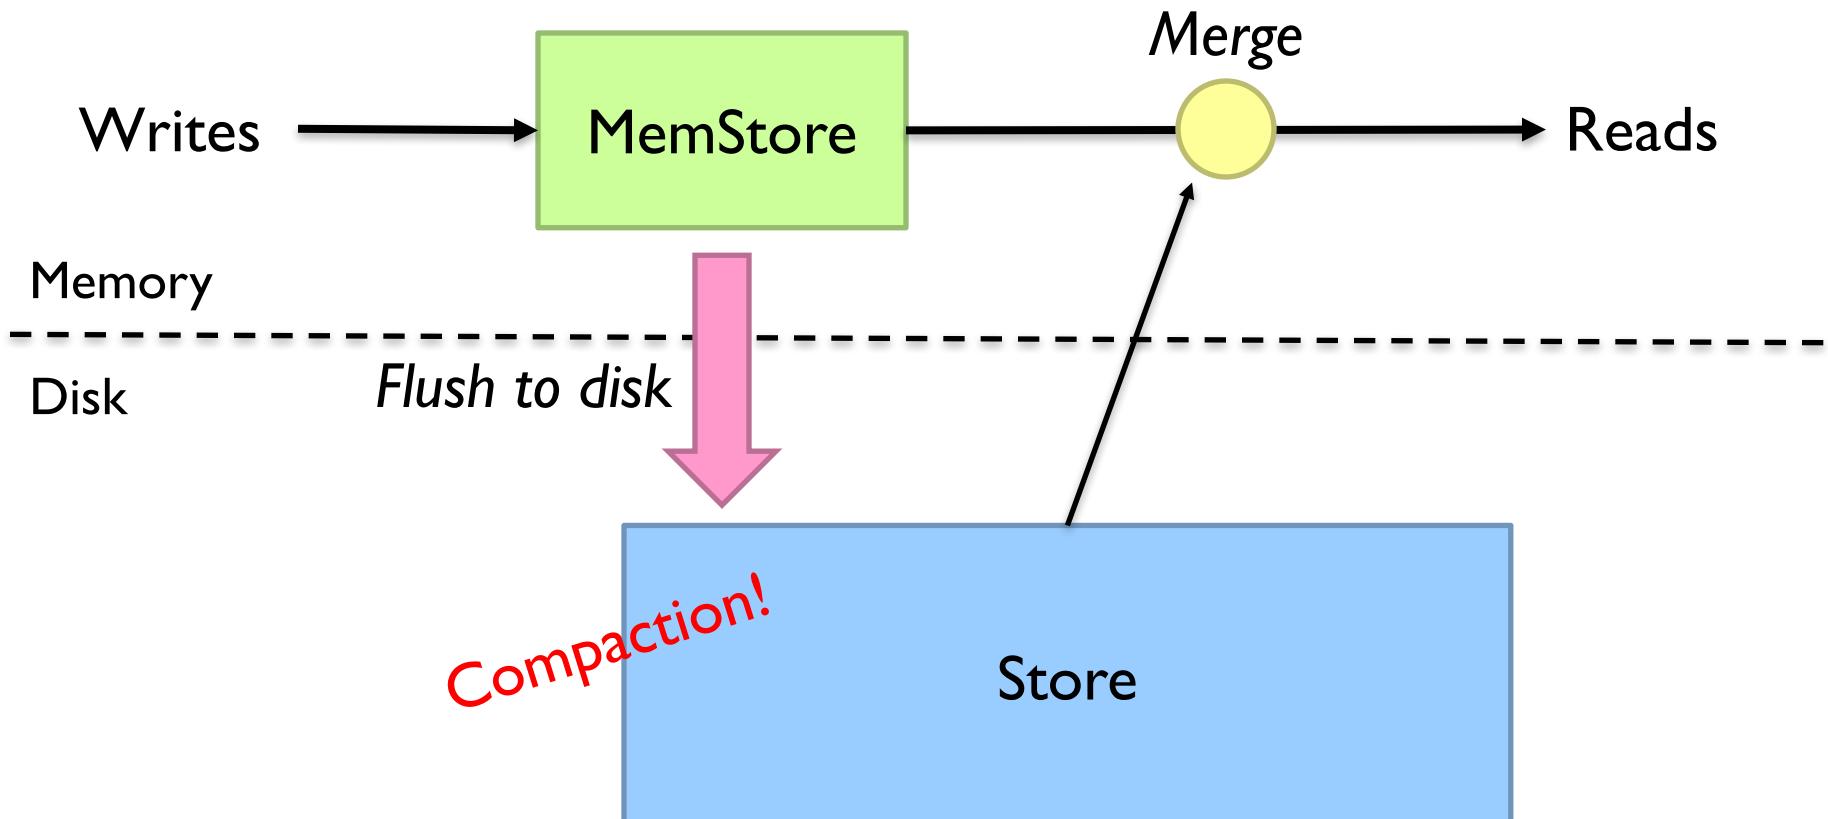
What's the next issue?

Log Structured Merge Trees



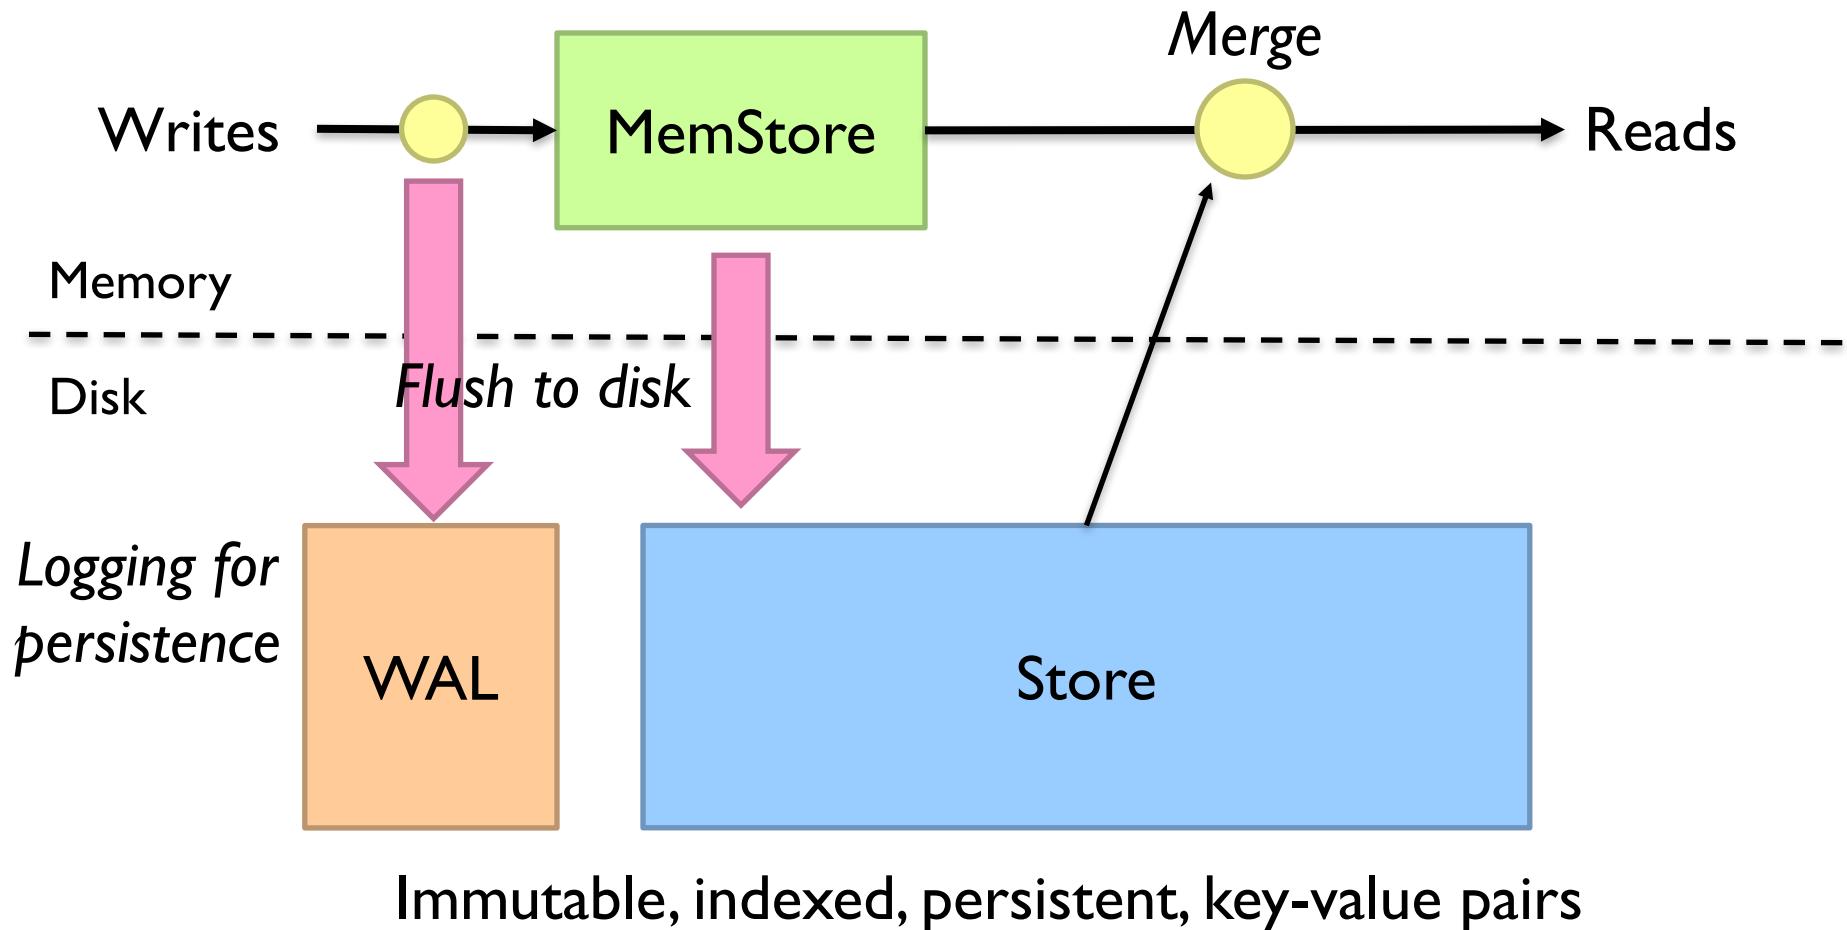
Immutable, indexed, persistent, key-value pairs

Log Structured Merge Trees



Immutable, indexed, persistent, key-value pairs

Log Structured Merge Trees

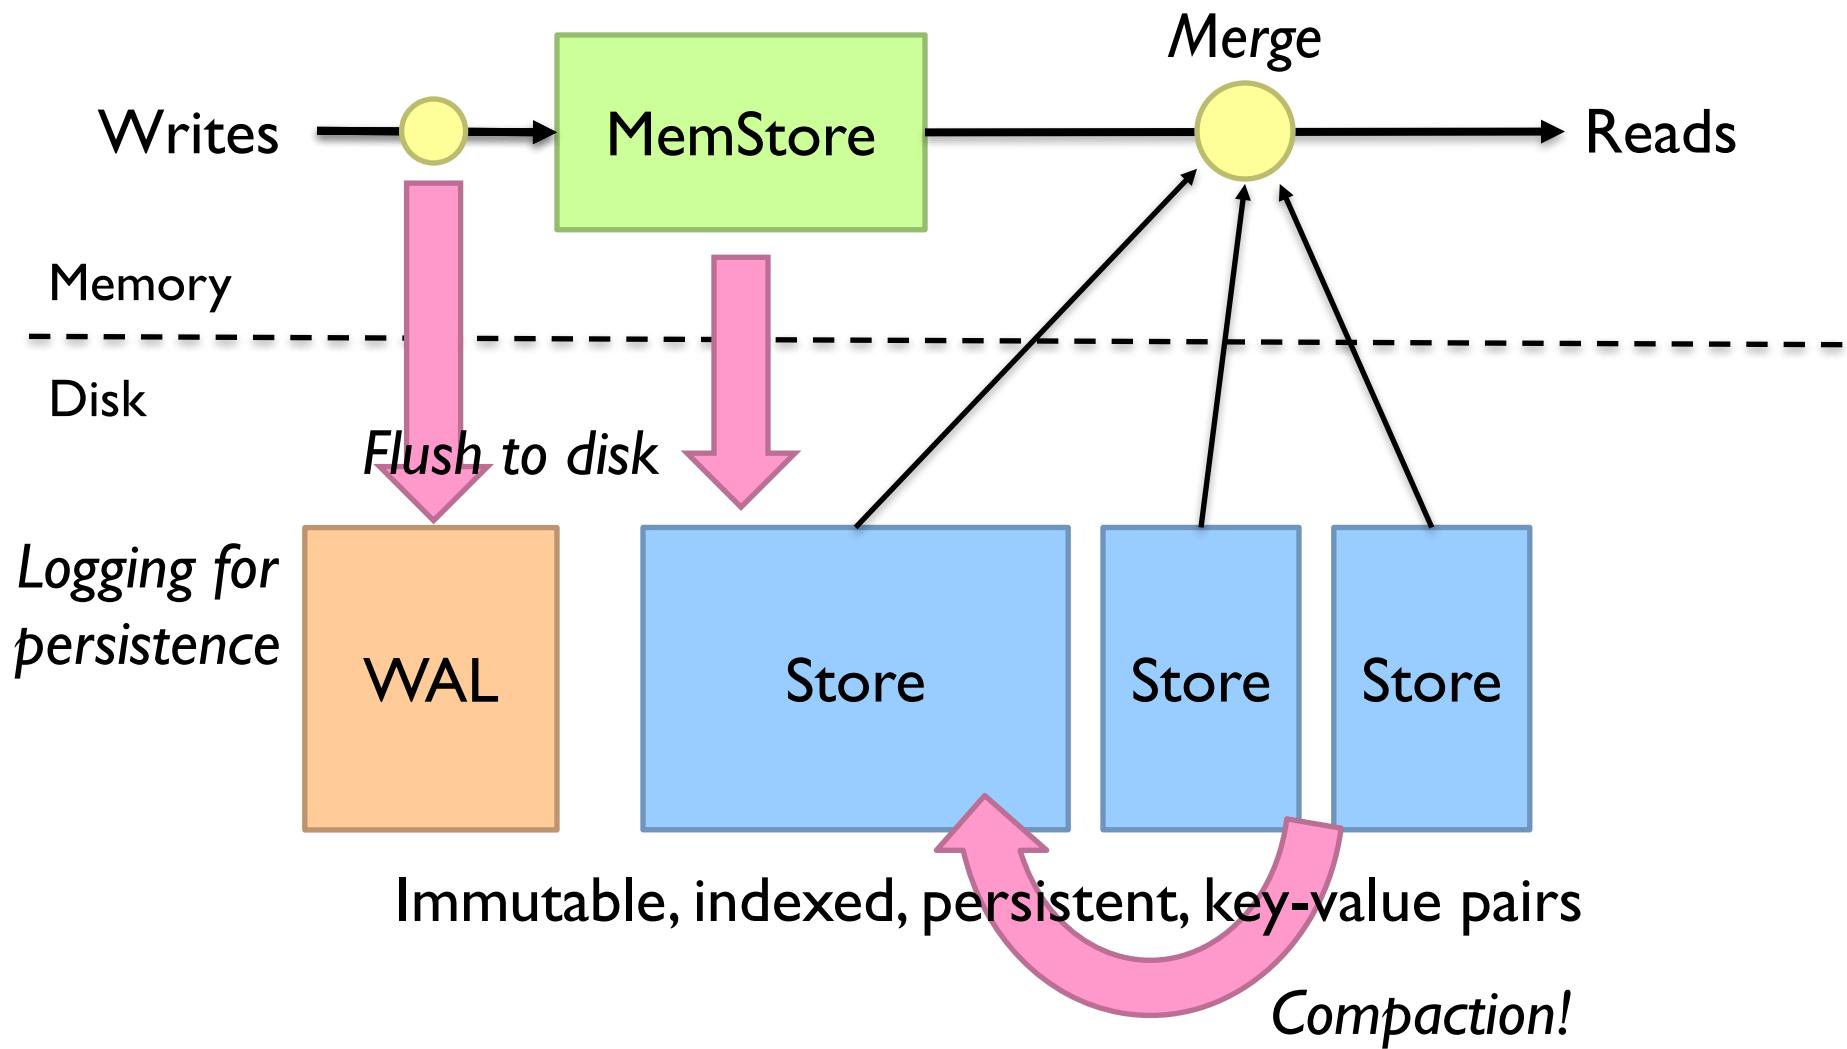


Immutable, indexed, persistent, key-value pairs

One final component...

Log Structured Merge Trees

The complete picture...



Log Structured Merge Trees

The complete picture...

Okay, now how do we build a distributed version?

HBase Bigtable building blocks

HDFS
GFS

HFile SSTable

Tablet Region

Regions Server
Tablet Server

Chubby
Zookeeper

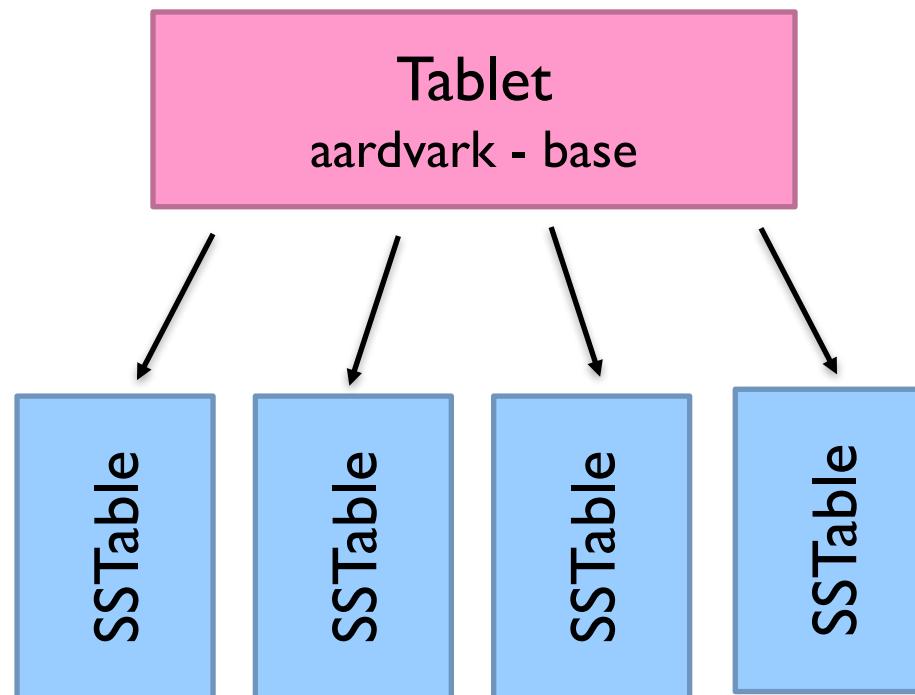
SSTable *HFile*

Persistent, ordered immutable map from keys to values
Stored in GFS: replication “for free”

Supported operations:
Look up value associated with key
Iterate key/value pairs within a key range

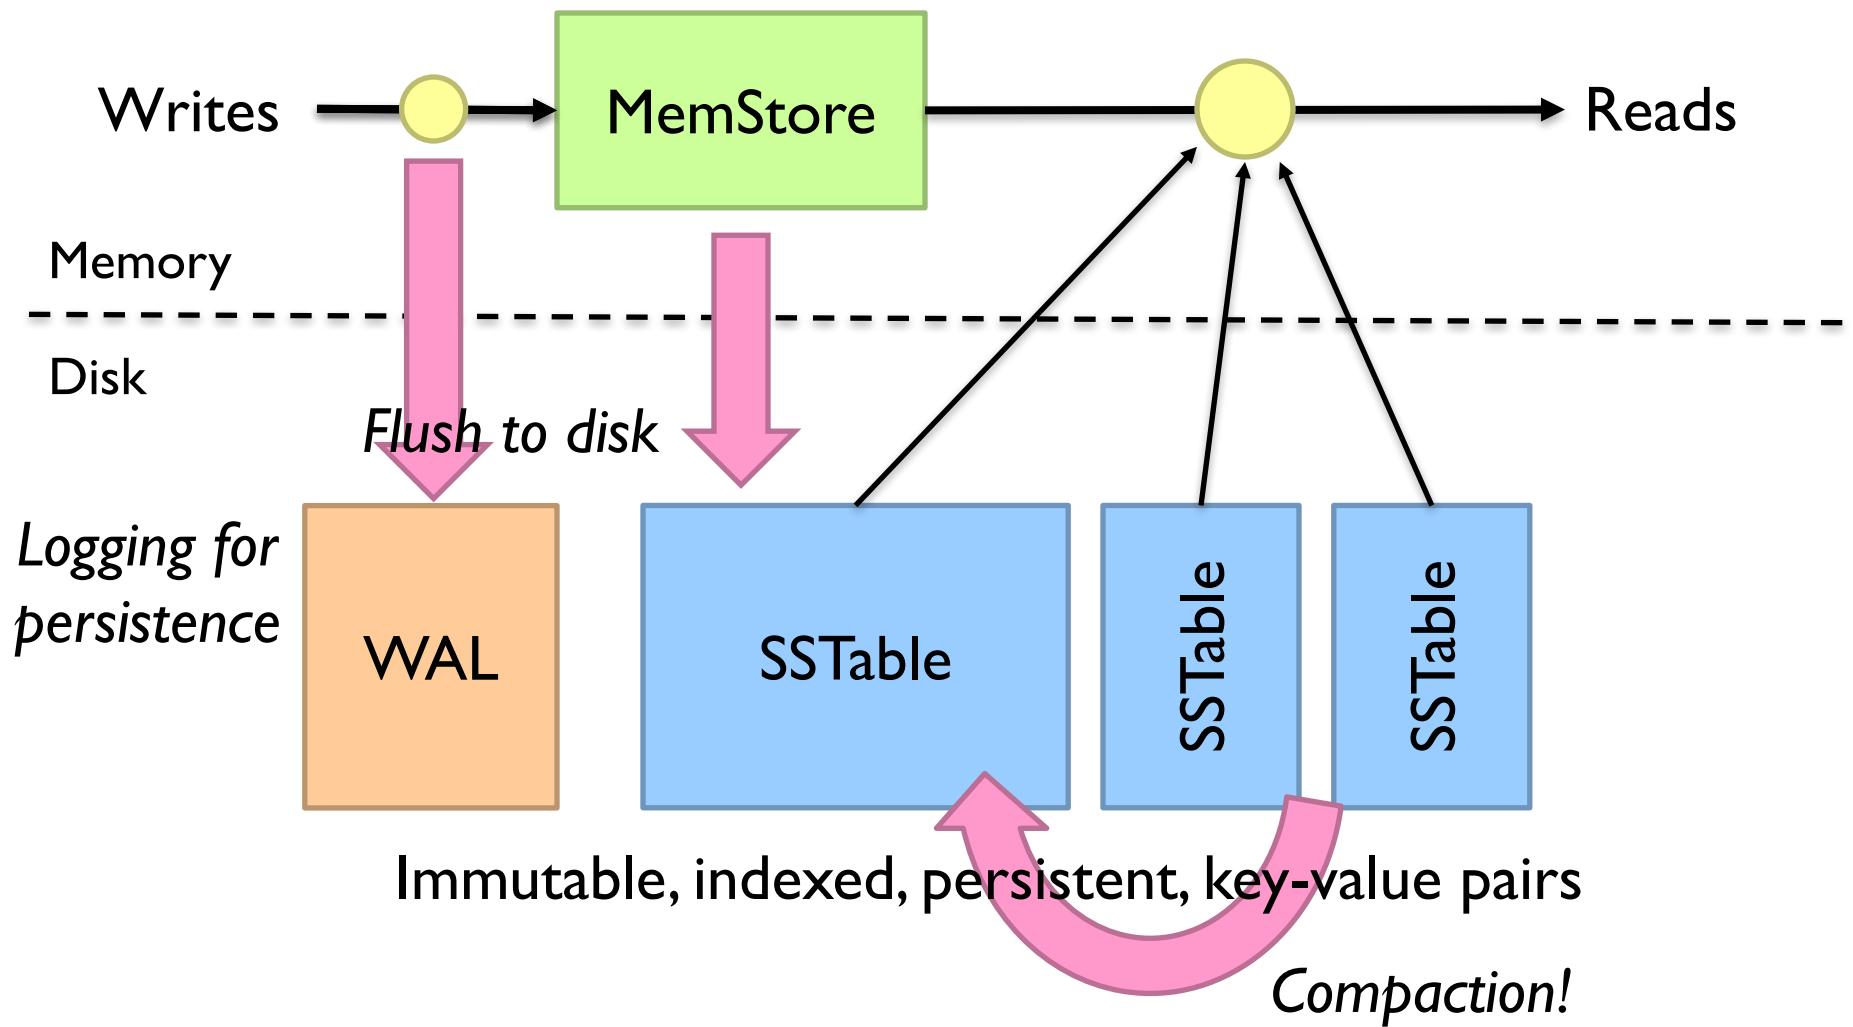
Region Tablet

Dynamically partitioned range of rows
Comprised of multiple SSTables



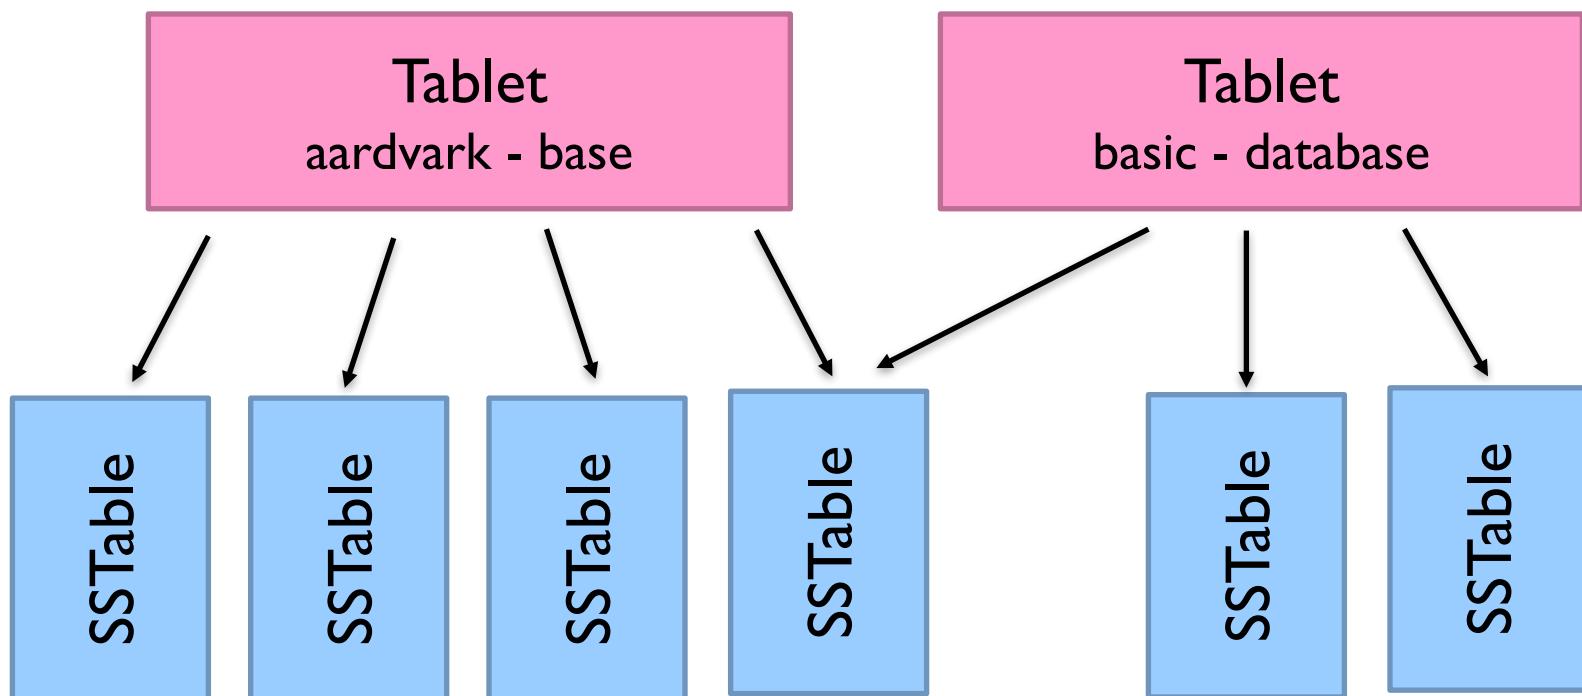
Region Server

Tablet Server



Table

Comprised of multiple tablets
SSTables can be shared between tablets



Region

Region Server

Tablet to Tablet Server Assignment

Each tablet is assigned to one tablet server at a time

Exclusively handles read and write requests to that tablet

What happens when a tablet grows too big?

What happens when a tablet server fails?

We need a lock service!

HBase Bigtable building blocks

HDFS
GFS

HFile SSTable

Tablet Region

Regions Server
Tablet Server

Chubby
Zookeeper

Architecture

Client library

Bigtable master **HMaster**

Tablet servers

Regions Servers

Bigtable Master

Roles and responsibilities:

Assigns tablets to tablet servers

Detects addition and removal of tablet servers

Balances tablet server load

Handles garbage collection

Handles schema changes

Tablet structure changes:

Table creation/deletion (master initiated)

Tablet merging (master initiated)

Tablet splitting (tablet server initiated)

Compactions

Minor compaction

Converts the memtable into an SSTable

Reduces memory usage and log traffic on restart

Merging compaction

Reads a few SSTables and the memtable, and writes out a new SSTable

Reduces number of SSTables

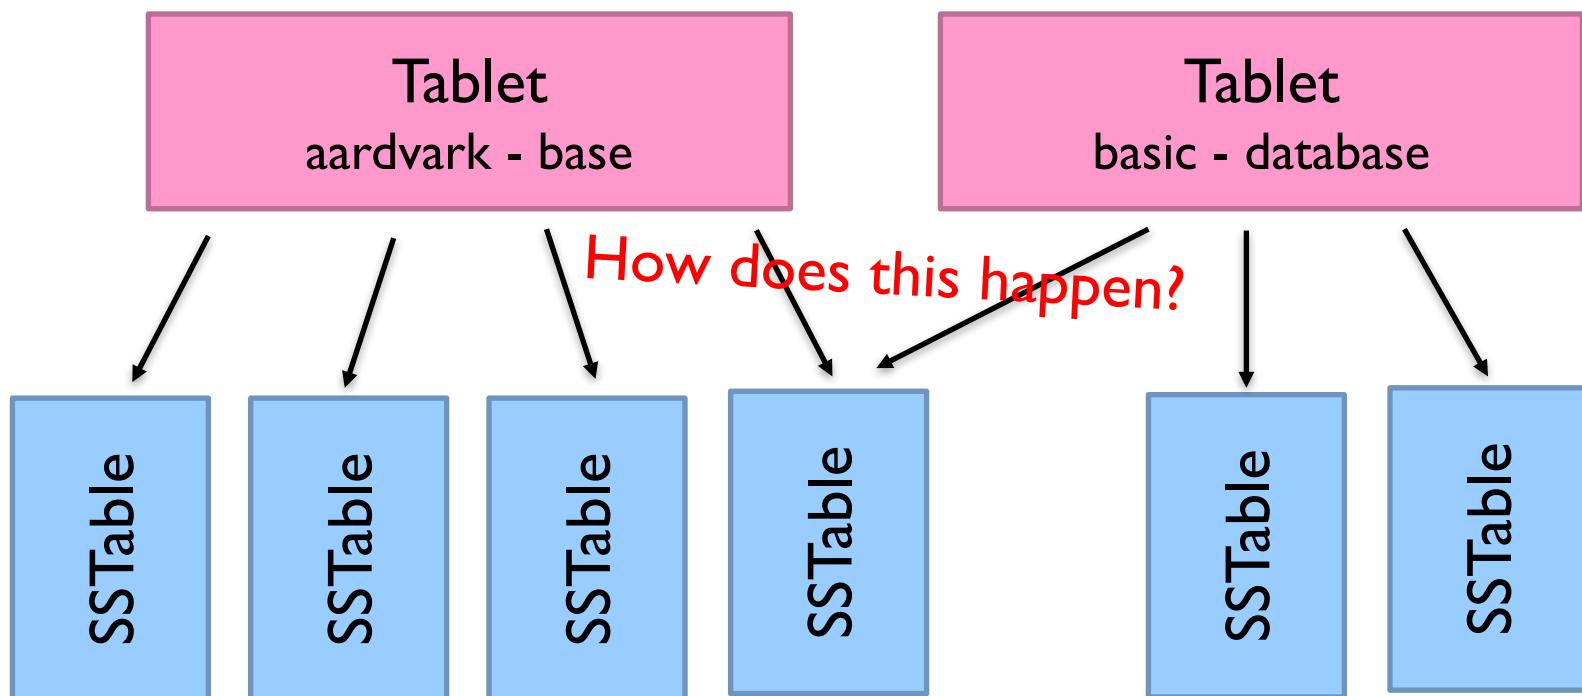
Major compaction

Merging compaction that results in only one SSTable

No deletion records, only live data

Table

Comprised of multiple tables
SSTables can be shared between tablets



Three Core Ideas

Keeping track of the partitions?

Partitioning (sharding)

To increase scalability and to decrease latency

Consistency?

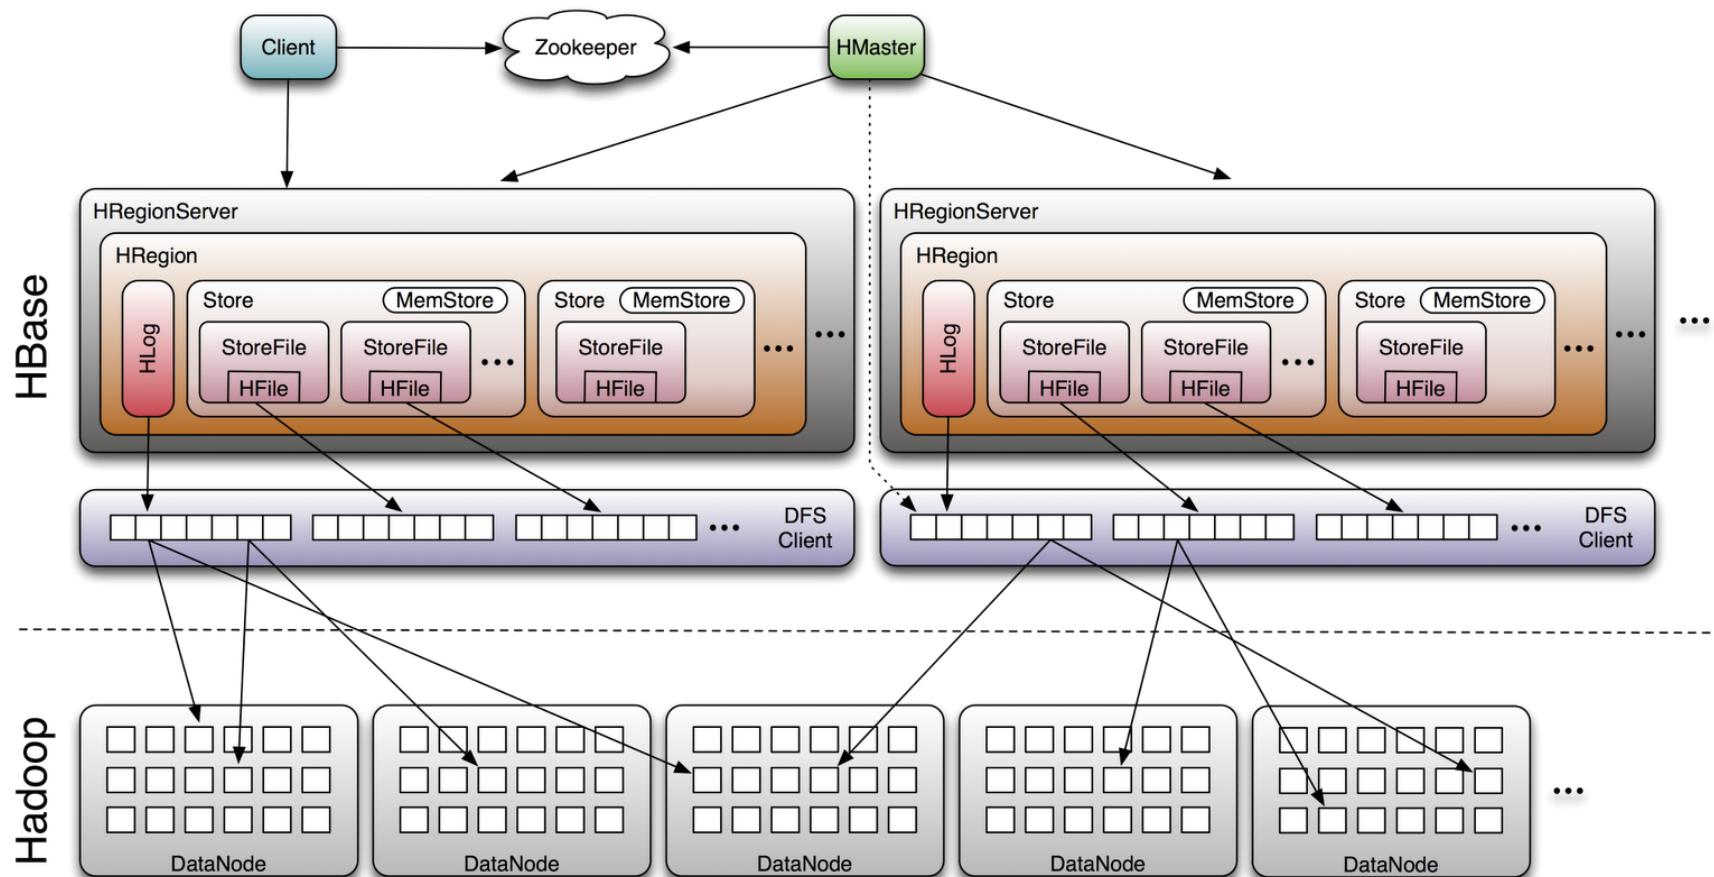
Replication

To increase robustness (availability) and to increase throughput

Caching

To reduce latency

HBase



A photograph of a traditional Japanese rock garden. In the foreground, a gravel path is raked into fine, parallel lines. Several large, dark, irregular stones are scattered across the garden. A small, shallow pond is visible in the middle ground, surrounded by more stones and some low-lying green plants. In the background, there are more stones, some small trees, and the wooden buildings of a residence with tiled roofs.

Questions?