



Big Data Infrastructure

CS 489/698 Big Data Infrastructure (Winter 2017)

Week 12: Real-Time Data Analytics (1/2)
March 28, 2017

Jimmy Lin
David R. Cheriton School of Computer Science
University of Waterloo

These slides are available at <http://lintool.github.io/bigdata-2017w/>



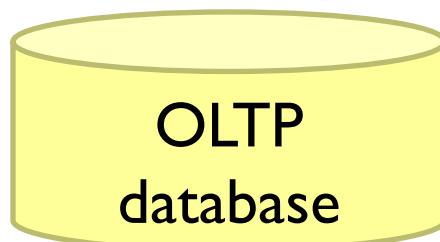
This work is licensed under a Creative Commons Attribution-Noncommercial-Share Alike 3.0 United States
See <http://creativecommons.org/licenses/by-nc-sa/3.0/us/> for details



users

Frontend

Backend



OLTP
database

ETL

(Extract, Transform, and Load)



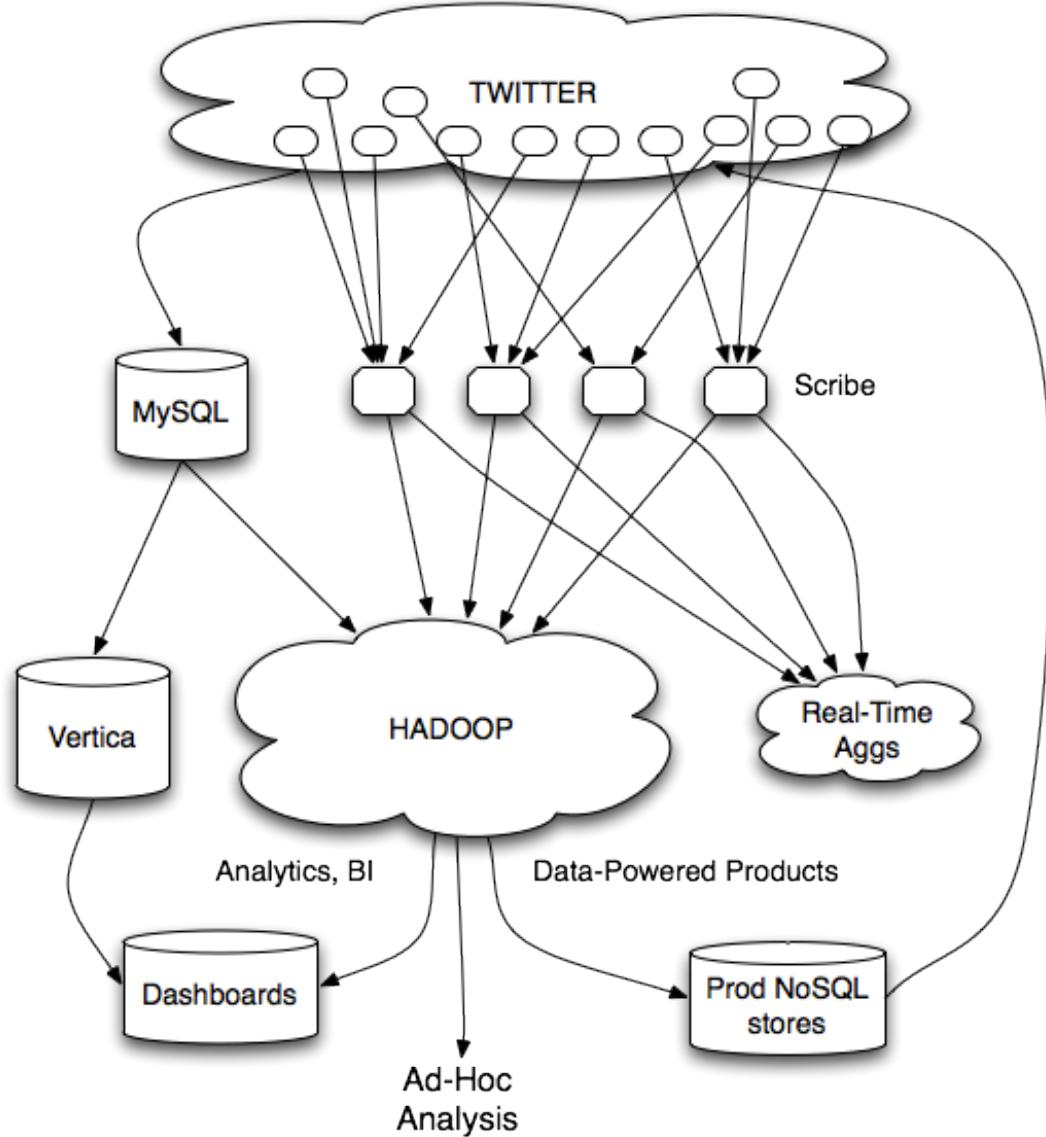
Data
Warehouse

BI tools

analysts

My data is a
day old...

Meh.



Twitter's data warehousing architecture



Tweets

Mishne et al. Fast Data in the Era of Big Data: Twitter's Real-Time Related Query Suggestion Architecture. SIGMOD 2013.

@lintool

TWEETS

1,647

FOLLOWING

253

FOLLOWERS

6,565

Compose new Tweet...

Who to follow · Refresh · View all



plotly @plotlygraphs

[Follow](#)

Promoted



Brad Anderson @boorad

Followed by Florian Leibert ...

[Follow](#)

Sheila Morrissey @sheilaMorr

[Follow](#)

Popular accounts · Find friends

Trends · Change

#Olympics Promoted

Ukraine

#ConfessYourUnpopularOpinion

Venny

#PremioLoNuestro

cloudera

Struggling with complex data of Data Science 2/20 to rehi

Promoted by Cloudera

Expand



Retweeted by Nitin Madnani

**Clinton Paquin** @clintonpaquin

Simply stated, "The only problem is muscle memory" @TheChange

[View conversation](#)**The Hill** @thehill · 1h

Republicans take debt ceiling

[View summary](#)

Retweeted by Alex Feinberg

Popehat @Popehat · 10h

In a world in which few things feed does.

Expand

**The Hill** @thehill · 1hBoehner: I'd rather kill myself than raise the minimum wage trib.al/jZikEus by @mollyhooper and @BobCusack[View summary](#)**CNN Breaking News** @cnnbrk · 1h

Ukrainian Pres. says he has begun work on 3 key opposition demands: New elections, return to old constitution, formation of a unity gov's.

Expand

#Sochi2014

#SochiProblems

Sochi

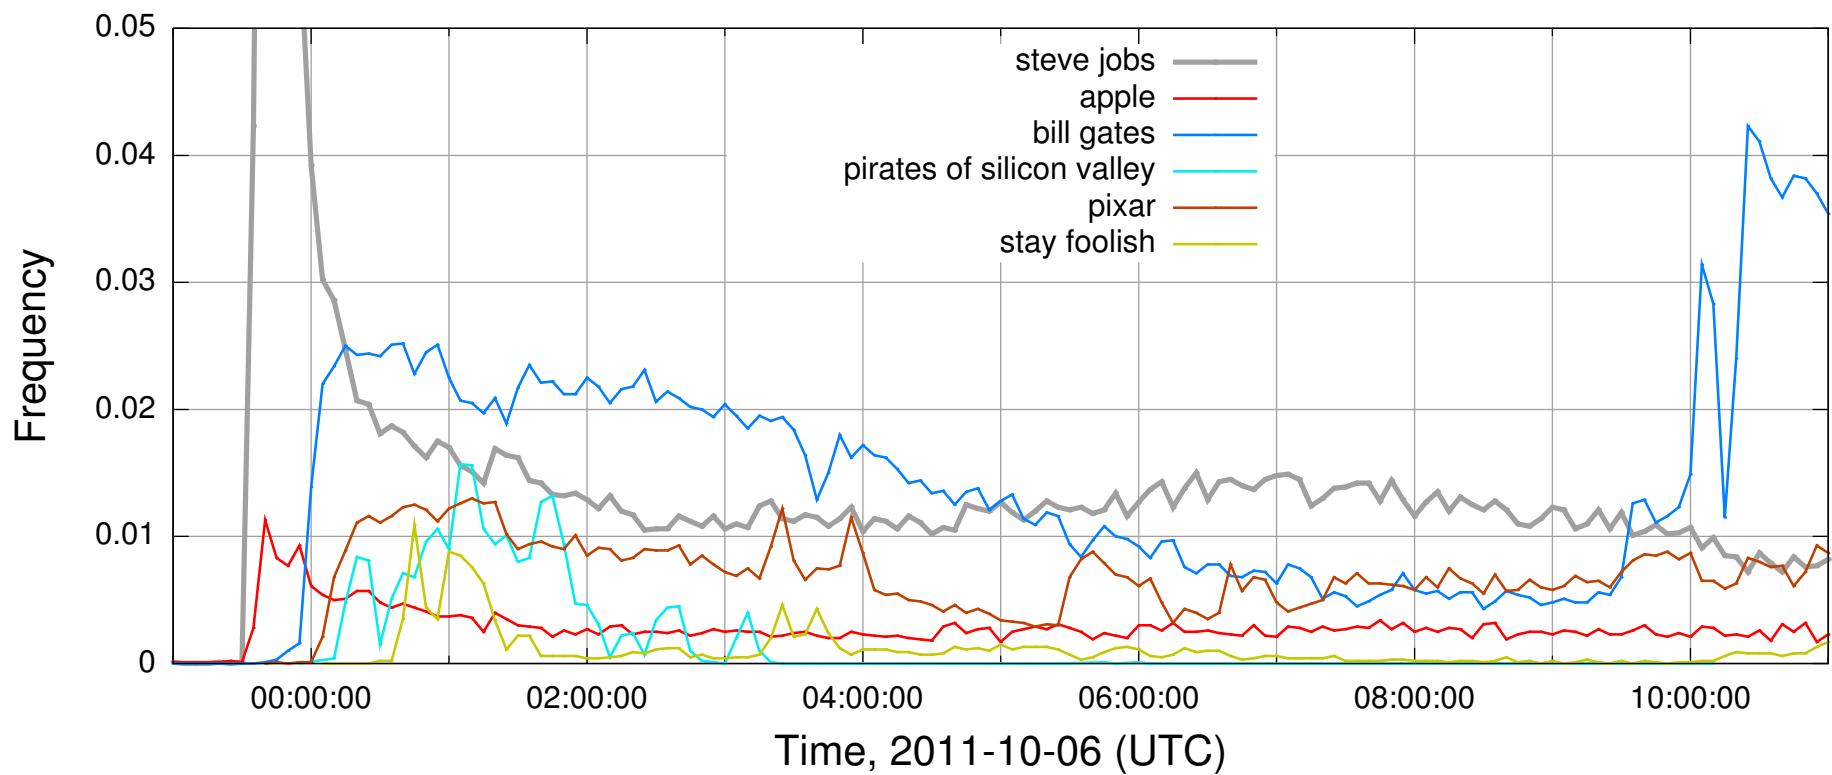
#SochiFail

**Sochi 2014**  @Sochi2014**Sochi Olympics 2014** @2014Sochi**Игры Сочи 2014**  @sochi2014_ru**Sochi Problems** @SochiProblem**NYT Olympics** @SochiNYT**Sochi Problems** @SochiProblems

Search all people for sochi

[Reply](#) [Retweet](#) [Favorite](#) [More](#)[Reply](#) [Retweet](#) [Favorite](#) [More](#)[Reply](#) [Retweet](#) [Favorite](#) [More](#)

Case Studies: Steve Jobs passes away



Initial Implementation

Algorithm: Co-occurrences within query sessions

Implementation: Pig scripts over query logs on HDFS

Problem: Query suggestions were several hours old!

Why?

Log collection lag

Hadoop scheduling lag

Hadoop job latencies

We need real-time processing!
(stay tuned next time for actual solution)

real-time

vs.

online

vs.

streaming

What is a data stream?

Sequence of items:

Structured (e.g., tuples)

Ordered (implicitly or timestamped)

Arriving continuously at high volumes

Sometimes not possible to store entirely

Sometimes not possible to even examine all items

Applications

Network traffic monitoring

Datacenter telemetry monitoring

Sensor networks monitoring

Credit card fraud detection

Stock market analysis

Online mining of click streams

Monitoring social media streams

What exactly do you do?

“Standard” relational operations:

Select

Project

Transform (i.e., apply custom UDF)

Group by

Join

Aggregations

What else do you need to make this “work”?

Issues of Semantics

Group by... aggregate

When do you stop grouping and start aggregating?

Joining a stream and a static source

Simple lookup

Joining two streams

How long do you wait for the join key in the other stream?

Joining two streams, group by and aggregation

When do you stop joining?

What's the solution?

Windows

Windows restrict processing scope:

Windows based on ordering attributes (e.g., time)

Windows based on item (record) counts

Windows based on explicit markers (e.g., punctuations)

Windows on Ordering Attributes

Assumes the existence of an attribute that defines the order of stream elements (e.g., time)

Let T be the window size in units of the ordering attribute



Windows on Counts

Window of size N elements (sliding, tumbling) over the stream



Windows from “Punctuations”

Application-inserted “end-of-processing”

Example: stream of actions... “end of user session”

Properties

Advantage: application-controlled semantics

Disadvantage: unpredictable window size (too large or too small)

Streams Processing Challenges

Inherent challenges

Latency requirements

Space bounds

System challenges

Load balancing

Unreliable and out-of-order message delivery

Fault-tolerance

Consistency semantics (at most once, exactly once, at least once)

Common Techniques

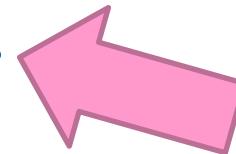


Streams Processing Challenges

Inherent challenges

Latency requirements

Space bounds



System challenges

Load balancing

Unreliable and out-of-order message delivery

Fault-tolerance

Consistency semantics (at most once, exactly once, at least once)

Space: The Final Frontier

Data streams are potentially infinite
Unbounded space!

How do we get around this?
Throw old data away
Accept some approximation

General techniques
Sampling
Hashing

Reservoir Sampling

Task: select s elements from a stream of size N with uniform probability

N can be very very large

We might not even know what N is! (infinite stream)

Solution: Reservoir sampling

Store first s elements

For the k -th element thereafter, keep with probability s/k
(randomly discard an existing element)

Example: $s = 10$

Keep first 10 elements

11th element: keep with $10/11$

12th element: keep with $10/12$

...

Reservoir Sampling: How does it work?

Example: $s = 10$

Keep first 10 elements

11th element: keep with $10/11$

If we decide to keep it: sampled uniformly by definition
probability existing item is discarded: $10/11 \times 1/10 = 1/11$
probability existing item survives: $10/11$

General case: at the $(k + l)$ th element

Probability of selecting each item up until now is s/k

Probability existing item is discarded: $s/(k+l) \times l/s = l/(k+l)$

Probability existing item survives: $k/(k+l)$

Probability each item survives to $(k + l)$ th round:

$$(s/k) \times k/(k+l) = s/(k+l)$$

Hashing for Three Common Tasks

Cardinality estimation

What's the cardinality of set S ?

How many unique visitors to this page?

HashSet HLL counter

Set membership

Is x a member of set S ?

Has this user seen this ad before?

HashSet Bloom Filter

Frequency estimation

How many times have we observed x ?

How many queries has this user issued?

HashMap CMS

HyperLogLog Counter

Task: cardinality estimation of set
`size()` → number of unique elements in the set

Observation: hash each item and examine the hash code

On expectation, 1/2 of the hash codes will start with 1

On expectation, 1/4 of the hash codes will start with 01

On expectation, 1/8 of the hash codes will start with 001

On expectation, 1/16 of the hash codes will start with 0001

...

How do we take advantage of this observation?

Bloom Filters

Task: keep track of set membership

$\text{put}(x) \rightarrow$ insert x into the set

$\text{contains}(x) \rightarrow$ yes if x is a member of the set

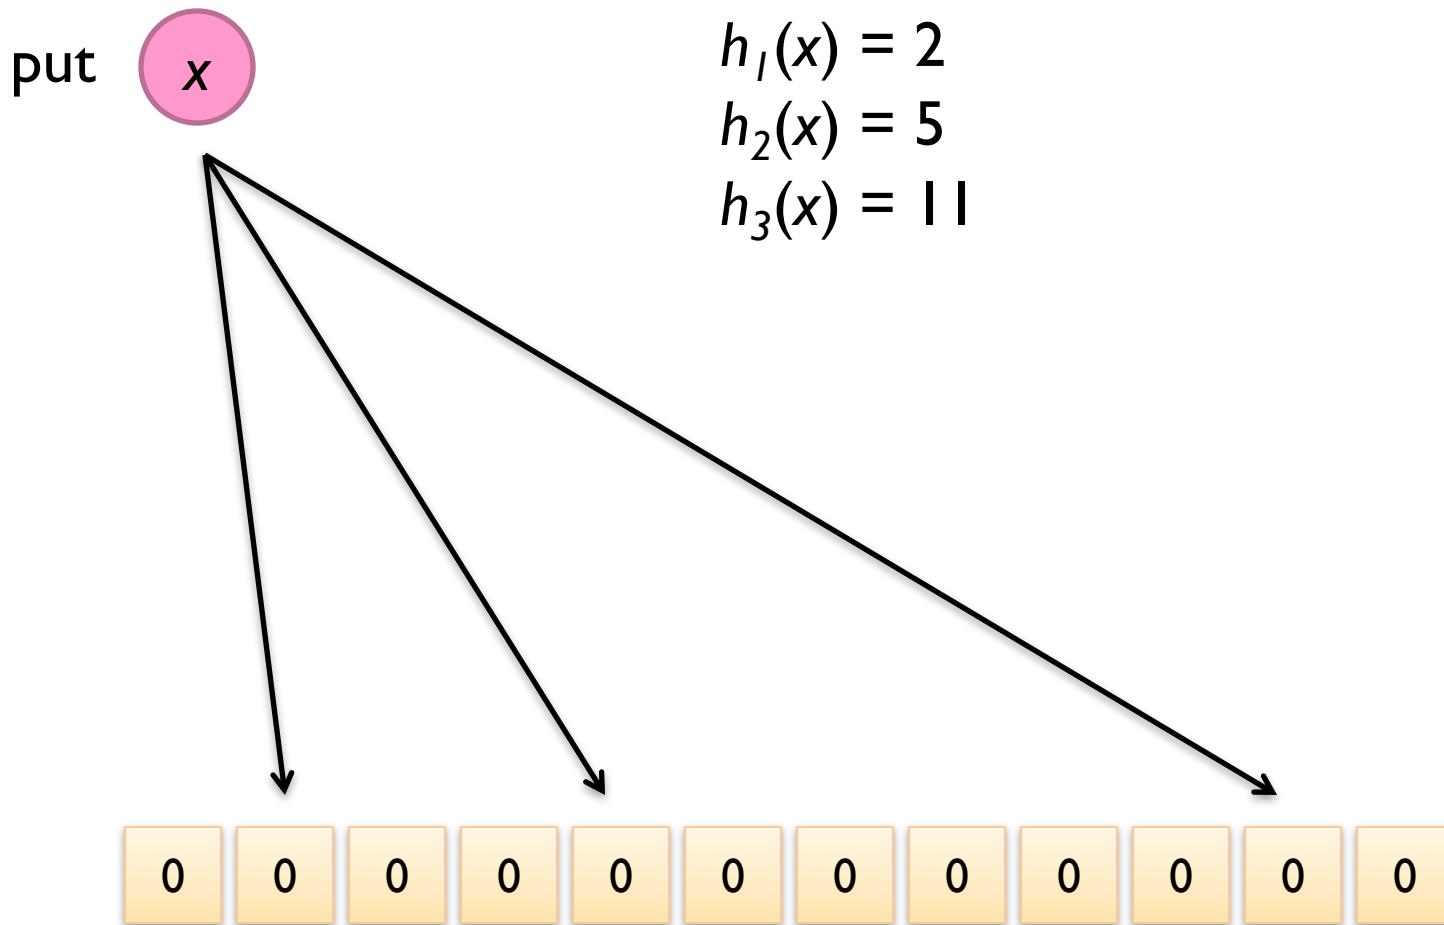
Components

m -bit bit vector

k hash functions: $h_1 \dots h_k$



Bloom Filters: put

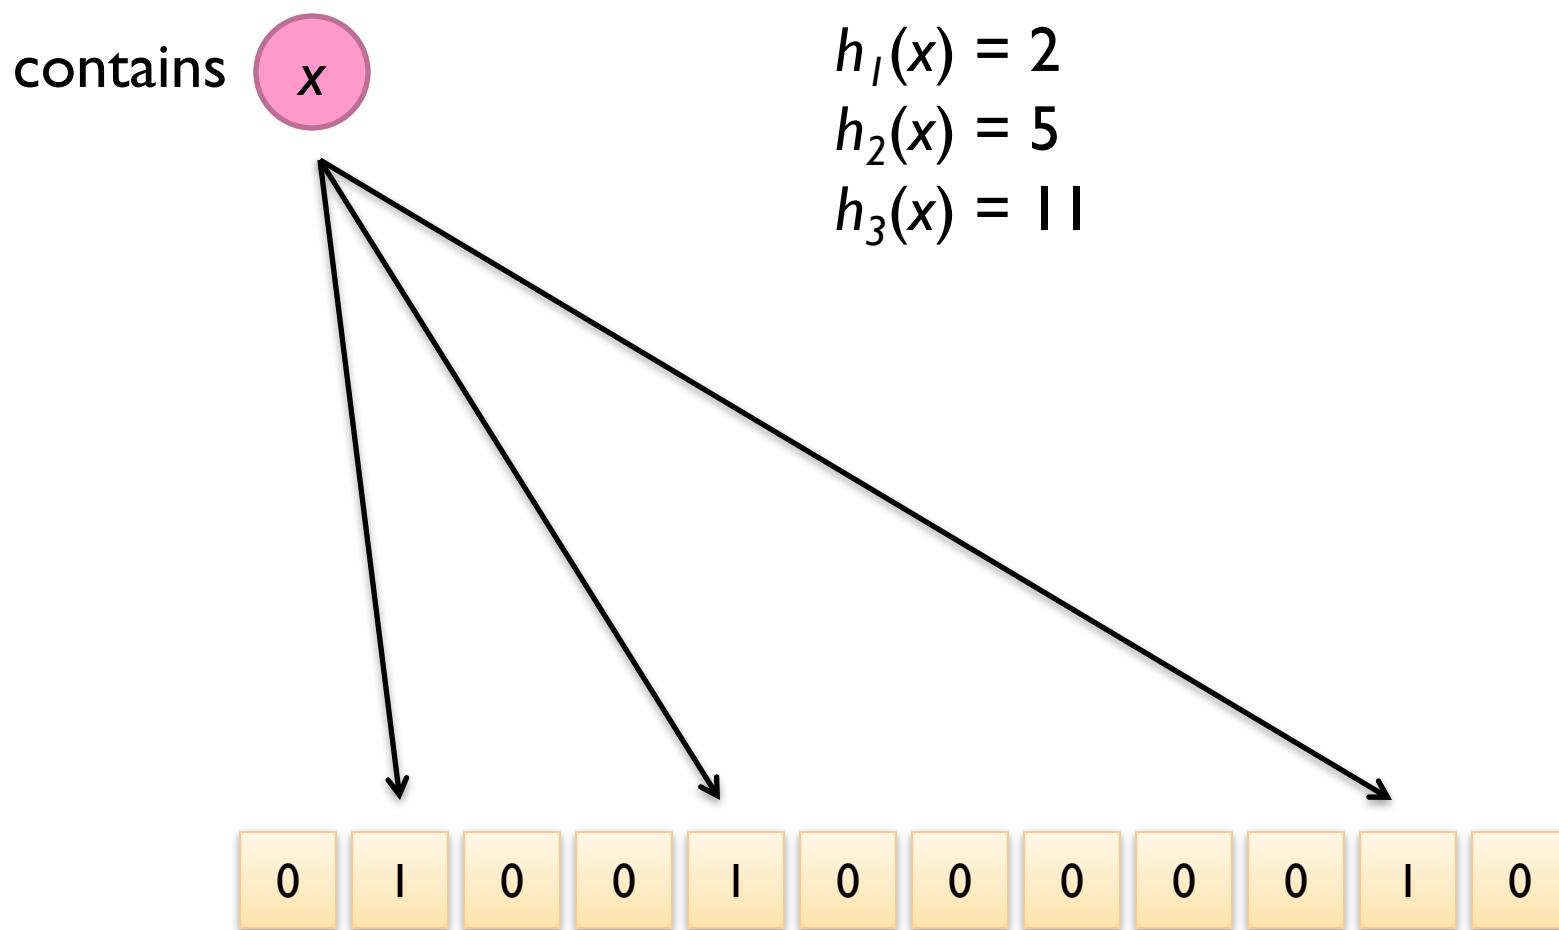


Bloom Filters: put

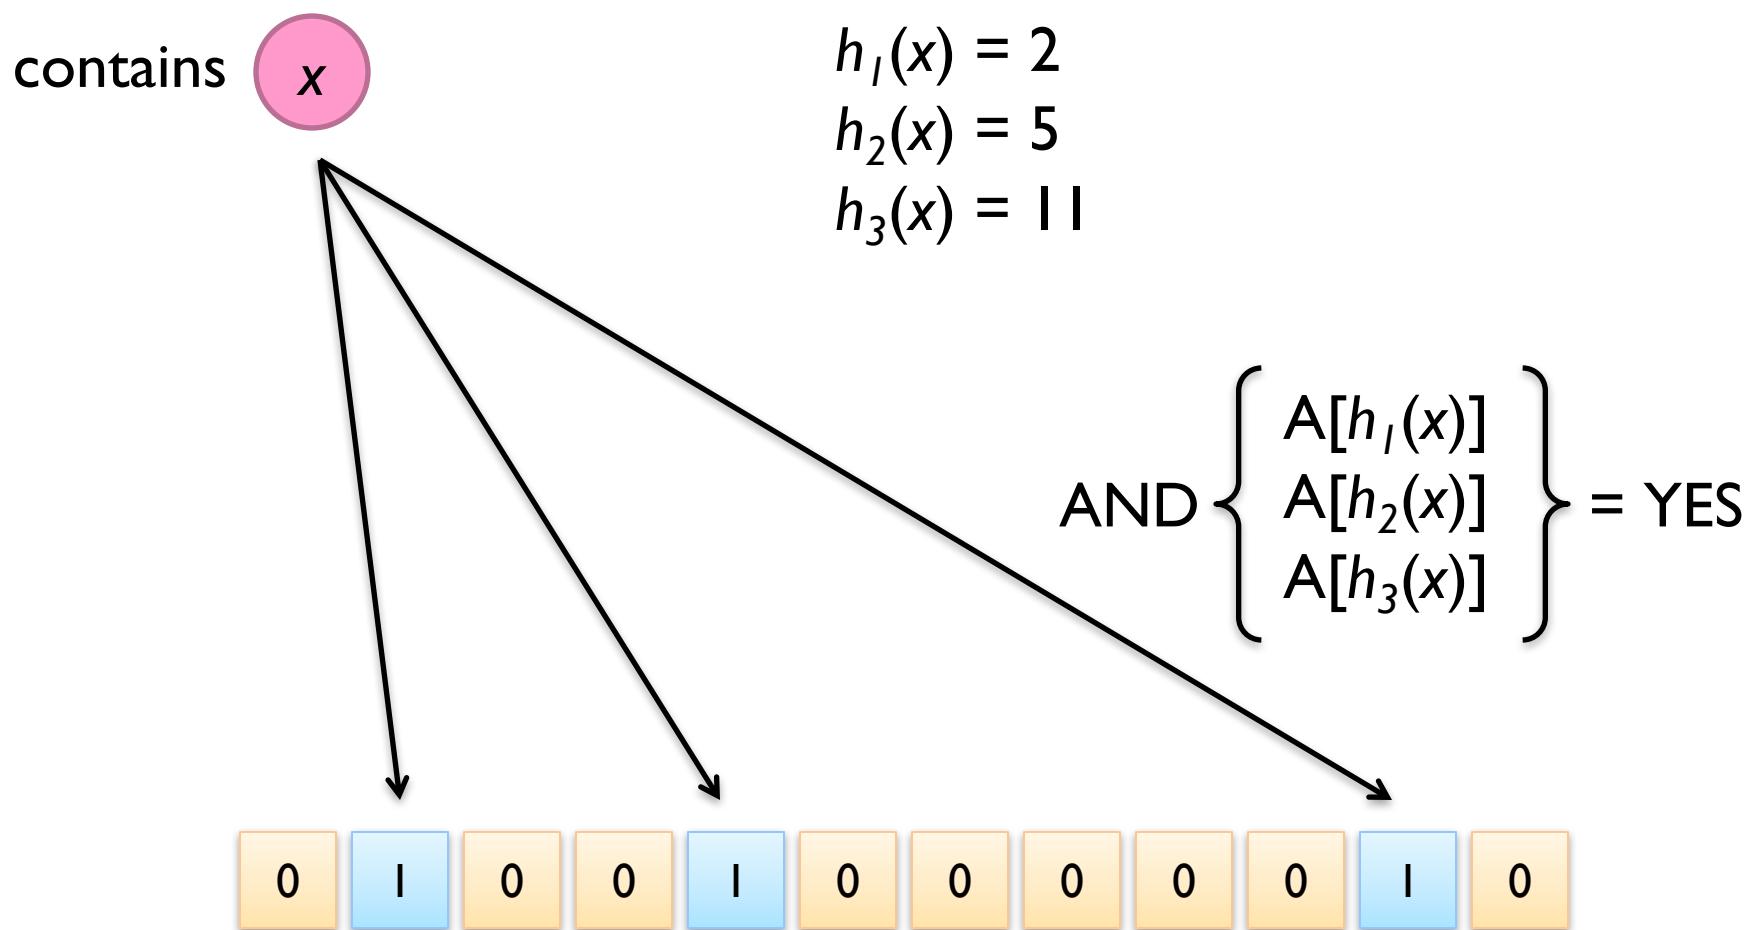
put



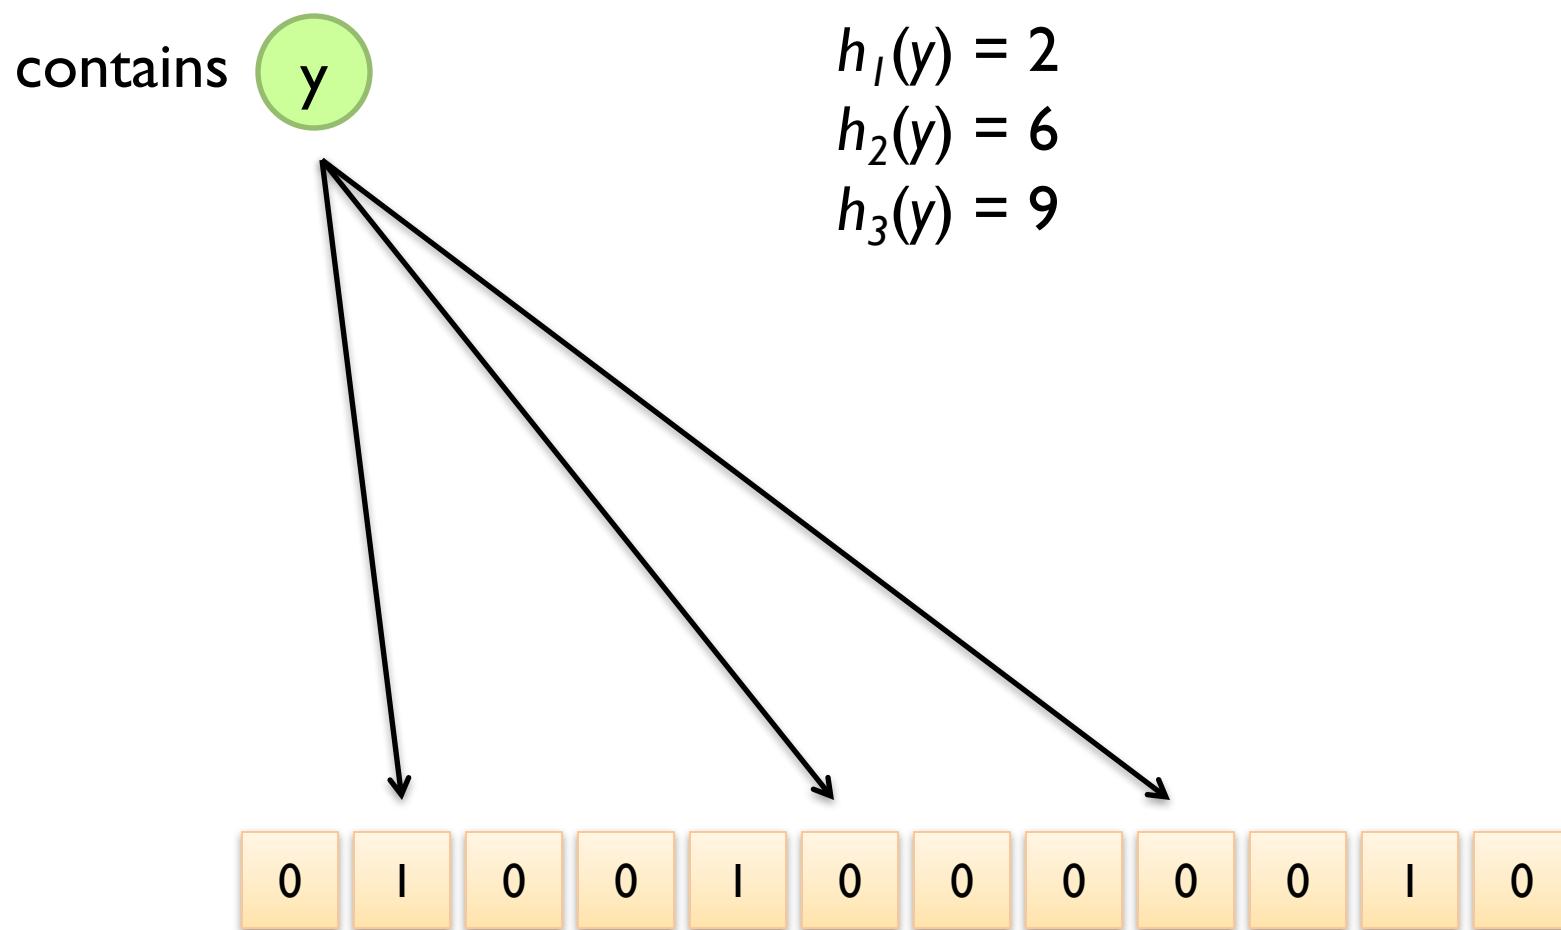
Bloom Filters: contains



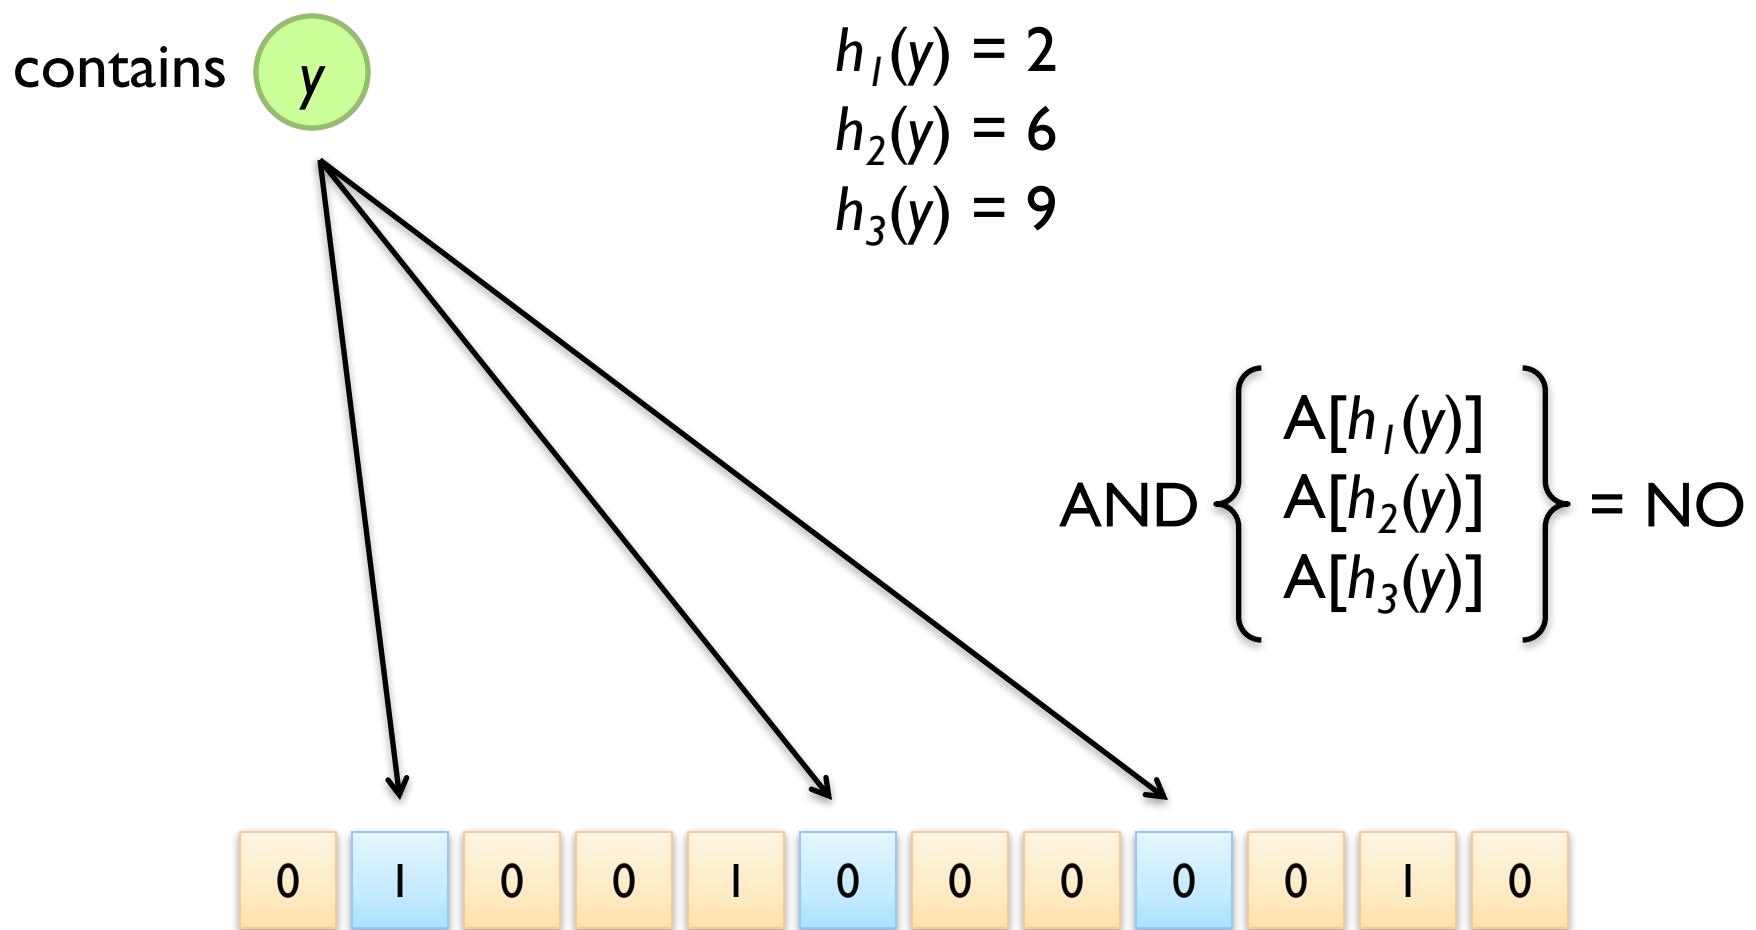
Bloom Filters: contains



Bloom Filters: contains



Bloom Filters: contains



What's going on here?

Bloom Filters

Error properties: `contains(x)`

False positives possible

No false negatives

Usage

Constraints: capacity, error probability

Tunable parameters: size of bit vector m , number of hash functions k

Count-Min Sketches

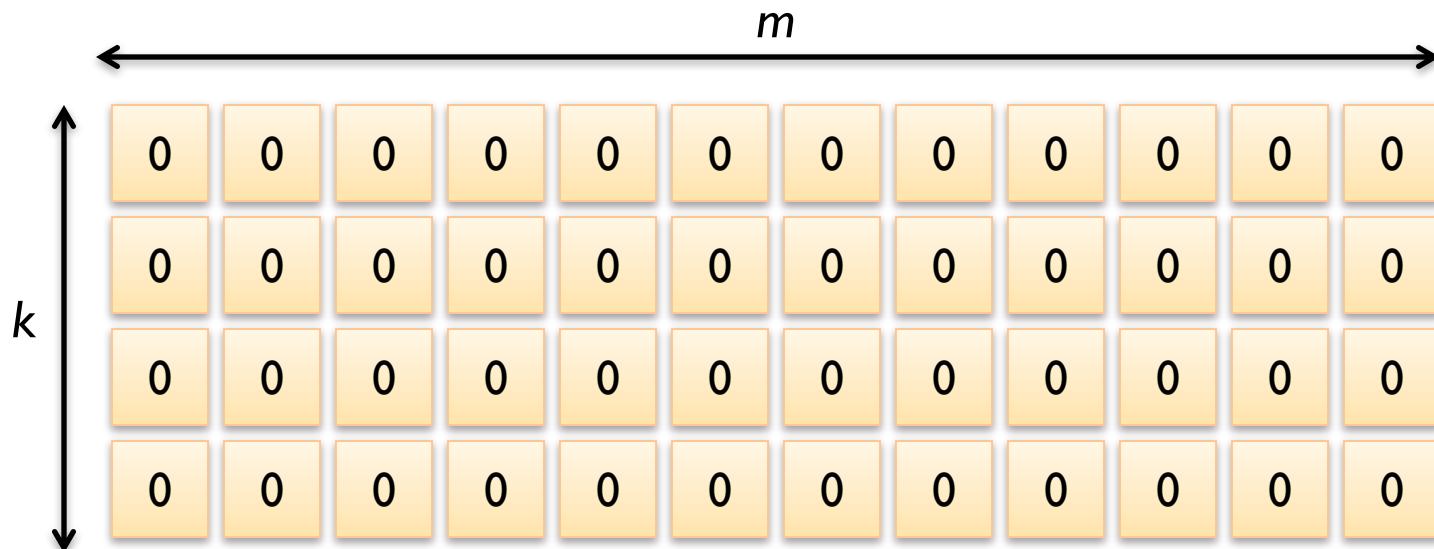
Task: frequency estimation

`put(x) → increment count of x by one`
`get(x) → returns the frequency of x`

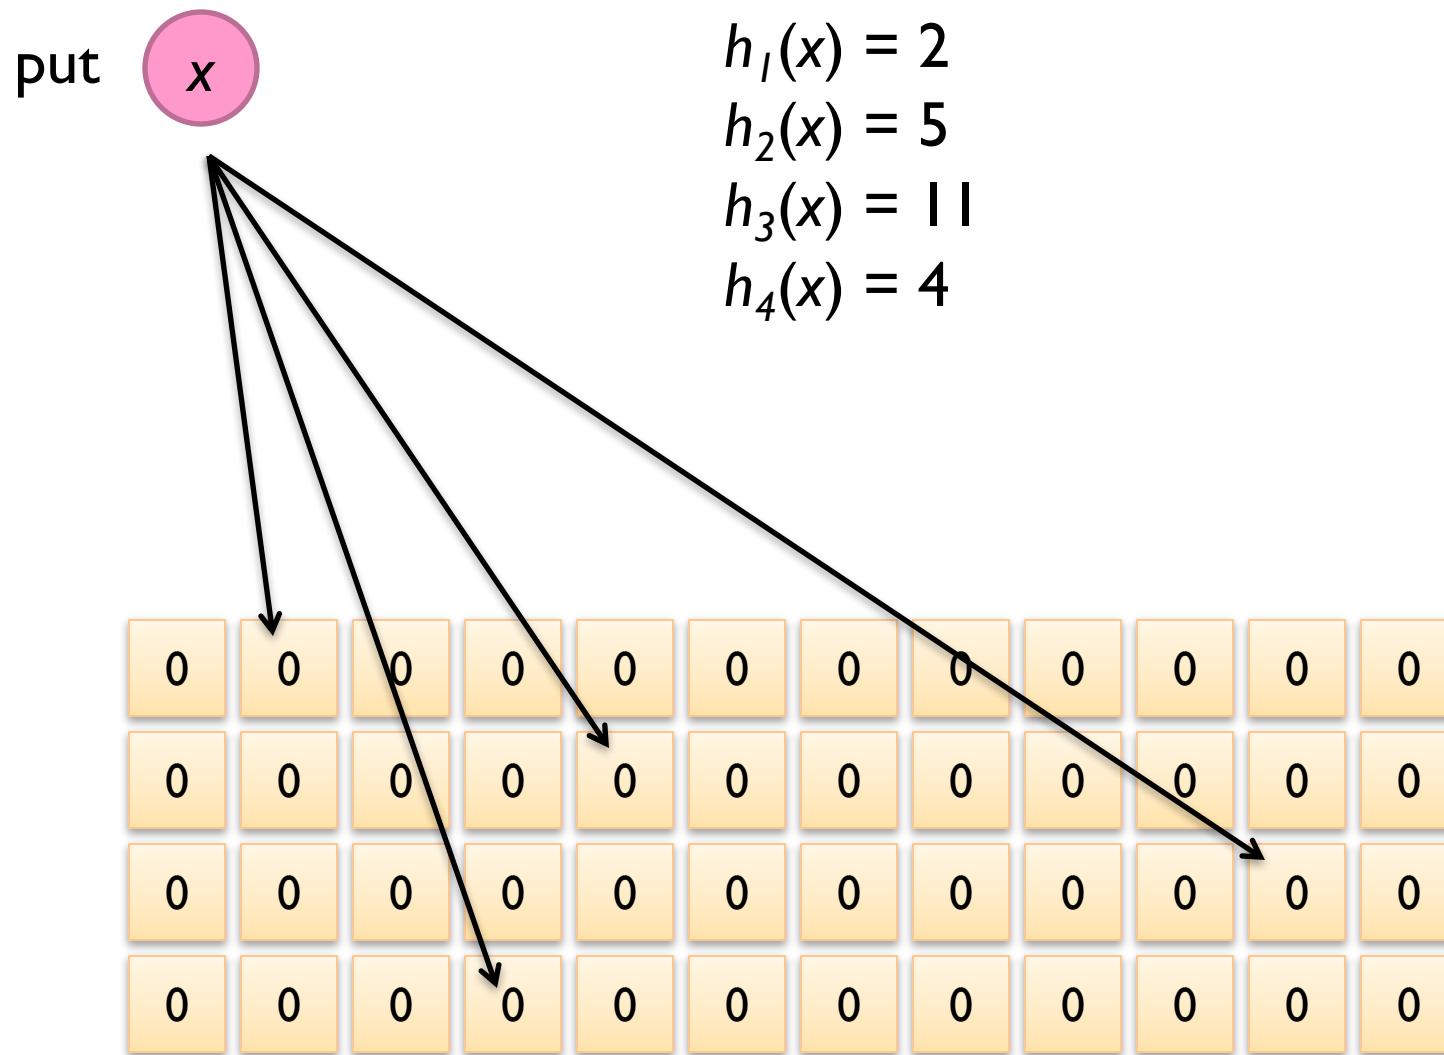
Components

m by k array of counters

k hash functions: $h_1 \dots h_k$



Count-Min Sketches: put

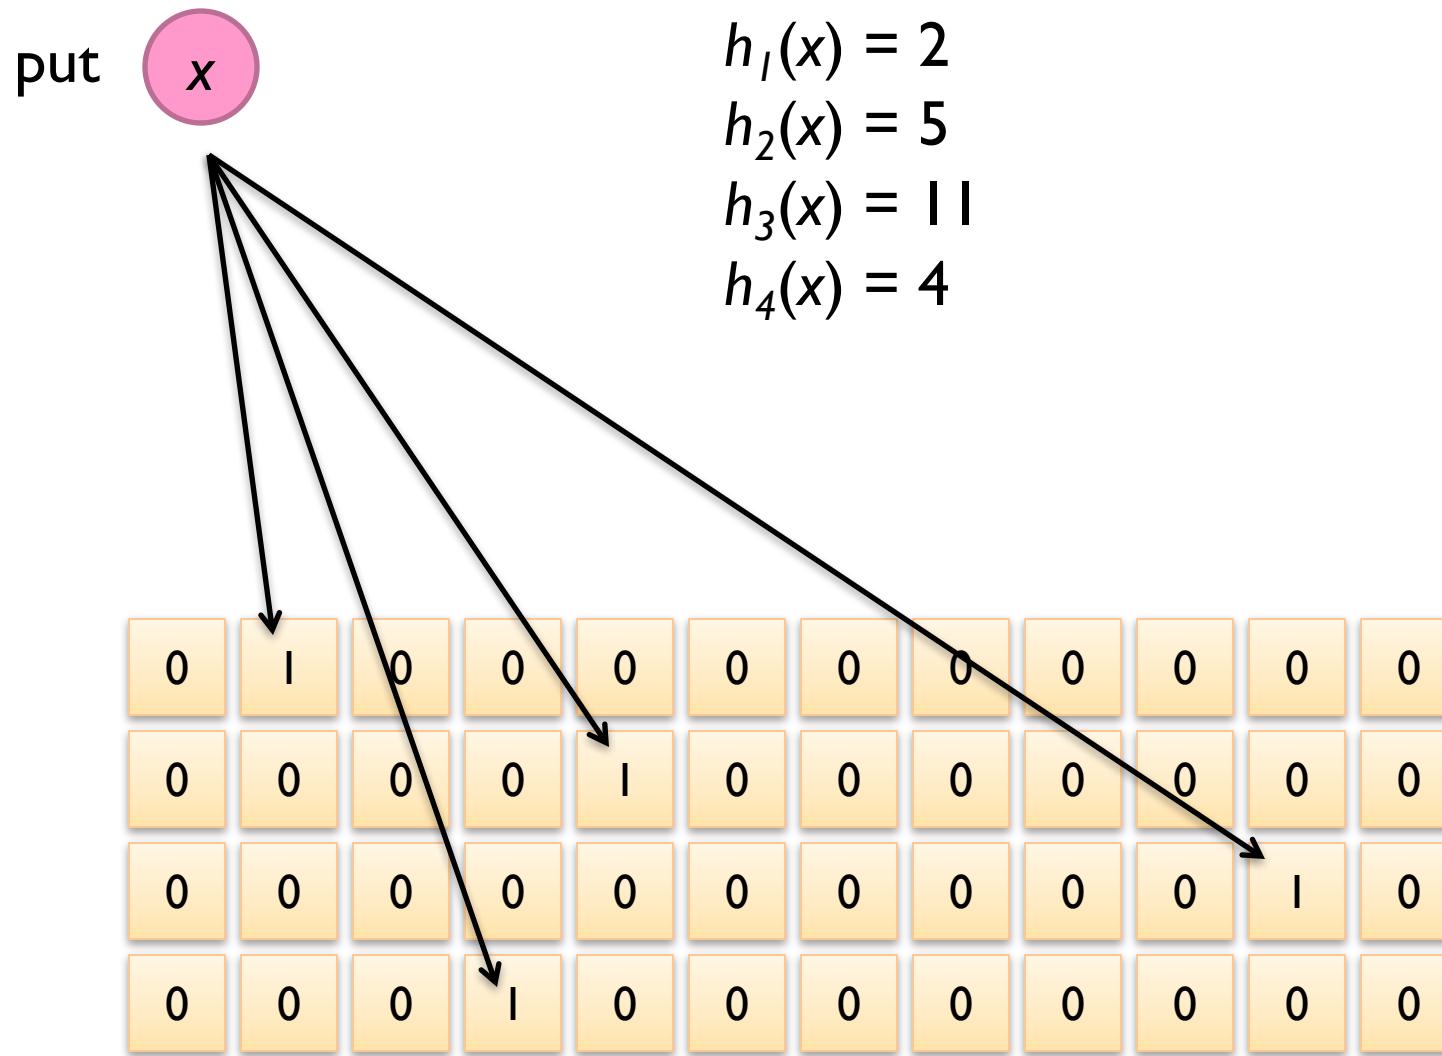


Count-Min Sketches: put

put

X

Count-Min Sketches: put

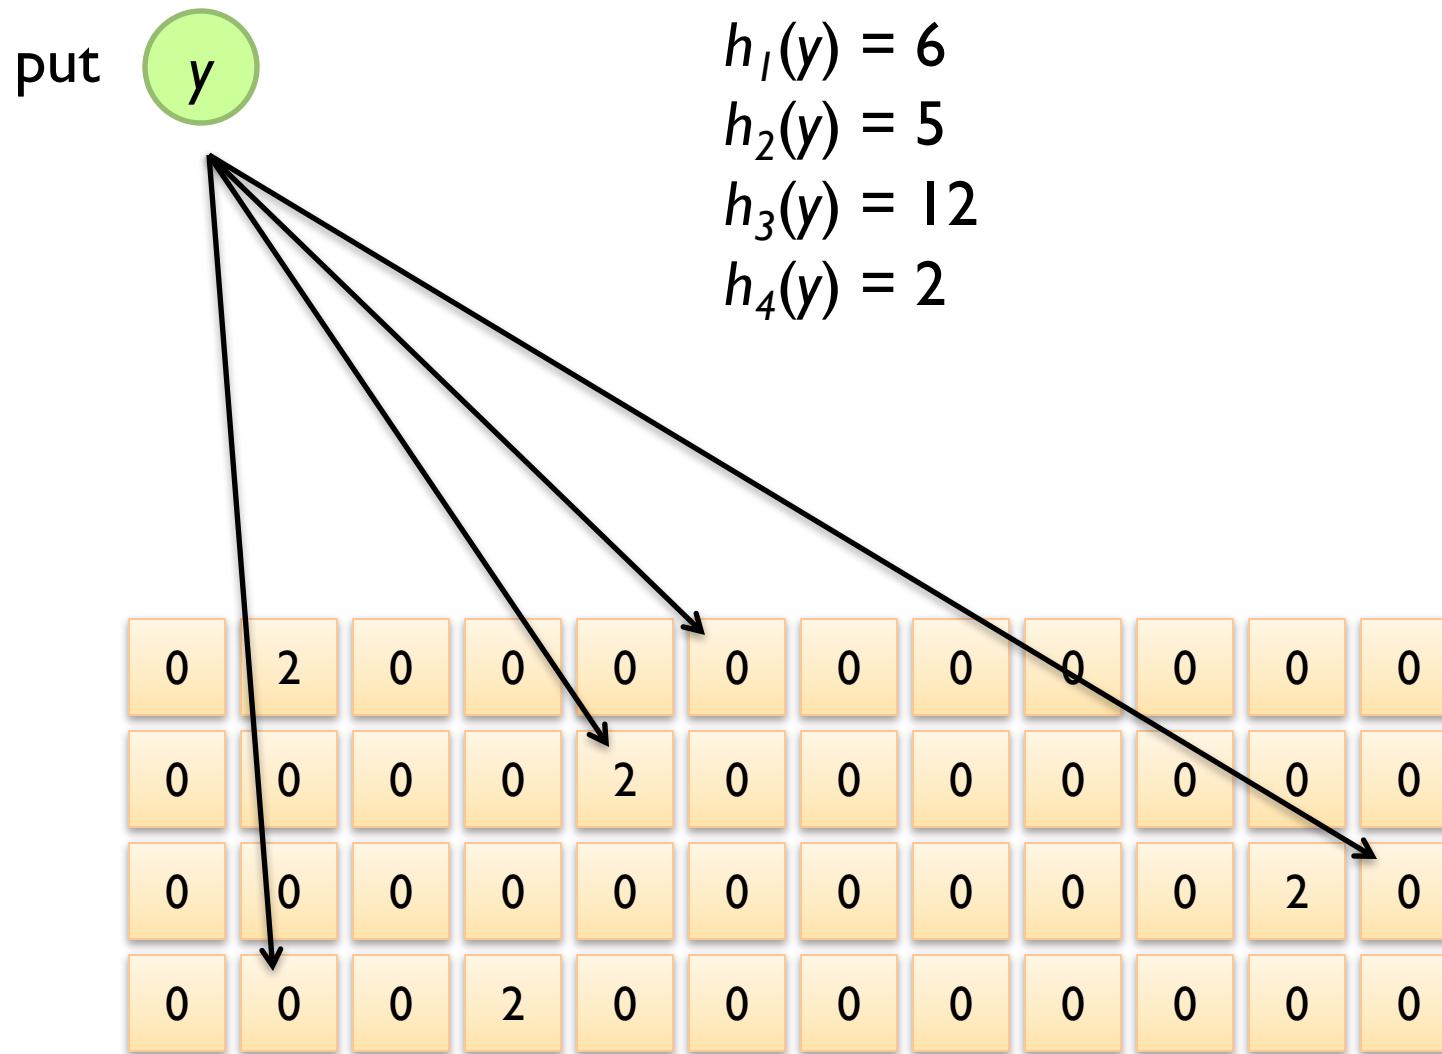


Count-Min Sketches: put

put



Count-Min Sketches: put

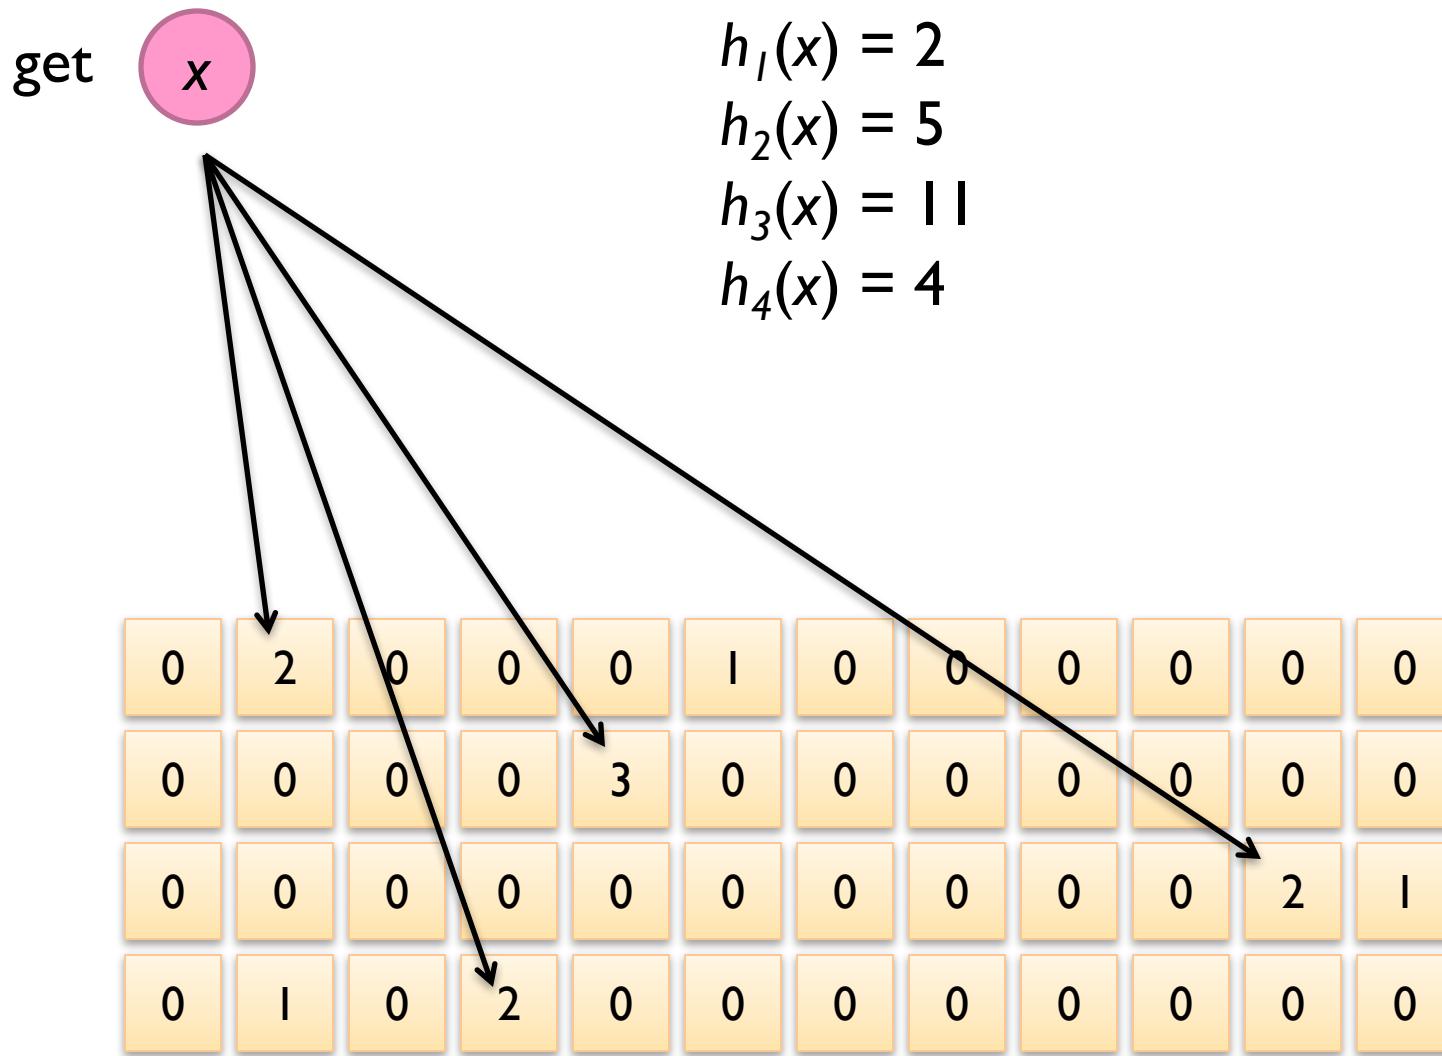


Count-Min Sketches: put

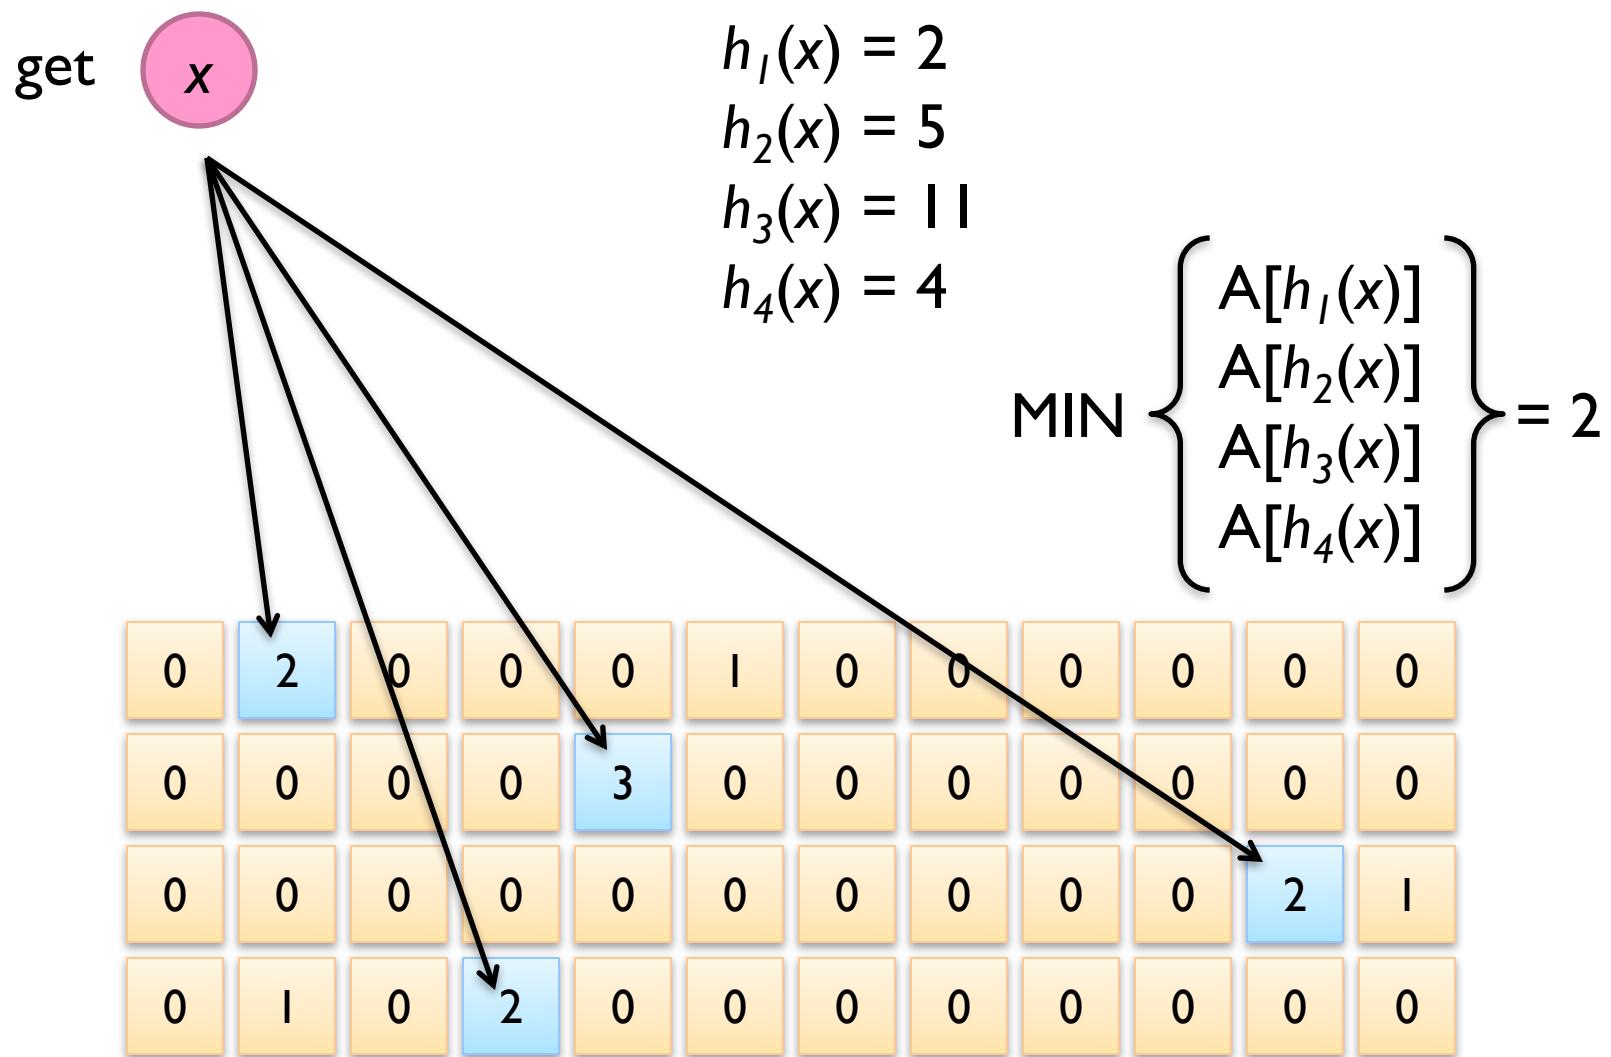
put

y

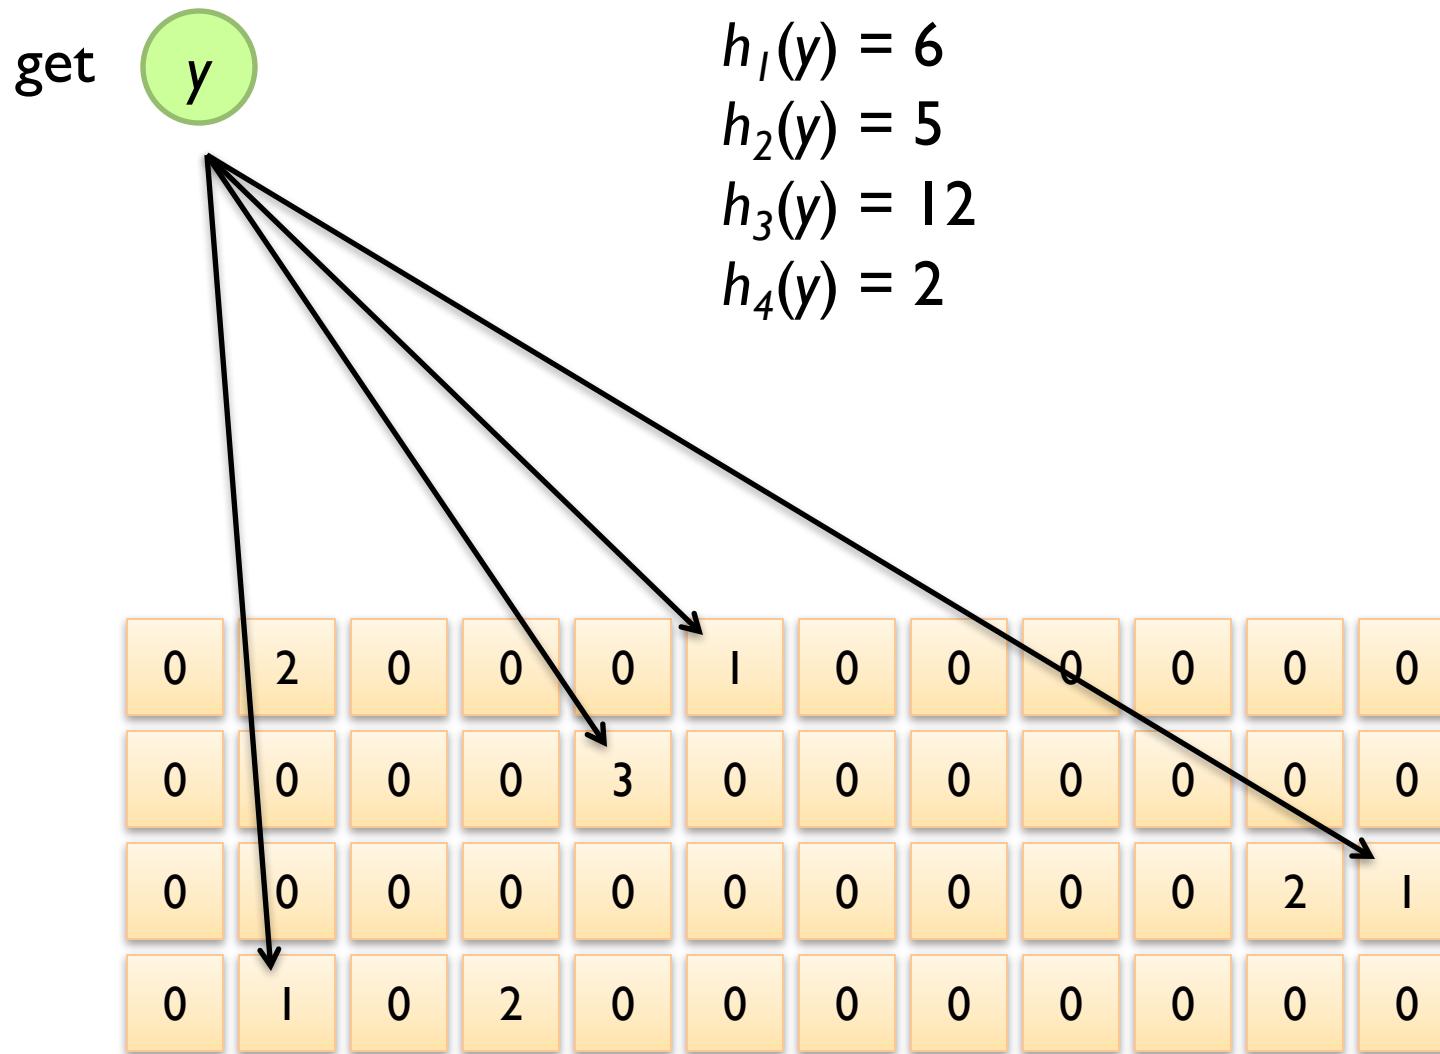
Count-Min Sketches: get



Count-Min Sketches: get



Count-Min Sketches: get



Count-Min Sketches

Error properties: $\text{get}(x)$

Reasonable estimation of heavy-hitters

Frequent over-estimation of tail

Usage

Constraints: number of distinct events, distribution of events, error bounds

Tunable parameters: number of counters m and hash functions k , size of counters

Hashing for Three Common Tasks

Cardinality estimation

What's the cardinality of set S ?

How many unique visitors to this page?

HashSet HLL counter

Set membership

Is x a member of set S ?

Has this user seen this ad before?

HashSet Bloom Filter

Frequency estimation

How many times have we observed x ?

How many queries has this user issued?

HashMap CMS



Next time:
Stream Processing Architectures

A photograph of a traditional Japanese rock garden. In the foreground, a gravel path is raked into fine, parallel lines. Several large, dark, irregular stones are scattered across the garden. A small, shallow pond is visible in the middle ground, surrounded by more stones and some low-lying green plants. In the background, there are more stones, some small trees, and the wooden buildings of a residence with tiled roofs.

Questions?